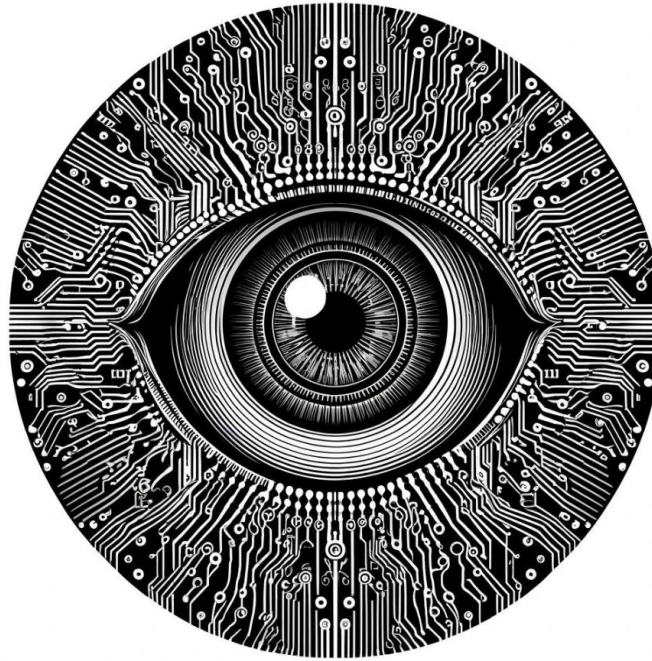


Digital Images in PIL and NumPy



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Understand image tensors as multidimensional arrays
- Gain familiarity with PIL images and NumPy arrays

The unsigned integer (uint8) datatype






Image from [Why Are There 8 Bits in a Byte?](#)

Digital image representation for grayscale images



what we see



what the computer “sees”

RGB images as “tensors” (stacks of matrices)




Image from [Google Earth](#)

What a human sees

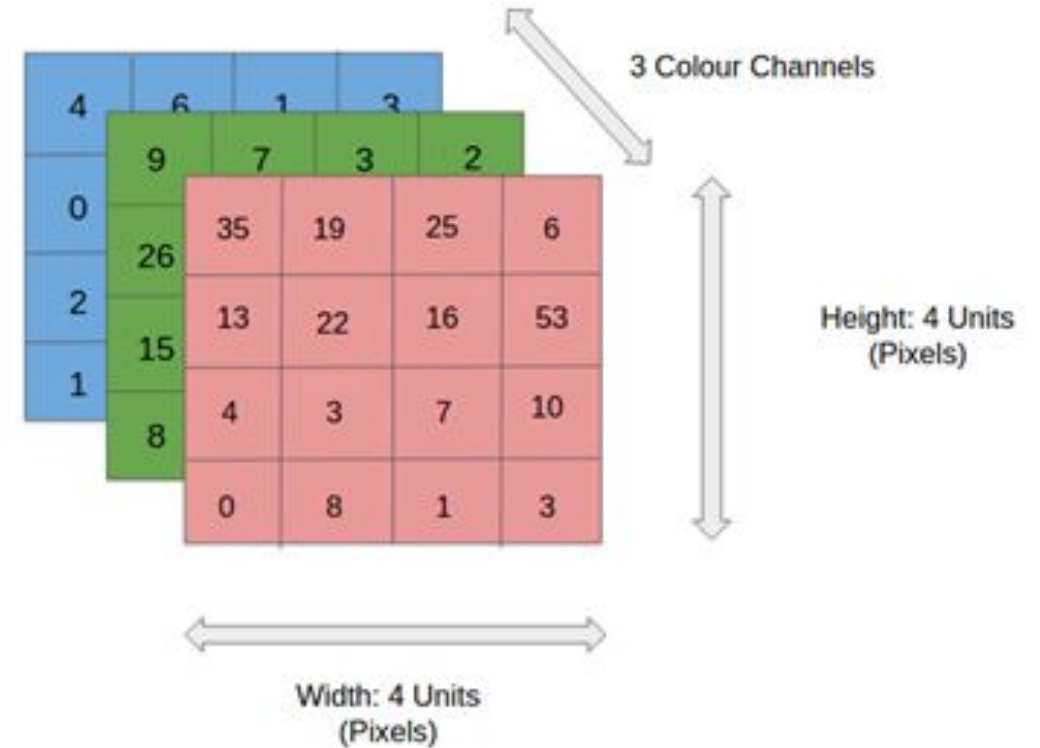


Image [source](#)

What the computer ‘sees’

The Python Image Library (PIL)



real size is $W = H = 28$

3

```
from PIL import Image
```

```
# Open image from file
```

```
pil_image = Image.open("path_to_image.jpg")
```

```
# Check the mode of the image, 'L' is for grayscale
```

```
# pil_image.get
```

```
print(pil_image.mode)
```

```
# Will print (28, 28), width first, height second
```

```
print(pil_image.size)
```

```
# Convert a NumPy ndarray to PIL
```

```
pil_image = Image.fromarray(np_array)
```

```
# Save image to storage as a PNG file
```

```
pil_image.save("path_to_image.png")
```

```
# Convert to numpy and upscale the image for visualization
```

```
import matplotlib.pyplot as plt
```

```
plt.imshow(pil_image)
```

NumPy's N-dimensional array (ndarray)



```
# Convert a PIL image to a Numpy array
```

```
np_array = np.array(pil_image)
```

```
# Check the shape of the array, prints (28, 28)
```

```
print(np_array.shape)
```

```
# Prints the data type, prints np.uint8
```

```
print(np_array.dtype)
```

```
# 255 minus current pixel values with NumPy's broadcasting
```

```
np_array_neg = 255 - np_array
```

```
# Show upscaled version on matplotlib
```

```
plt.imshow(np_array_neg)
```


Red Green Blue (RGB) Images in NumPy

```
import numpy as np

# Create a 3x3 RGB image array
rgb_image = np.array([
    # First row of pixels
    [[255, 0, 0],      # Pure red pixel
     [0, 255, 0],      # Pure green pixel
     [0, 0, 255]],     # Pure blue pixel

    # Second row of pixels
    [[255, 255, 0],    # Yellow pixel (red + green)
     [255, 0, 255],    # Magenta pixel (red + blue)
     [0, 255, 255]],   # Cyan pixel (green + blue)


    # Third row of pixels
    [[255, 255, 255],  # White pixel (all colors maximum)
     [128, 128, 128], # Gray pixel (all colors at half intensity)
     [0, 0, 0]],       # Black pixel (all colors minimum)
])
```



Indexing values in NumPy

```
# We select a row and a column
# Indexing starts at 0, as in standard Python
# We index first with height, then width
pixel = rgb_image[2, 0]

# Will show [255, 255, 255]
# Which index should we use to print magenta's values?
print(f"RGB values: {pixel}")
```



Intensities for RGB images in NumPy


Original Image



Red Channel



Green Channel



Blue Channel



```
import matplotlib.pyplot as plt
```


```
# Intensity image from the R channel  
plt.imshow(rgb_np_array[:, :, 0])
```

```
# Intensity image from the G channel  
plt.imshow(rgb_np_array[:, :, 1])
```


```
# Intensity image from the B channel  
plt.imshow(rgb_np_array[:, :, 2])
```

False colors for intensity images

Original Image




Red Channel



Green Channel



Blue Channel



```
import skimage.io as io
```

```
# Converts from JPEG to NumPy array  
rgb_np_array =  
io.imread("gujarat_image_url.jpeg")
```

```
# Images in NumPy are saved as  
# Height (H), Width(W), Channels (C)  
# Prints (1200, 1800, 3)  
print(rgb_np_array.shape)
```

```
# Prints uint8  
print(np_array.dtype)
```

```
# Intensity image with Red colormap  
plt.imshow(rgb_np_array[:, :, 0], cmap='Reds')  
plt.imshow(rgb_np_array[:, :, 1], cmap='Greens')  
plt.imshow(rgb_np_array[:, :, 2], cmap='Blues')
```

Image arithmetic in NumPy


```
import numpy as np
from PIL import Image

# Create arrays of of zeros
red = np.zeros((600, 600))
green = np.zeros((600, 600))
blue = np.zeros((600, 600))

# Set sections to maximum intensity
red[150:350, 150:350] = 255
green[200:400, 200:400] = 255
blue[250:450, 250:450] = 255

red_img = Image.fromarray(red).convert("L")
green_img = Image.fromarray(green).convert("L")
blue_img = Image.fromarray(blue).convert("L")

# Merge channels
square_img = Image.merge("RGB", (red_img, green_img, blue_img))
square_img.show()
```



Code and image from [source](#)

The dimensions of PIL and NumPy arrays



NumPy image tensors follow the format: $[H, W, C]$ where:

- H = height in pixels
- W = width in pixels
- C = channels (e.g., 1 for grayscale, 3 for RGB)

PIL image tensors follow the format $[W, H]$ with mode

- mode is "L" for grayscale images
- mode is "RGB" for color images

Summary

Digital images are represented as numeric pixel values

- uint8 is memory-efficient and commonly used to store images

We use different Python libraries for image processing

- PIL for image loading and saving to JPEG or PNG
- NumPy for numerical operations and visualization with matplotlib

PIL and NumPy use different formats to represent images

- Height, width, and number of channels is the standard for NumPy
- Width, height and mode is the standard for PIL

References

NumPy's N-dimensional array

- <https://numpy.org/doc/2.1/reference/arrays.ndarray.html>

Intro to PIL and NumPy for image manipulation

- <https://realpython.com/image-processing-with-the-python-pillow-library/>

SPONSORED BY THE