


Image Tensors in PyTorch



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Gain familiarity with `torch` tensors
- Understand the need for data type transformation and scaling
- Understand how to move `torch` tensors between devices

Images as tensors




Image from [Google Earth](#)

What a human sees





Image [source](#)

What the computer 'sees'

Why do we use PyTorch tensors?



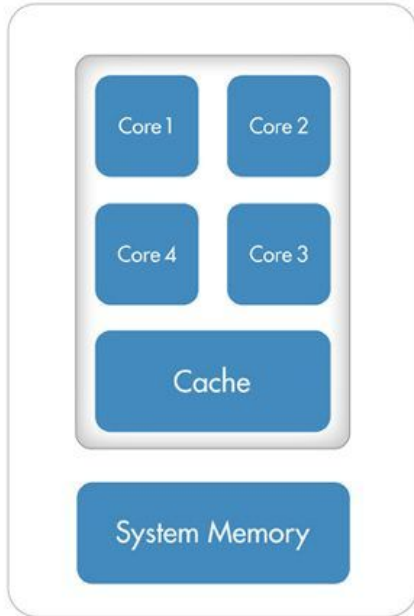
1. Can run on the Graphics Processing Unit (GPU) or Tensor Processing Unit (TPU) speeding up training and inference



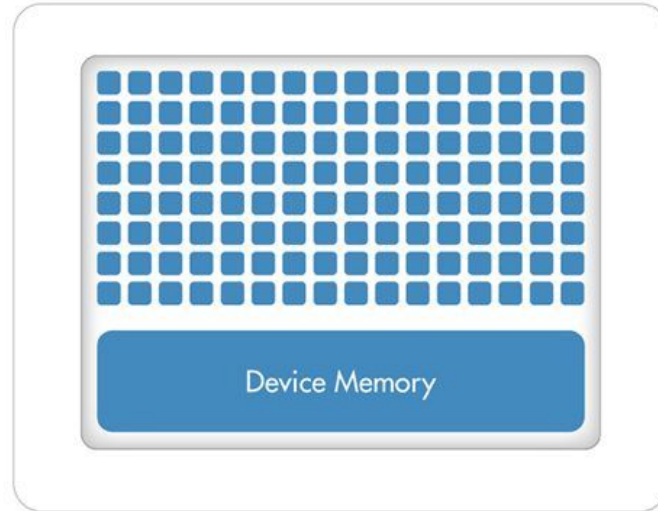
2. Do automatic gradient computation (autograd), enabling us to update the weights of models with automatically computed derivatives of the loss with respect to the weights

CPUs vs GPUs

CPU (Multiple Cores)



GPU (Hundreds of Cores)



Moving tensors to the GPU




```
import torch.nn as nn

# A model with a single linear layer
model = nn.Linear(in_features = 28 * 28,
                  out_features = 10)

device = 'cuda' if torch.cuda.is_available() else 'cpu'
# Move the model to the GPU
model = model.to(device)
torch_tensor_gray.to(device)

# Get prediction scores aka 'logits' (in GPU)
scores = model(torch_tensor_gray)
```

Moving tensors to the CPU for inspection/visualization



```
# Get prediction scores aka 'logits' (in GPU)
```

```
scores = model(torch_tensor_gray)
```

```
# Fails, scores are in the GPU
```

```
plt.show(scores)
```

```
# Moves a copy of the tensor to the cpu
```

```
# scores.cpu() is acceptable too
```

```
plt.show(scores.to('cpu'))
```

Speedup on matrix multiplication

Running matrix multiplication benchmarks...

Device: Tesla P100-PCIe-16GB

Matrix Size	CPU Time (s)	GPU Time (s)	Speedup
32x32	0.0000 ± 0.0000	0.0000 ± 0.0000	0.3x
2048x2048	0.0588 ± 0.0024	0.0026 ± 0.0001	23.0x
8192x8192	3.8277 ± 0.1194	0.1285 ± 0.0022	29.8x

[Quick benchmark on Kaggle](#)

Autograd to compute derivatives

```
import torch
```

```
# Create weight tensor with requires_grad=True
```

```
w1 = torch.tensor([4.0], requires_grad=True)
```

```
# Input features and ground truth do not require gradients (do not change)
```

```
x1 = torch.tensor([5.0], requires_grad=False)
```

```
y = torch.tensor([6.0], requires_grad=False)
```

```
# Let's create a simple computation graph
```

```
L = torch.abs(w1 * x1 - y)
```

```
# Will show 14.0, requires_grad = True
```

```
print(f"l = {L.item()}, requires_grad = {L.requires_grad}")
```

```
# Compute gradients
```


```
L.backward()
```

```
# Access the gradient
```


```
print(f"dL/dw1 = {w1.grad}") # Will be 5.0
```

$$L = |w_1 x_1 - y|$$

$$\frac{\partial L}{\partial w_1} = \begin{cases} x_1 & \text{if } w_1 x_1 - y \geq 0, \\ -x_1 & \text{if } w_1 x_1 - y < 0. \end{cases}$$



The dimensions of a torch image tensor



torch image tensors follow the format: [N, C, H, W] where:

- N = batch size (number of images)
- C = channels (e.g., 1 for grayscale, 3 for RGB)
- H = height in pixels
- W = width in pixels

*# Tensors shown to neural networks **include** the batch size*

```
print(torch_tensor_gray.unsqueeze(0).shape)
```

prints torch.Size([1, 1, 28, 28])

```
print(torch_tensor_gray.dim())
```

prints 3, the "rank" (# of dimensions)

PyTorch (aka torch) tensors

ToImage() converts NumPy arrays with H, W, C format to torch's C, W, H

```
transform = transforms.ToImage()  
image_tensor = transform(np_array)
```

Will print 3, 1200, 1800

```
print(image_tensor.shape)
```

Will print 1200, 1800, 3

```
print(image_tensor.permute(1, 2, 0).shape)
```


Error, matplotlib expects NumPy format

```
plt.imshow(image_tensor)
```

Channel order needs to be permuted to use matplotlib, as it expects NumPy's channel order

```
plt.imshow(image_tensor.permute(1, 2, 0))
```

Numerical stability



```
original_values = torch.tensor([255, 204,  
170], dtype=torch.uint8)
```

```
torch.exp(original_values)
```

$$\text{softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}} \quad \text{for } i = 1, \dots, n$$

Numerical stability example

Let's simulate a single pixel intensity in RGB

```
original_values = torch.tensor([255, 204, 170], dtype=torch.uint8)
print(f"Original values: {original_values}")
```

Convert to float and scale

```
float_values = original_values.float() / 255
print(f"Scaled values: {float_values}")
```

Example of multiplication stability

```
print(f"Original exp: {torch.exp(original_values)}") # Very large! inf!
print(f"Scaled exp: {torch.exp(float_values)}") # Stay small, bitte
```

Numerical stability

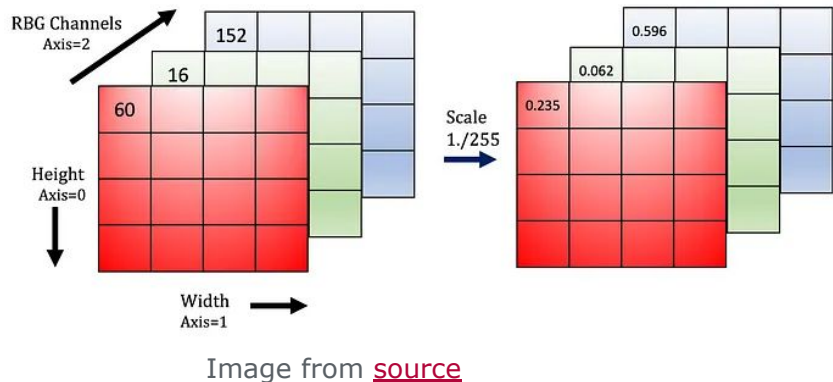
Original values: tensor([255, 204, 170], dtype=torch.uint8)

Scaled values: tensor([1.0000, 0.8000, 0.6667])

Original exp: tensor([inf, inf, inf])

Scaled exp: tensor([2.7183, 2.2255, 1.9477])

torch tensors with rescaled floating point values



1.0



0.5



0.0

```
# We convert the input to float and rescale it to  
# 0,1 to avoid overflow and get stable operations  
transform = transforms.Compose([  
    transforms.ToImage(),  
    transforms.ToDType(torch.float32, scale=True)])
```

```
# This transform would also work with a NumPy array  
image_tensor = transform(pil_image)
```

```
# Prints (1.0, 0.0)  
print(image_tensor.max(), image_tensor.min())
```

Summary

Tensors in PyTorch can run on the GPU and do automatic gradient computation

- We use `tensor_name.to(device)` to move tensors between CPU and GPU memory
- We use `result.backward()` to compute gradients on tensors with `requires_grad = True`

Common pitfall: different standards of tensor dimensions in NumPy vs PyTorch

- We use the B, C, H, W order in PyTorch; H, W, C in NumPy. The batch dimension corresponds to the number of tensors that we process at once on the device (GPU or CPU). The `toImage()` transform allows us to put the channel dimension on its proper position

Conversion to float and scaling tensors helps against numerical errors

- Operations like `torch.exp()` overflow without scaling

References

torch.Tensor.to

- <https://pytorch.org/docs/stable/generated/torch.Tensor.to.html>

ToImage

- <https://pytorch.org/vision/0.19/generated/torchvision.transforms.v2.ToImage.html>

SPONSORED BY THE



Federal Ministry
of Education
and Research