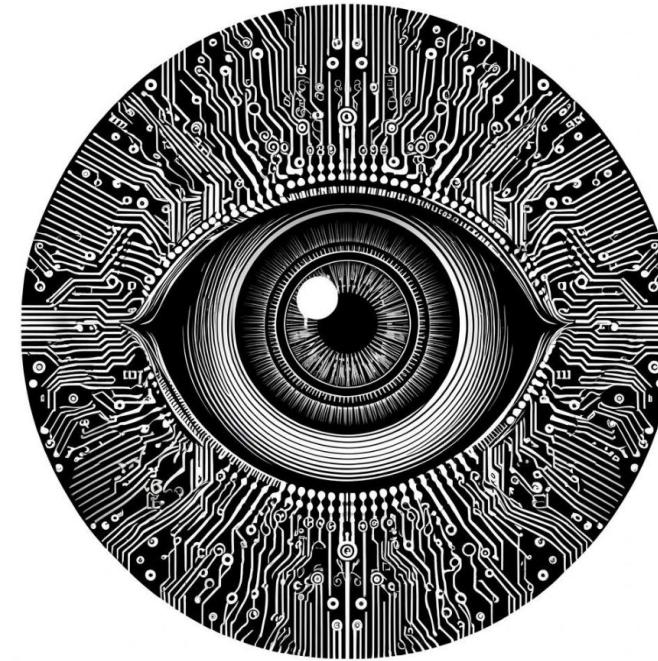


# Introduction to Neural Networks



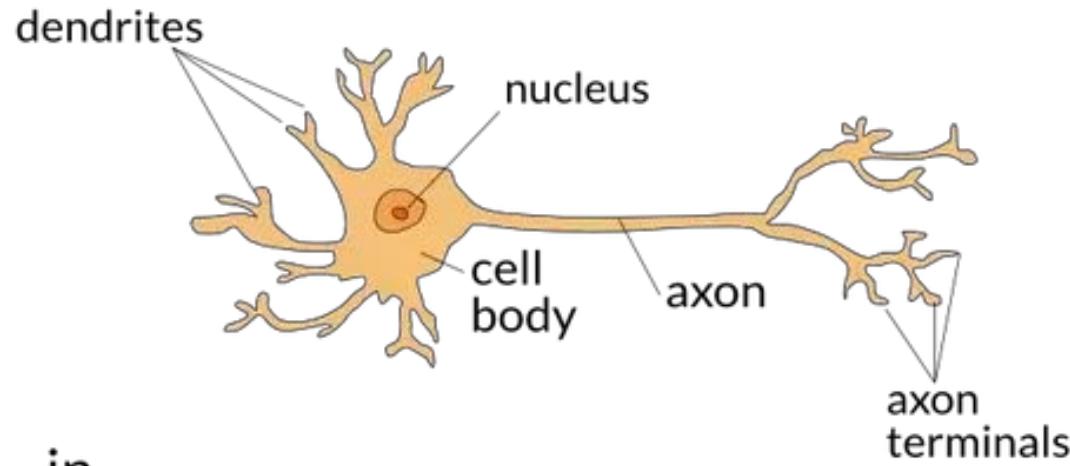
**Antonio Rueda-Toicen**

SPONSORED BY THE

# Learning goals

- Explain how neural networks learn during training
- Identify the components of input units, hidden layers, and output units
- Apply concepts of weights, gradients and loss
- Understand the purpose of validation and test data

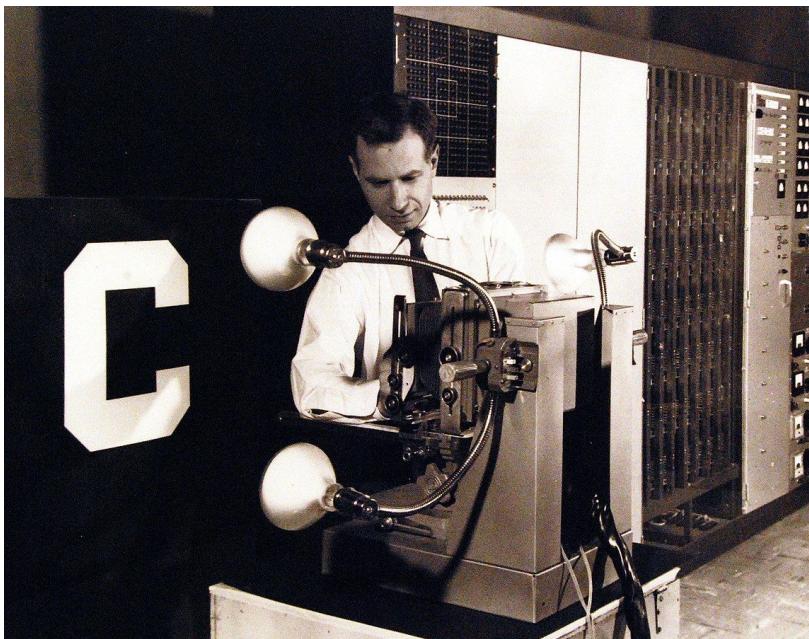
# What is an artificial neural network?



- Artificial “neurons” are also called “processing units”
- We can think of them as functions that map  $X$  inputs to  $\hat{Y}$  outputs

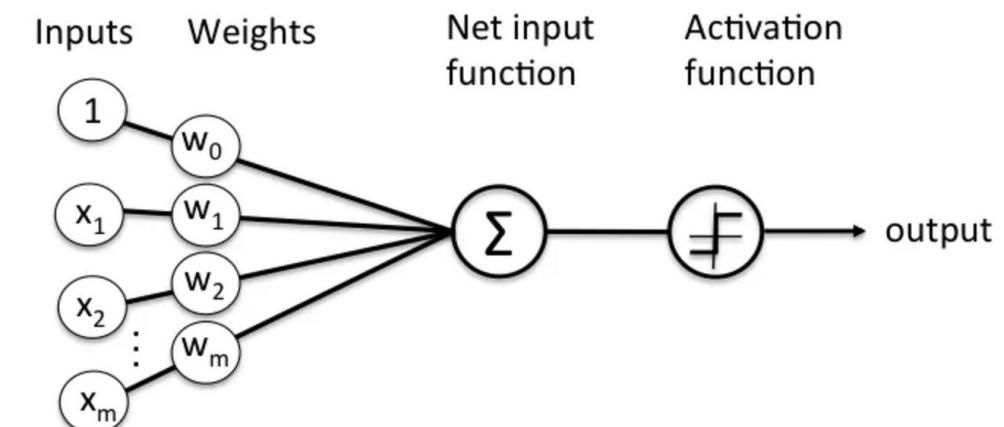
Image from [source](#)

# The perceptron



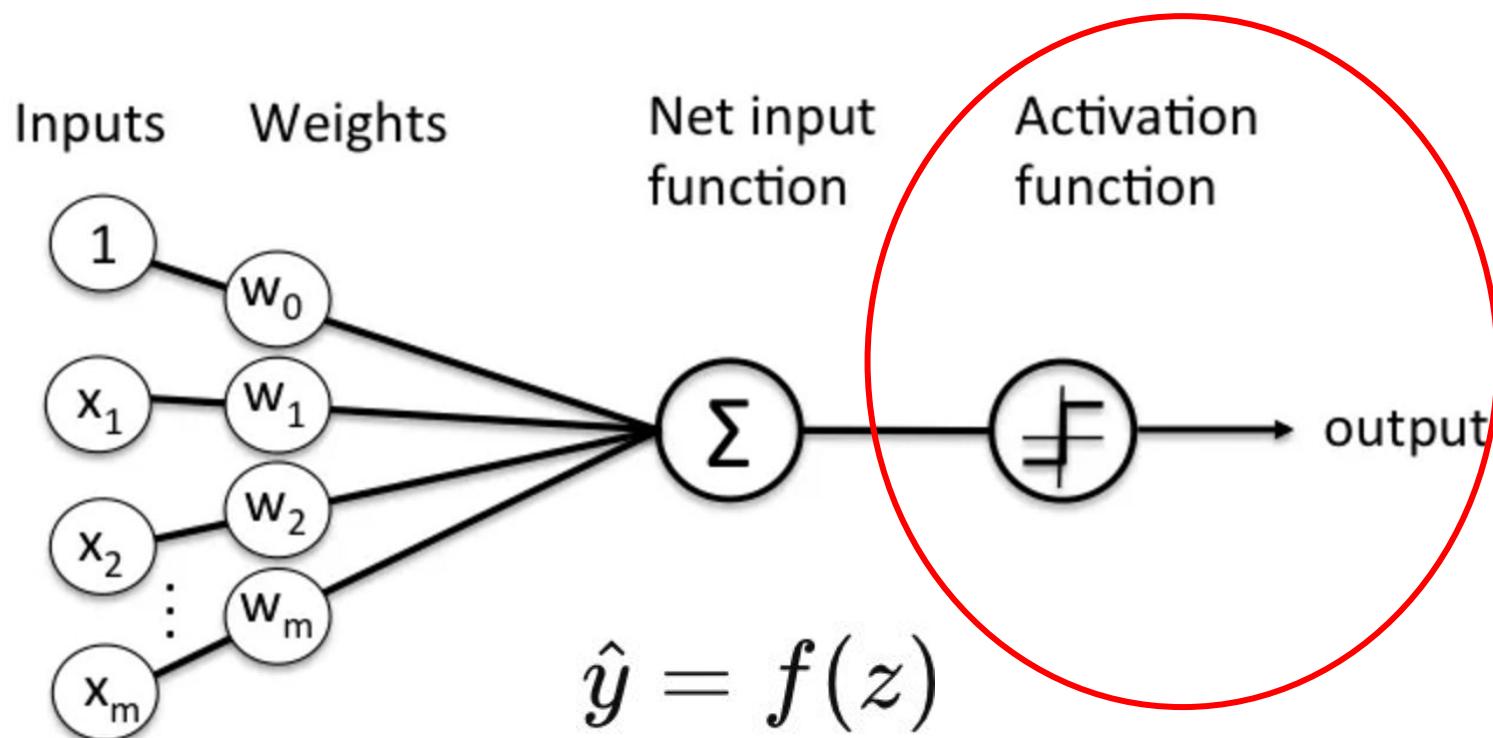
Images: Mark I Perceptron at Cornell University (1961) [source](#)

Image from [source](#)



$$\hat{y} = f \left( \sum_{j=0}^m w_j x_{ij} \right)$$

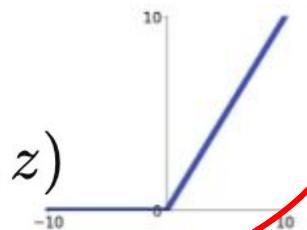
# A modern fully connected layer



$$z = \sum_{i=0}^m w_i x_i$$

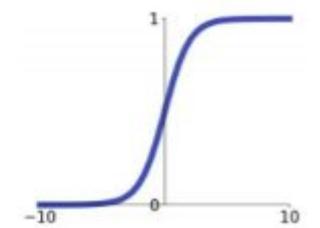
**ReLU**

$$f(z) = \max(0, z)$$



**Sigmoid**

$$f(z) = \frac{1}{1 + e^{-z}}$$



# Making a perceptron learn the price of a vintage car

training



**1964 Chevrolet Corvette Stingray**

Image [source](#)

$y = 83999$  EUR (real price, target value)

$\hat{y} = 84000$  EUR (a possible prediction)

Absolute error =  $|\hat{y} - y| = 1$  EUR

validation



Image [source](#)

testing



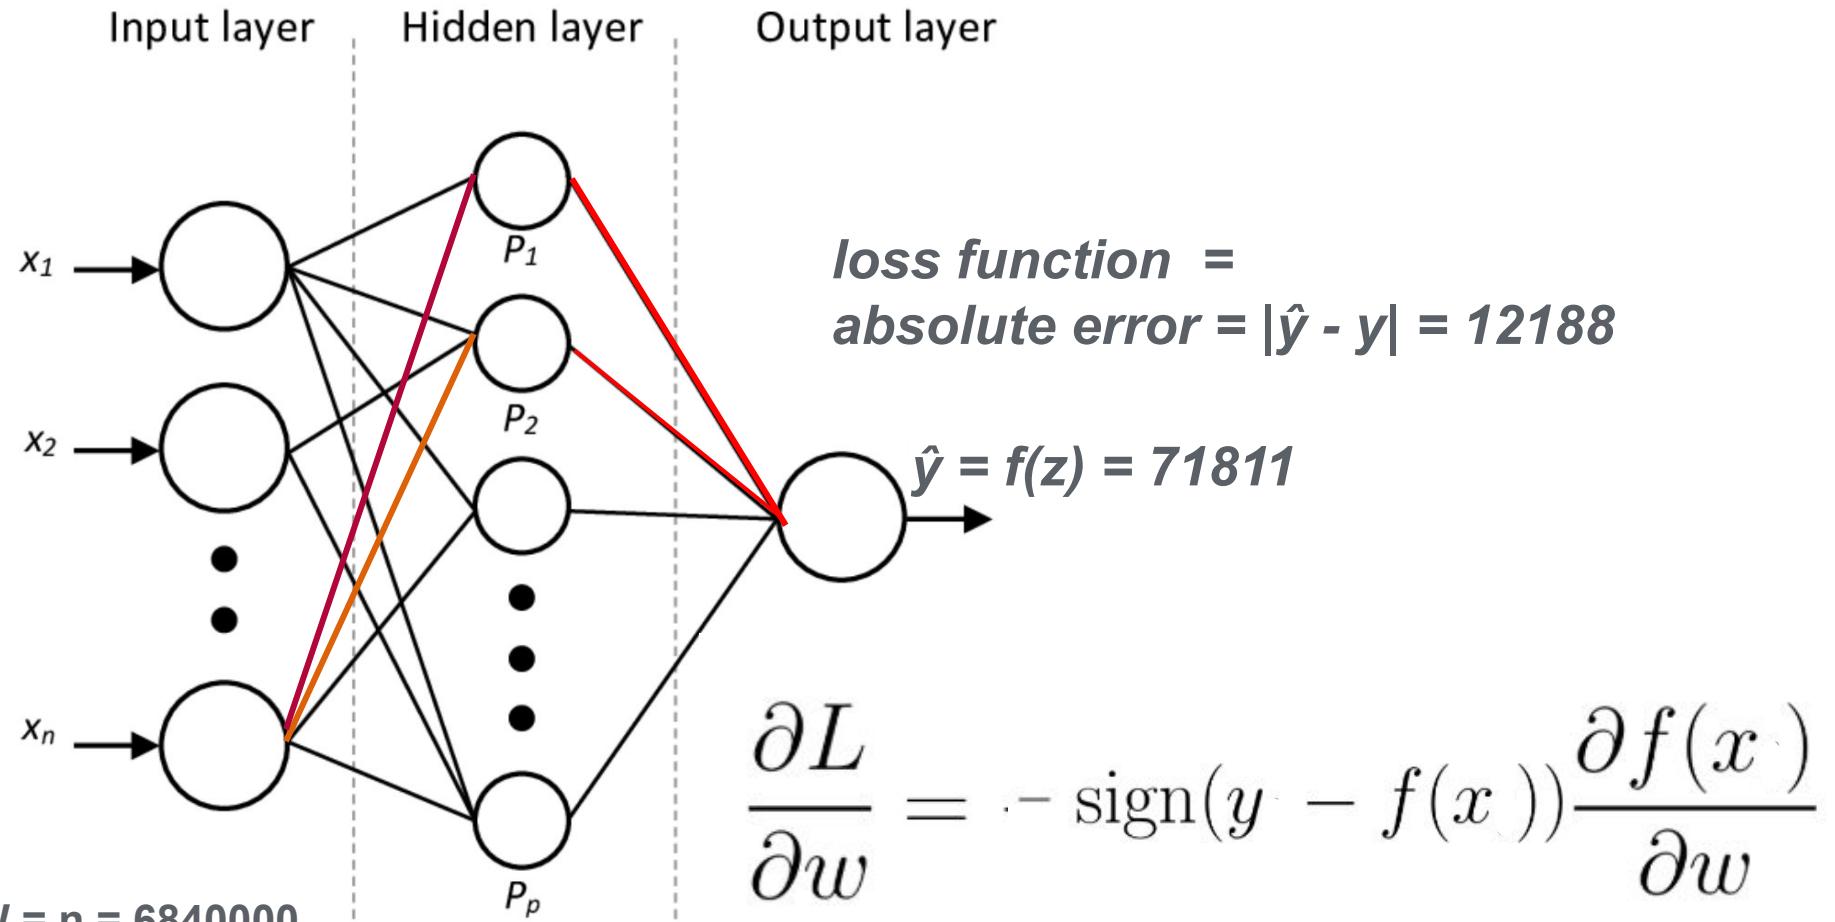
Image [source](#)

# Training a fully connected network (multilayer perceptron)

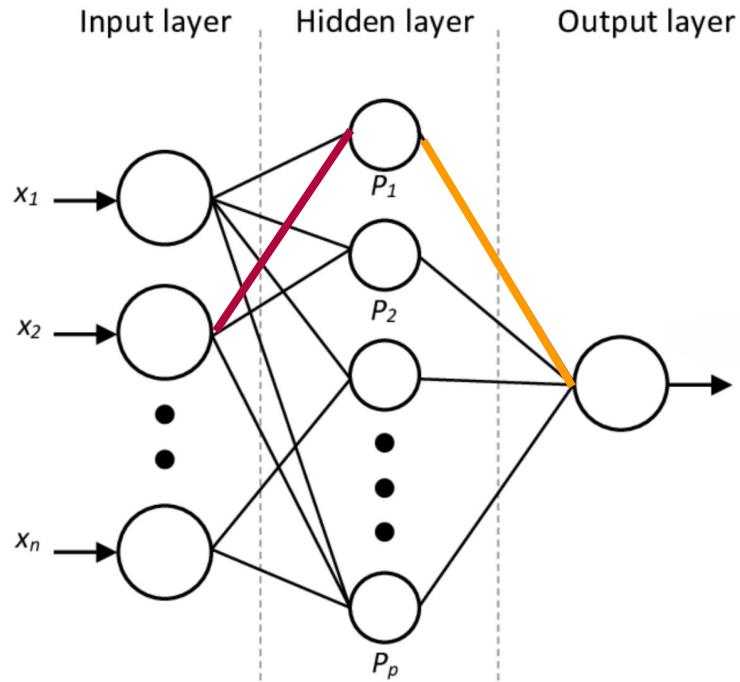
$y = 83999 \text{ EUR}$



$H = 1900$   
 $W = 1200$   
 $C = 3 \text{ (RGB)}$     $C * H * W = n = 6840000$



# Understanding network weights and their updates



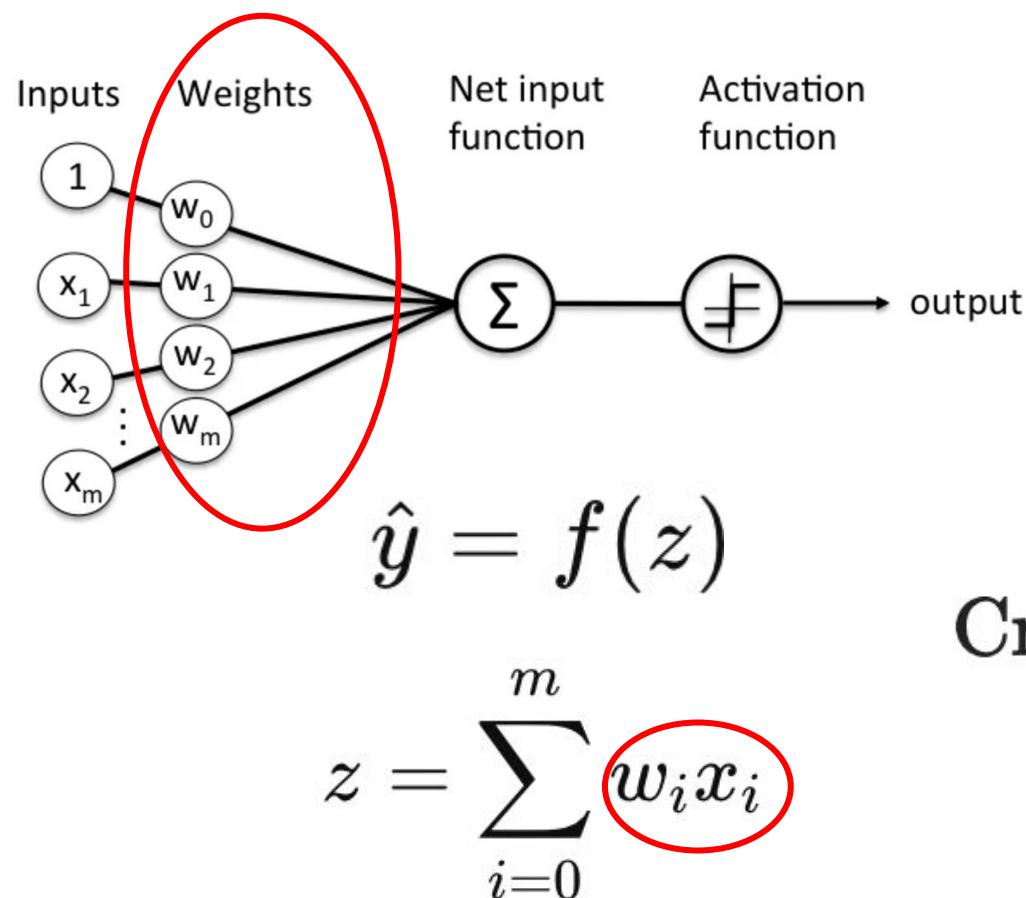
$w_{ij}$  = weight from neuron  $i$  to neuron  $j$

learning\_rate = step size for updates, a constant like 0.01

$L$  = loss function

$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$

# Modern perceptrons use loss functions as feedback



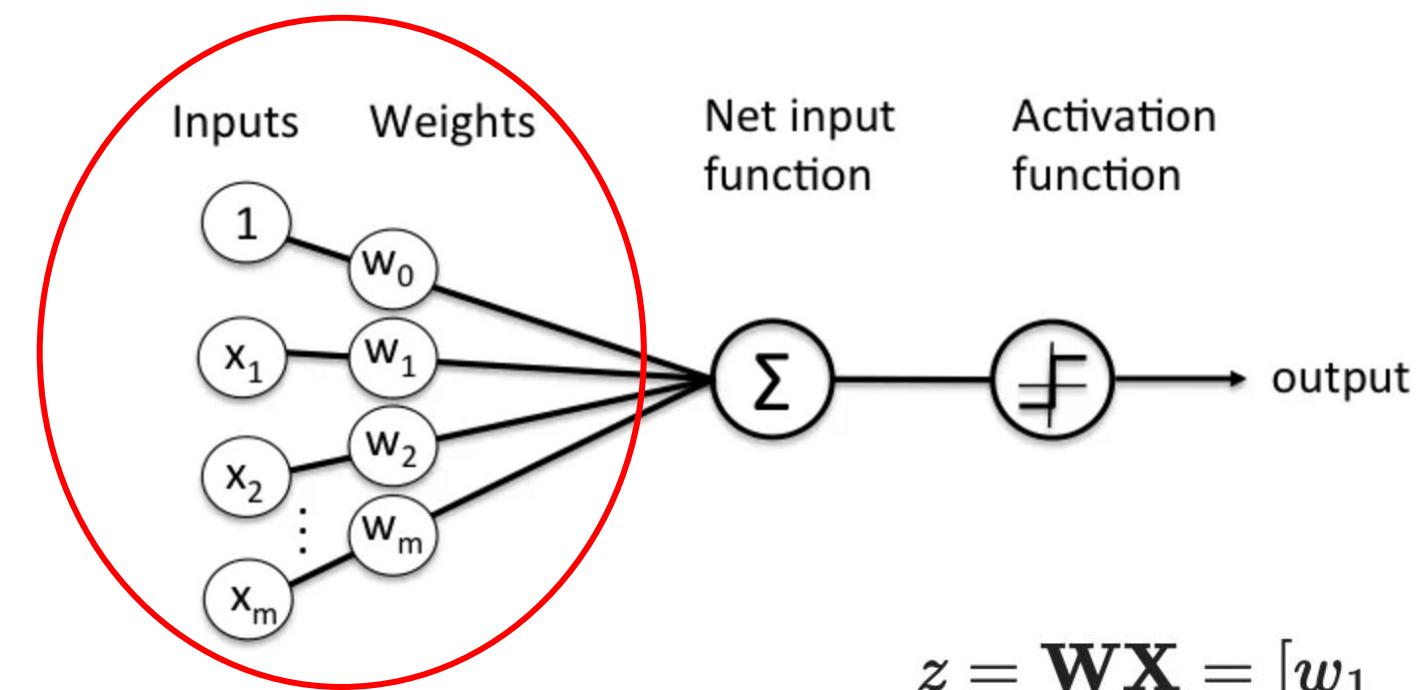
Regression

$$\text{Mean Absolute Error} = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

Classification

$$\text{Cross Entropy} = -\frac{1}{N} \sum_{i=1}^N \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c})$$

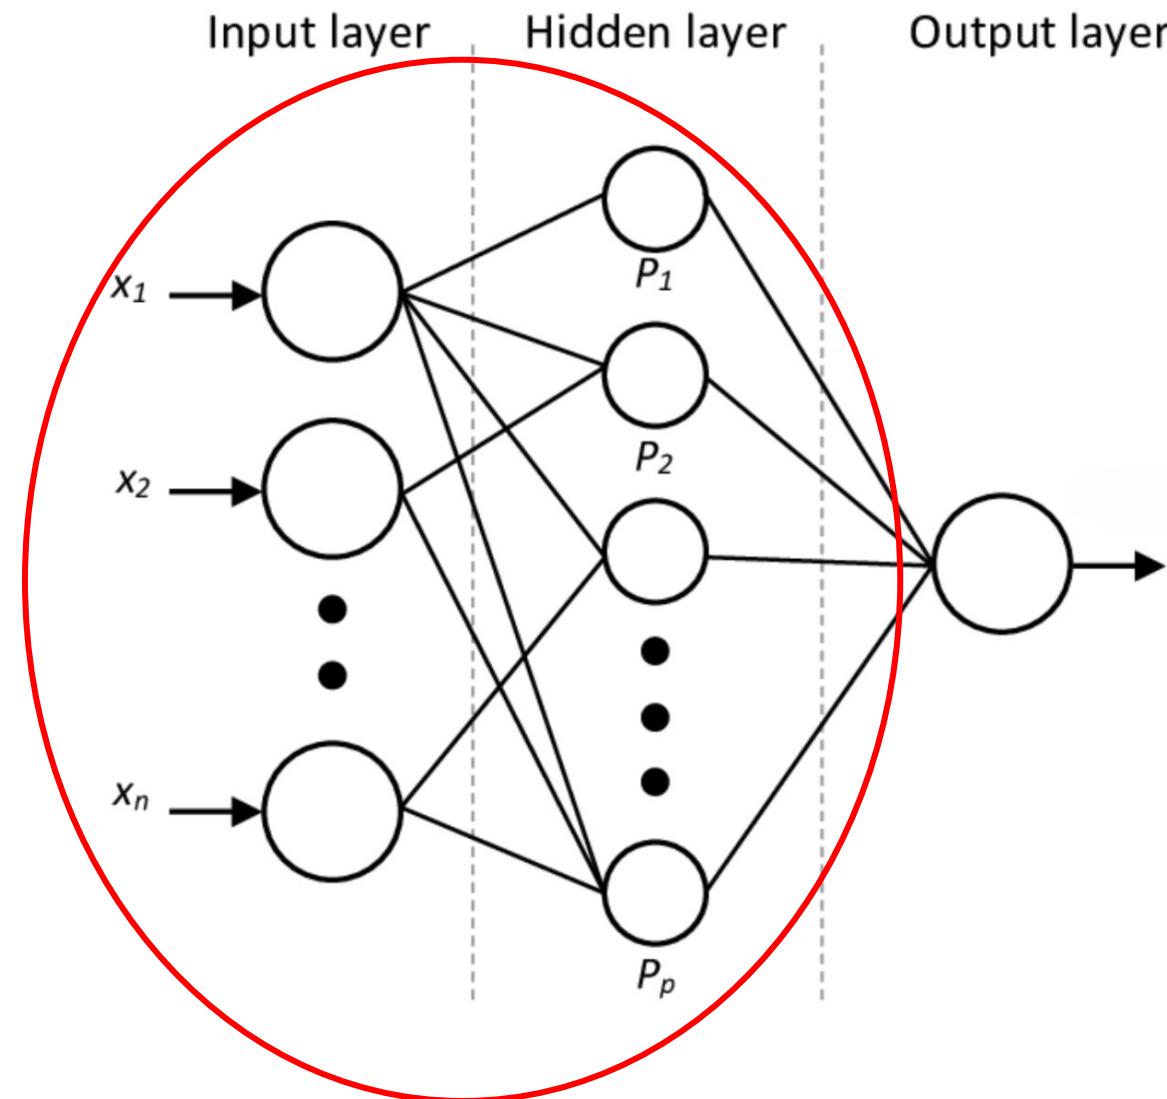
# Matrix multiplication: features times weights



$$z = \mathbf{W}\mathbf{X} = [w_1 \quad w_2 \quad \dots \quad w_m]$$

$$\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} = \sum_{i=1}^m w_i x_i$$

# Dimensions of the weight matrix define the output's



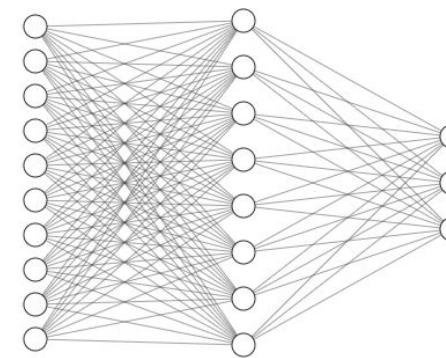
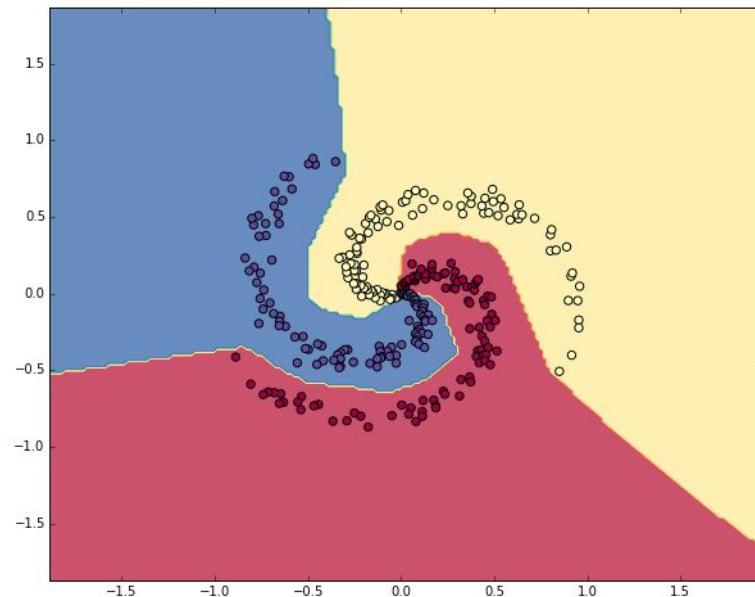
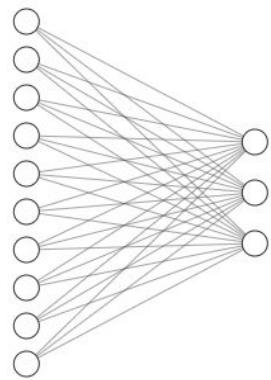
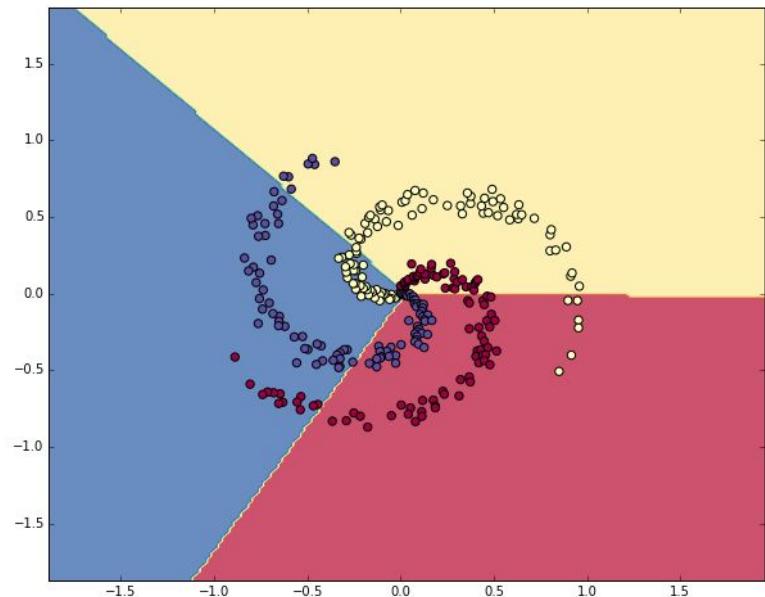
$$P = WX$$

$$\begin{bmatrix} p_1 \\ p_2 \\ \vdots \\ p_p \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{p1} & w_{p2} & \cdots & w_{pn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

For example:

$$p = 4 \quad n = 3$$

# Why do we add hidden layers?



# Validating and testing the model

sample to decide when to stop training: validation



[www.motorcarclassics.com](http://www.motorcarclassics.com)

$$y = 82999$$

Image [source](#)

sample to use as a final benchmark: test



$$y = 82699$$

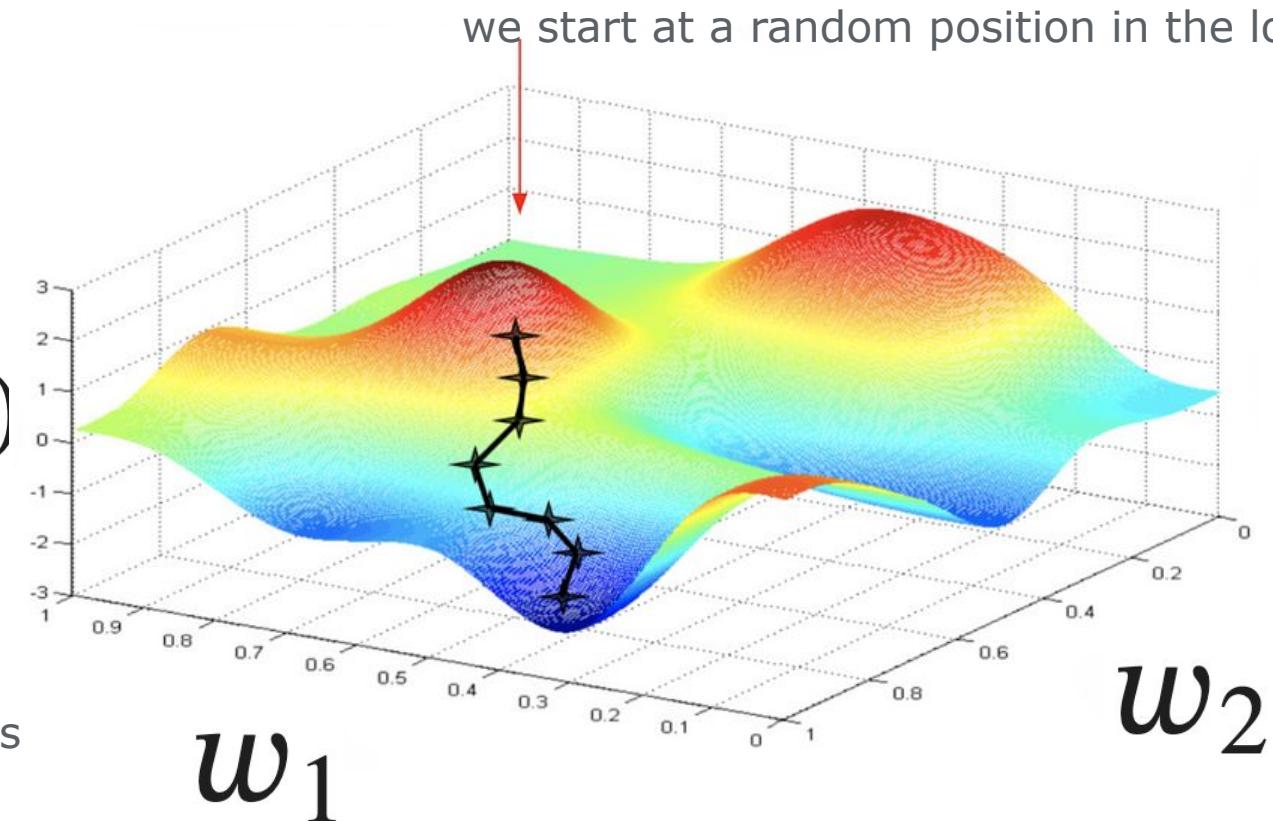
Image [source](#)

# Gradient descent

$$L(w_1, w_2)$$



training image  
role: update weights



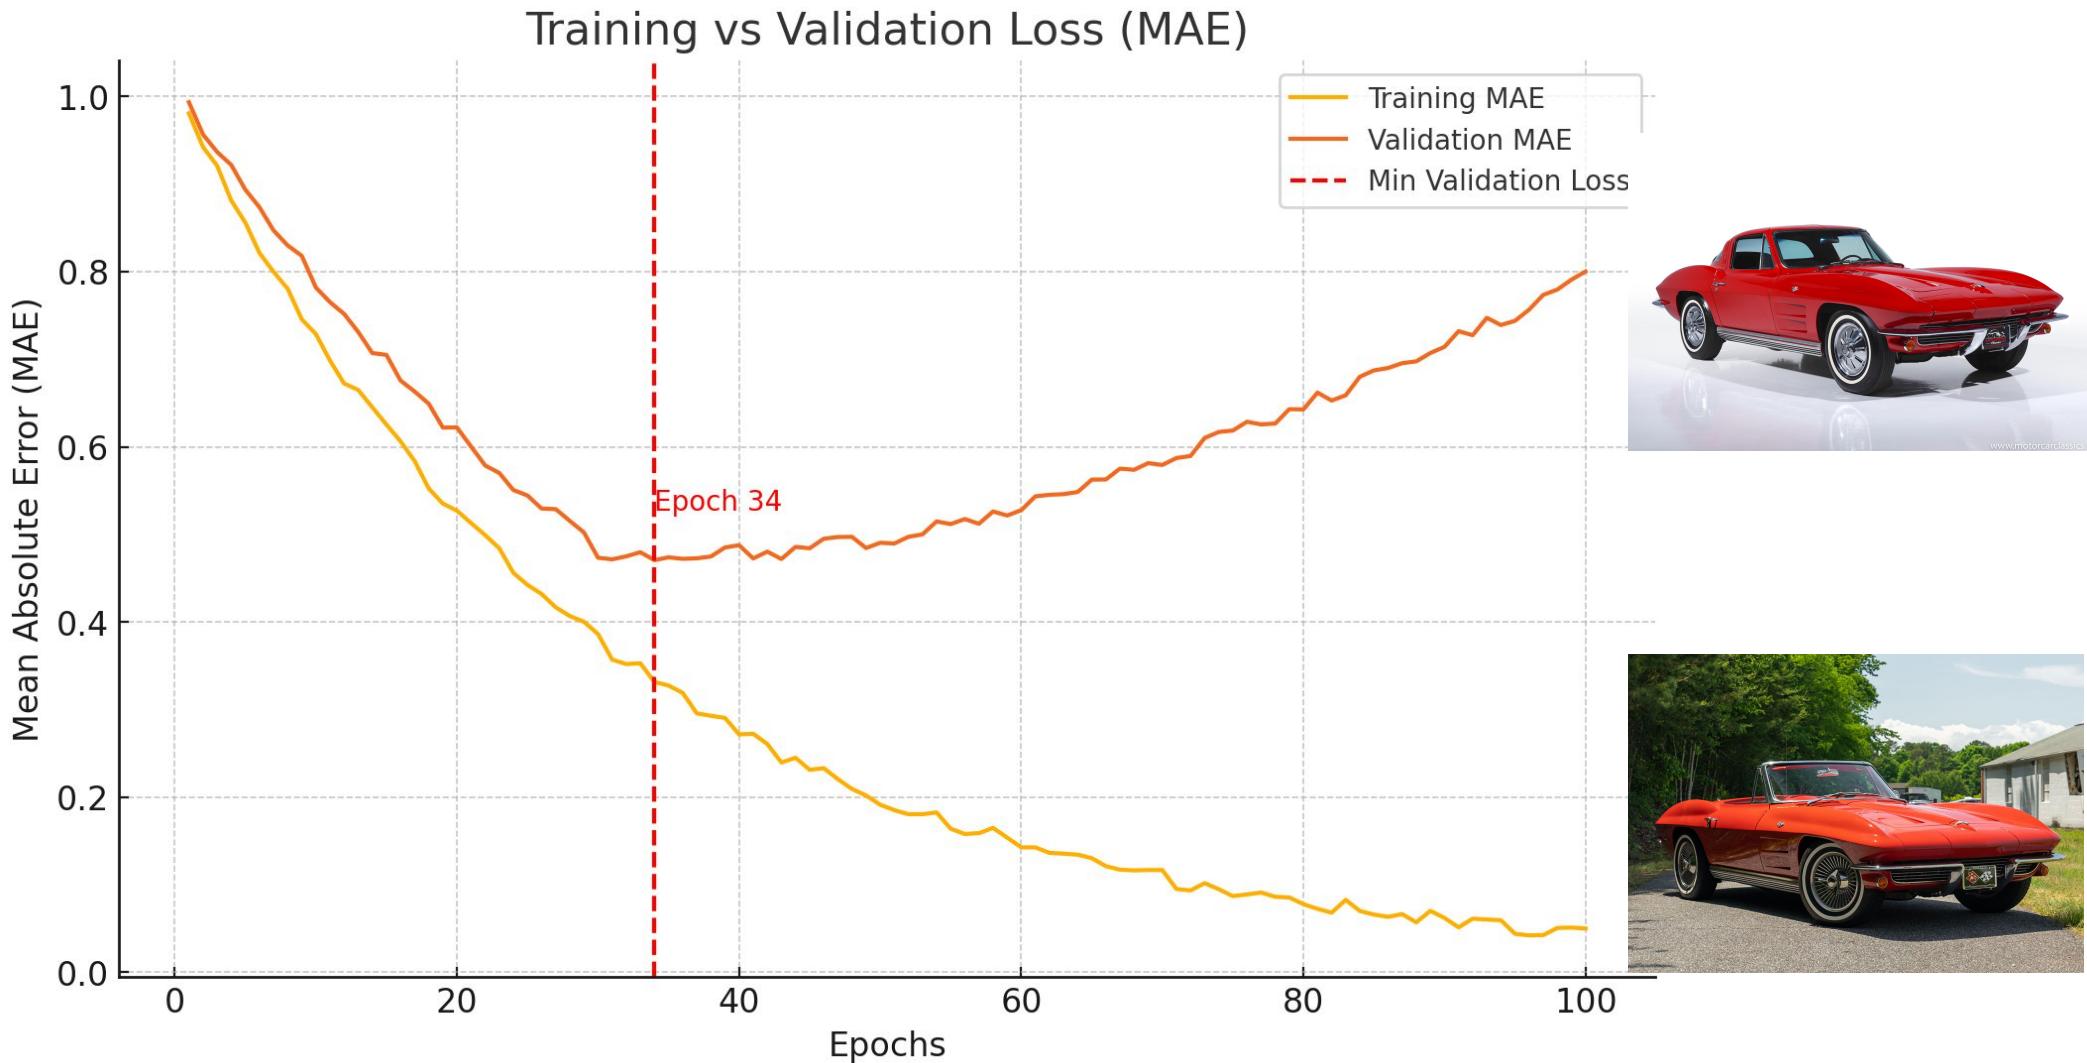
we start at a random position in the loss landscape



validation image  
role: evaluate best weights

$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$

# Training vs validation loss



# The test set gives the benchmark of generalization



$$y = 82699$$

$$\hat{y} = 84015$$

$$\text{test set error} = |y - \hat{y}| = 1316$$

Image from <https://xkcd.com/1838/>

# Summary

**Neural networks are simplified mathematical models of biological neurons**

- They are functions with trainable parameters (aka ‘weights’) that we stack in layers with variable number of processing units (aka ‘neurons’)
- The size of our weight matrices define the dimensions of our hidden and output layers

**We use loss functions and gradients to update the weights of a network**

- Loss functions provide feedback to the model
- The weight update is an iterative process (gradient descent)

**Deciding on training settings (aka “hyperparameters”) is an explorative process**

- We use validation sets to decide how and when to stop training. We use tests to provide final benchmarks

SPONSORED BY THE

# Further reading

## On the origin of neural networks

- <https://nautil.us/the-man-who-tried-to-redeem-the-world-with-logic-2885/>

## Annotated history of modern AI and deep learning by Jürgen Schmidhuber

- <https://people.idsia.ch/~juergen/deep-learning-history.html>

## Frank Rosenblatt's Perceptron

- <https://news.cornell.edu/stories/2019/09/professors-perceptron-paved-way-ai-60-years-to-o-soon>