


# Building a Multilayer Perceptron for Regression in PyTorch



**Antonio Rueda-Toicen**

SPONSORED BY THE



Federal Ministry  
of Education  
and Research


# Learning goals

- Implement a fully connected network in PyTorch with `nn.Linear` and `nn.Sequential` syntax
- Explore the necessary agreements between processing units and input size
- Understand the motivations for image input resizing
- Code input units, hidden layers, and output units
- Examine the representation of weights, gradients, and loss in PyTorch

# A minimal network

input layer

hidden layer



```
import torch.nn as nn
```

```
# Input layer to hidden layer (3 -> 4)
```

```
i_h_layer = nn.Linear(in_features=3, out_features=4, bias = False)
```

```
# Hidden layer to output layer (4 -> 1)
```

```
h_o_layer = nn.Linear(in_features=4, out_features=1, bias = False)
```

```
# Connecting the layers, notice that 'hidden layer' is implicit
```

```
model = nn.Sequential(  
    i_h_layer,  
    nn.ReLU(),  
    h_o_layer,  
)
```

```
# The input has as many entries as we have input units
```

```
input_data = torch.tensor([1.0, 2.0, 3.0], dtype=torch.float32)
```

```
# Feedforward pass
```

```
prediction = model(input_data)
```

```
# The value that we want to match
```

```
target = torch.tensor([11.], dtype=torch.float32)
```

```
# Our measure of error
```

```
loss_function = nn.L1Loss()
```

```
# Evaluate how good our prediction is
```

```
loss_value = loss_function(target, prediction)
```

# Inspecting nn.Linear

```
import torch
import torch.nn as nn

# Define a linear layer
layer = nn.Linear(in_features=4, out_features=1, bias=True)
```

```
# Examine parameters, weights and bias
for name, param in layer.named_parameters():
    print(f"{name}: {param.shape}")
    print(param)
```

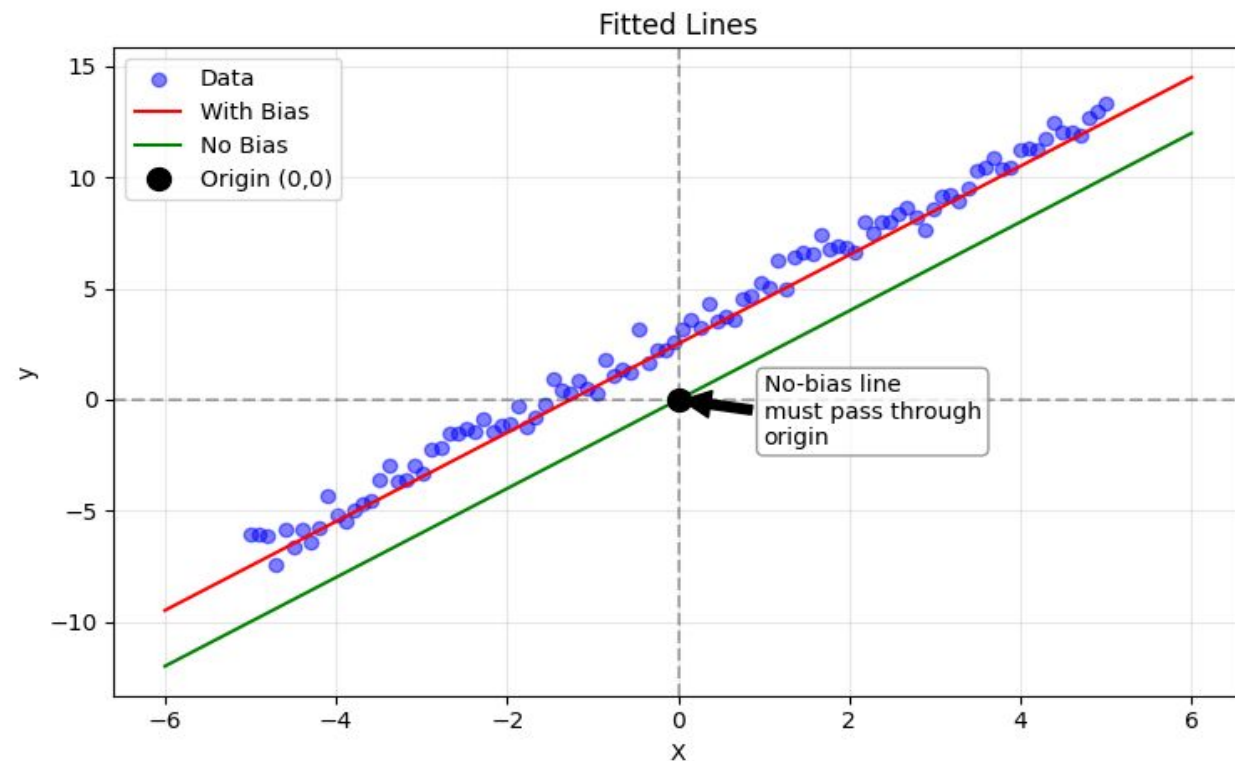

```
# Forward pass
x = torch.randn(5, 4) # Batch of 5 samples, 4 features each
output = layer(x)
```

```
# Batch of 5 samples, single output value
print(output.shape)
```

$$z = \mathbf{WX} = \begin{bmatrix} w_1 & w_2 & \cdots & w_m \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} = \sum_{i=1}^m w_i x_i$$

```
weight: torch.Size([1, 4])
Parameter containing:
tensor([[ -0.0576,  0.3385, -0.3604, -0.1795]], requires_grad=True)
bias: torch.Size([1])
Parameter containing:
tensor([0.3461], requires_grad=True)
```

# Bias vs no bias



$$\text{Mean Squared Error (MSE) Loss} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

[A Colab demo notebook](#)

# Inspecting nn.Sequential


*# Connecting the layers,  
# notice that 'hidden layer' is implicit*

```
model = nn.Sequential(  
    nn.Linear(in_features=3, out_features=4, bias = False),  
    nn.ReLU(),  
    nn.Linear(in_features=4, out_features=1, bias = False)  
)
```

*# Forward pass*

```
x = torch.randn(5, 3) # Batch of 5 samples, 3 features each  
output = model(x)
```

*# Batch of 5 outputs, each one having one regression target value  
print(output.shape) # prints torch.Size([5, 1])*



# Fully connected networks (aka multilayer perceptrons)

$y = 83999 \text{ EUR}$




$H = 1920$

$W = 1282$

$C = 3 \text{ (RGB)}$

$C * H * W = n = 7381440$







# Resizing the input to reduce the number of parameters



**H = 1920**  
**W = 1282**  
**C = 3 (RGB)**



```
transforms.Compose([  
    transforms.ToImage(),  
    transforms.Resize((224, 224)),  
    transforms.ToDtype(torch.float32, scale=True)  
])
```




**H = 224**  
**W = 224**  
**C = 3 (RGB)**



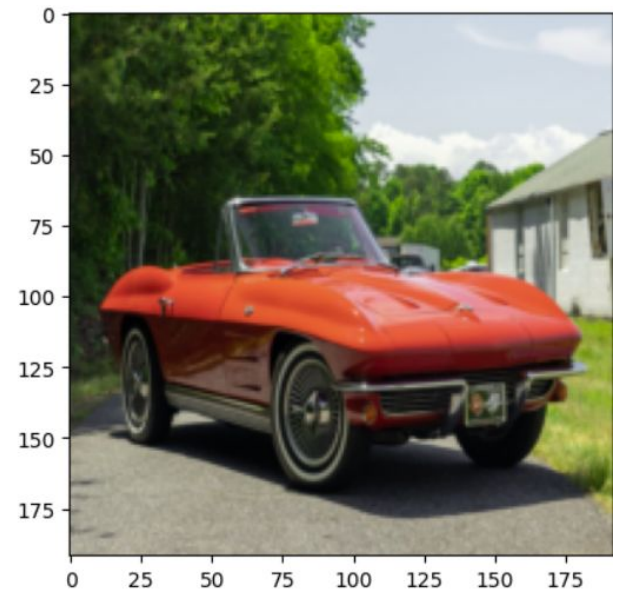
## CenterCrop to capture more of the object of interest



H = 1920  
W = 1282  
C = 3 (RGB)




```
transforms.Compose([  
    transforms.ToImage(),  
    transforms.Resize((256, 256)),  
    transforms.CenterCrop(192),  
    transforms.ToDtype(torch.float32, scale=True)  
])
```



H = 192  
W = 192  
C = 3 (RGB)

## Stochastic gradient descent - the weight update rule



$w_{ij}$  = weight from neuron  $i$  to neuron  $j$

learning\_rate = step size for updates, a constant like 0.01

$L$  = loss function

$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$

# Implementing the network

*# Number of units on the hidden layer, this number is arbitrary*

`p = 100`

*# We use nn.Sequential to concatenate PyTorch modules in order*

```
multilayer_perceptron=nn.Sequential(
```

```
    nn.Flatten(), # Flattens the images within a batch
```

```
    nn.Linear(in_features= C * H * W, # This matches the image shape
```

```
              out_features= p,
```

```
              bias=True),
```


```
    nn.ReLU(),
```

```
    nn.Linear(in_features=p,
```

```
              out_features=1, # We are predicting a single value
```

```
              bias=True),
```

```
)
```



**H = 192**  
**W = 192**  
**C = 3 (RGB)**

**y = 83999 EUR**

## Implementing the weight update rule

```
def update_rule(weight, grad, learning_rate):  
    return nn.Parameter((weight - learning_rate * grad))  
  
learning_rate = 0.1  
new_weights_0 = update_rule(multilayer_perceptron[0].weight,  
                             multilayer_perceptron[0].weight.grad,  
                             learning_rate)  
  
new_weights_2 = update_rule(multilayer_perceptron[2].weight,  
                             multilayer_perceptron[2].weight.grad,  
                             learning_rate)  
  
# We reassign the weights to our model  
multilayer_perceptron[0].weight = new_weights_0  
multilayer_perceptron[2].weight = new_weights_2
```



Image from ideogram.ai

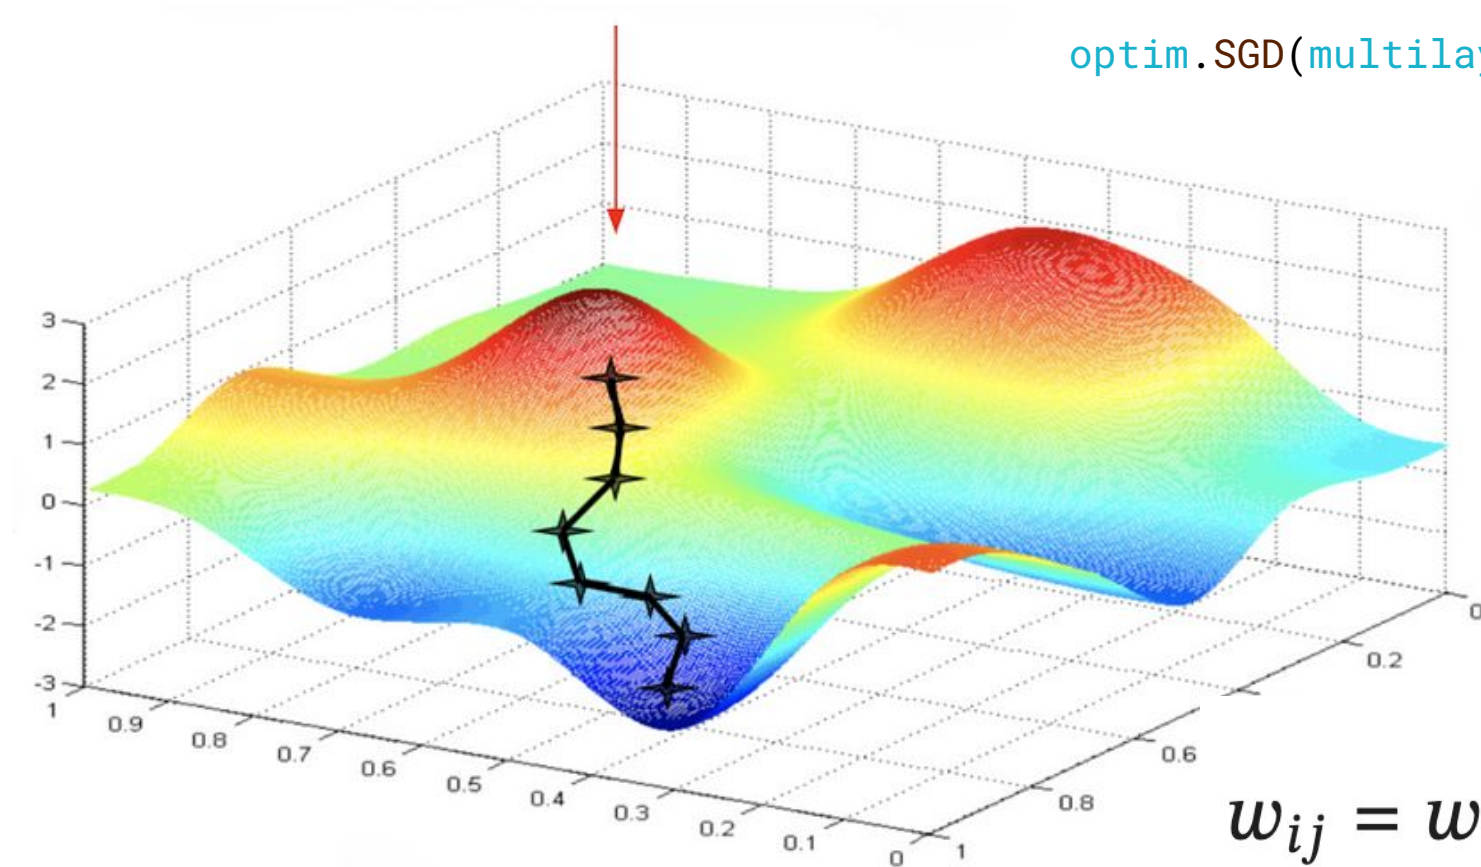
$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$



## Training with an optimizer

```
import torch.optim as optim
```

```
optim.SGD(multilayer_perceptron.parameters(), lr=1e-2)
```



$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$

# Training with an optimizer

```
import torch
import torch.nn as nn
import torch.optim as optim
```

```
# Define an Stochastic Gradient Descent optimizer
optimizer = optim.SGD(multilayer_perceptron.parameters(), lr=1e-2)
```

```
# Define a loss function
criterion = nn.L1Loss()
```

```
# Forward pass
pred = multilayer_perceptron(x)
```

```
# Compute loss
loss = criterion(pred, y)
```


```
# Backprop
optimizer.zero_grad()
loss.backward()
```

```
# Update weights
optimizer.step()
```

$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$



# The importance of the learning rate



# Summary

## PyTorch building blocks

- Multilayer perceptrons in PyTorch combine `Linear` layers, activation functions, and loss functions. The framework handles gradient computation through autograd.
- We can use the `Sequential` syntax to wrap up components, this is an alternative to subclassing `nn.Module`

## Input processing

- Image input requires preprocessing through resizing and normalization to create manageable feature tensors. The input size determines the network architecture.

## Model training

- Training combines forward passes, loss computation, backpropagation, and weight updates. The learning rate controls optimization stability and convergence speed.

# Further reading

## What is torch.nn really?

- [https://pytorch.org/tutorials/beginner/nn\\_tutorial.html](https://pytorch.org/tutorials/beginner/nn_tutorial.html)

## Tensorflow playground

- <https://playground.tensorflow.org/>

SPONSORED BY THE