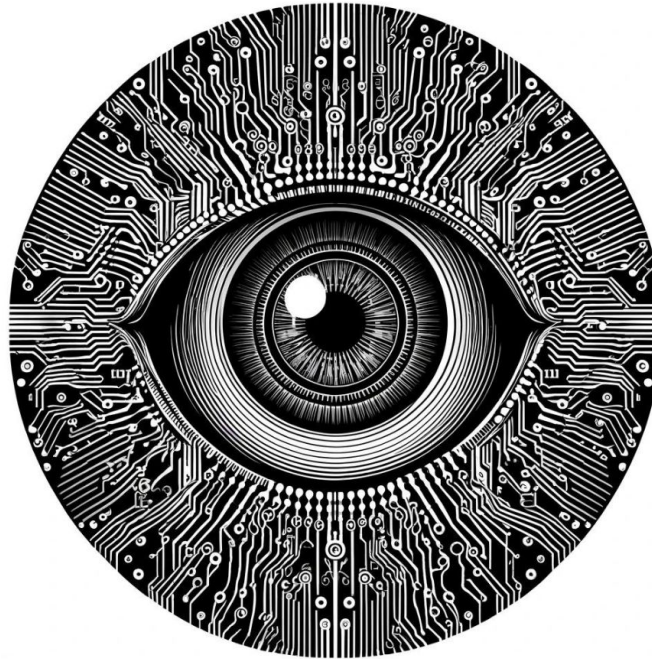


# Fundamentals of Convolutions



**Antonio Rueda-Toicen**

SPONSORED BY THE



Federal Ministry  
of Education  
and Research

# Learning goals

- Understand the role of convolutions in producing image features
- Evaluate effect of stride, padding, and filter size on output images

# Convolution filters create new images






Image from [source](#)



$$\text{Sharpening Kernel} = \begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

$$\text{Image Patch (Input)} = \begin{bmatrix} 10 & 20 & 30 \\ 40 & 50 & 60 \\ 70 & 80 & 90 \end{bmatrix}$$

$$\text{Output} = (0 \cdot 10) + (-1 \cdot 20) + (0 \cdot 30) + (-1 \cdot 40) + (5 \cdot 50) + (-1 \cdot 60) + (0 \cdot 70) + (-1 \cdot 80) + (0 \cdot 90)$$

$$\text{Output} = 50$$

# Convolutions were ‘traditionally’ handcrafted

Input Image




Blurred Image with Convolution Filter



$$\begin{bmatrix} 0.111 & 0.222 & 0.111 \\ 0.222 & 0.444 & 0.222 \\ 0.111 & 0.222 & 0.111 \end{bmatrix}$$

Gaussian blur kernel

# Handcrafted convolution kernels



Vertical edge detection:

$$K_v = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

Horizontal edge detection:

$$K_h = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

# The receptive field size controls how many neighboring pixels are considered

Input Image



Blurred Image (3x3)



Blurred Image (5x5)




3x3 vs 5x5 uniform blur kernels

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}$$

$$\begin{bmatrix} \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \\ \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} & \frac{1}{25} \end{bmatrix}$$



# Padding and stride effects



Convolution of 3x3 and stride = 1 without padding

Effect: the output loses one pixel on each dimension

# Border padding

zeros



border




reflection



Image from [the fastai documentation](#)



# “Automatic” feature learning



$$\begin{pmatrix} -0.64 & -0.74 & 0.91 \\ -0.96 & 0.48 & -0.46 \\ 0.56 & 0.67 & -0.08 \end{pmatrix}$$


$$\begin{pmatrix} -0.89 & 0.40 & 0.12 \\ -0.45 & 0.38 & 0.67 \\ 0.91 & -0.71 & -0.35 \end{pmatrix}$$

$$\begin{pmatrix} 0.65 & 0.39 & -0.75 \\ 0.36 & -0.57 & 0.34 \\ 0.61 & 0.69 & 0.43 \end{pmatrix}$$

$$\begin{pmatrix} -0.64 & -0.74 & 0.91 \\ -0.96 & 0.48 & -0.46 \\ 0.56 & 0.67 & -0.08 \end{pmatrix}$$



After 3x3 Conv (Channel 1/3)  
[1, 3, 224, 224]



After 1x1 Conv (Channel 1/2)  
[1, 2, 224, 224]



These weights are adjusted to minimize the loss function

$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic})$$

After 3x3 Conv (Channel 2/3)  
[1, 3, 224, 224]



After 1x1 Conv (Channel 2/2)  
[1, 2, 224, 224]



Convolutions of 3x3 and stride = 1 padded with zeros


**Effect:** the output preserves the original image size while producing a “different-looking” image (based on the values of the weights)

# ReLU adds non-linearity to activations

Input with Padding




After Conv2d



After ReLU



# Downsampling images with stride $> 1$ convolutions




Convolution of 3x3 and stride = 2 padded with zeros

Effect: the output is downsampled to about half its size

Input Image  
Shape: 224x224



Conv2d Output (3x3 kernel, stride 2)  
Shape: 112x112






# Convolutions are concatenated

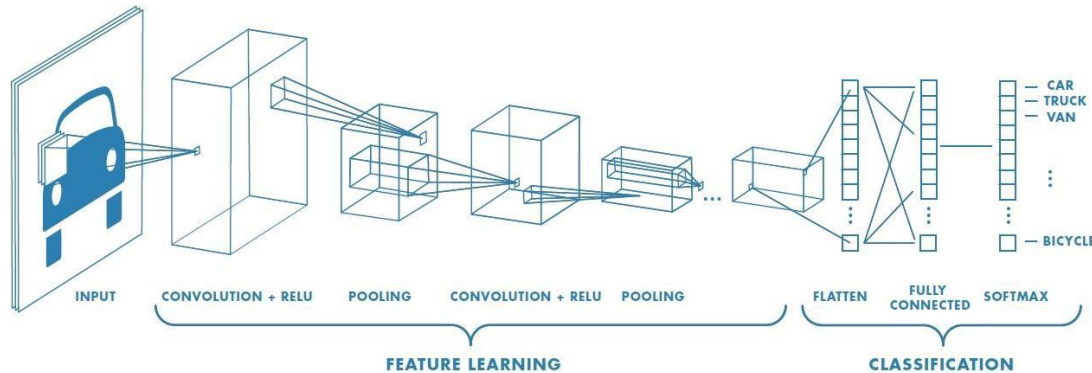

Original Image  
Shape: torch.Size([224, 224])



After First Conv + ReLU (stride = 2)  
Shape: torch.Size([112, 112])

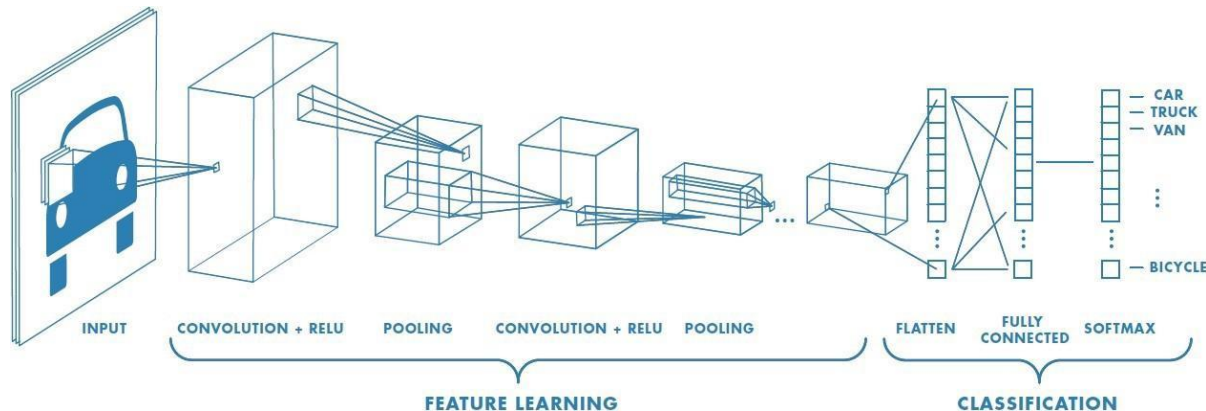



After Second Conv + ReLU (stride = 2)  
Shape: torch.Size([56, 56])

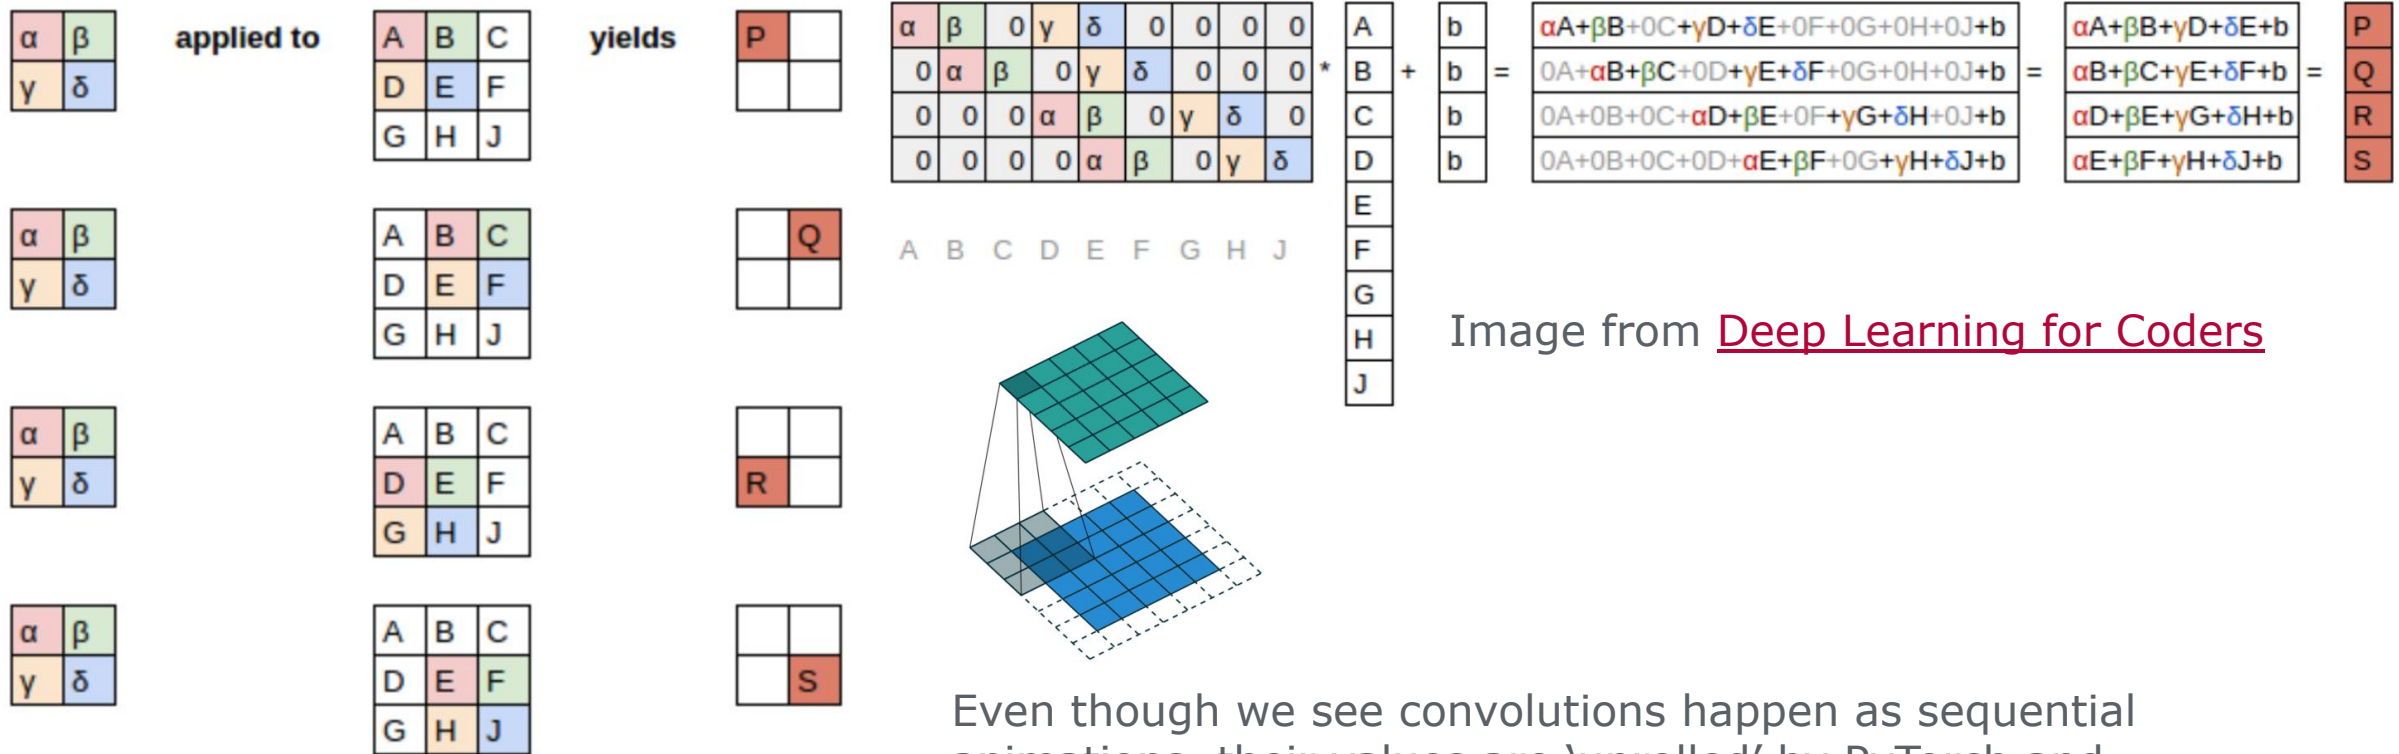


$$\mathcal{L} = - \sum_{i=1}^N \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic})$$

# We use convolutions as “feature extractors” for classifiers



# ‘Unrolling’ convolutions



Even though we see convolutions happen as sequential animations, their values are ‘unrolled’ by PyTorch and applied as parallel matrix multiplications **simultaneously**



# Summary

## Convolutions and feature learning

- Convolutional layers create new images (“new features”) through learnable weights
- Convolutions were traditionally hand-crafted, now we use the neural network to set their weights
- Classifiers flatten these learned features before feeding them to a feedforward network
- Convolutions are “unrolled” by PyTorch in order to do parallel processing

## Effects of receptive field size, padding, and stride

- The size of the receptive field influences how many pixels we consider
- Convolutions with padding preserve spatial dimensions
- Strided convolutions allow us to do learnable downsampling

SPONSORED BY THE

# Further reading and references

## A guide to convolution arithmetic for deep learning

- <https://arxiv.org/abs/1603.07285>

## Network in network (1x1 convolutions)

- <https://arxiv.org/abs/1312.4400>

## Hypercolumns for object segmentation and fine-grained localization

- [https://openaccess.thecvf.com/content\\_cvpr\\_2015/papers/Hariharan\\_Hypercolumns\\_for\\_Object\\_2015\\_CVPR\\_paper.pdf](https://openaccess.thecvf.com/content_cvpr_2015/papers/Hariharan_Hypercolumns_for_Object_2015_CVPR_paper.pdf)