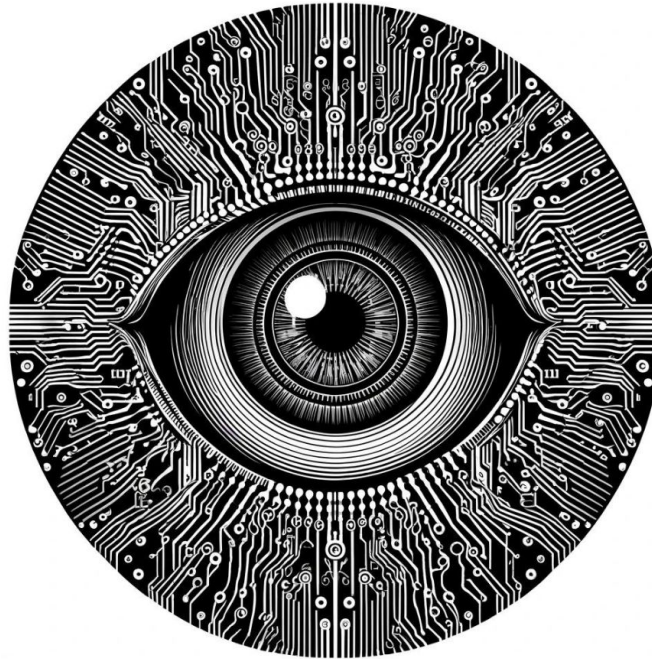


Pooling in Neural Networks



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Explore image downsampling and reduction of network parameters with mean, max, and global pooling

Architecture of a convolutional network with pooling

```
nn.Sequential(  
    nn.Conv2d(3, 16, kernel_size=5, stride=1, padding=2),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=2, stride=2),  
  
    nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),  
    nn.ReLU(),  
    nn.MaxPool2d(kernel_size=2, stride=2),  
  
    nn.Flatten(),  
    nn.Linear(32 * 8 * 8, 120),  
    nn.ReLU(),  
    nn.Linear(120, 84),  
    nn.ReLU(),  
    nn.Linear(84, 10), # 10 classes, we are working with CIFAR 10  
)
```

airplane

automobile

bird

cat

deer

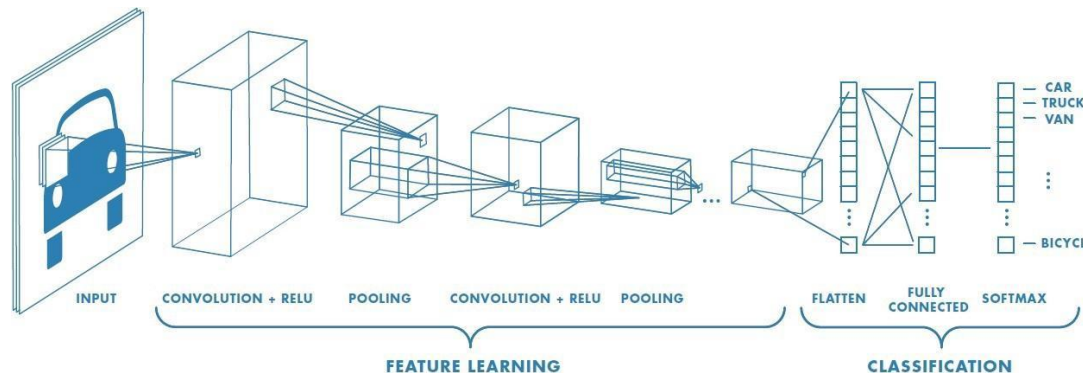

dog

frog

horse


ship

truck



The CIFAR-10 dataset


Max Pooling




Input Image
Shape: 224x224



Max Pooled Image (4x4 Kernel)
Shape: 56x56



Mean Pooling



Input Image
Shape: 224x224




Mean Pooled Image (2x2 Kernel)
Shape: 112x112





Global pooling aka adaptive pooling

```
import torch.nn as nn
```

```
class GlobalPooling(nn.Module):  
    def __init__(self):  
        super().__init__()  
        # Computes the average value of each activation map  
        self.avg_pool = nn.AdaptiveAvgPool2d((1, 1))  
        # Selects the max value of each activation map  
        self.max_pool = nn.AdaptiveMaxPool2d((1, 1))  
  
    def forward(self, x):  
        avg_pooled = self.avg_pool(x).flatten(1)  
        max_pooled = self.max_pool(x).flatten(1)  
        # Concatenates both poolings to produce a feature vector  
        return torch.cat([avg_pooled, max_pooled], dim=1)
```



Feature vector



Using global pooling to reduce parameters


```
import torch
import torch.nn as nn

model = nn.Sequential(
    nn.Conv2d(3, 16, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),


    nn.Conv2d(16, 32, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),

    GlobalPooling(), # This replaces the Flatten layer

    nn.Linear(64, 120),
    nn.ReLU(),
    nn.Linear(120, 84),
    nn.ReLU(),
    nn.Linear(84, 10)
)
```



Feature vector



Summary

Pooling is a way to compress information

- Pooling allows us to do lossy compression while retaining important visual features
- Convolutions with stride > 1 achieve a similar effect at the cost of a higher number of parameters

Global pooling is an alternative to flattening activation maps

- We can create feature vectors (aka embeddings) using the global mean and/or max pooling operations

Further reading and references

A guide to convolution arithmetic for deep learning

- <https://arxiv.org/abs/1603.07285>

Network in network (1x1 convolutions)

- <https://arxiv.org/abs/1312.4400>

Hypercolumns for object segmentation and fine-grained localization

- https://openaccess.thecvf.com/content_cvpr_2015/papers/Hariharan_Hypercolumns_for_Object_2015_CVPR_paper.pdf