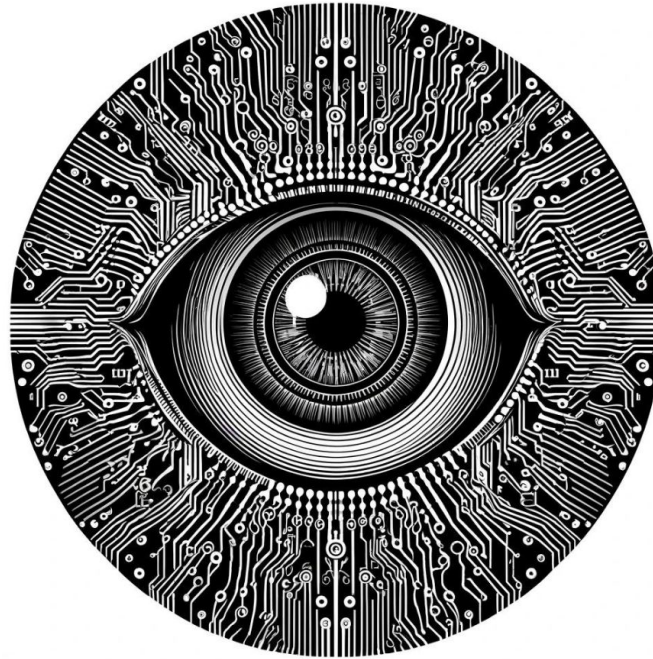


Transfer Learning and Fine-Tuning Pretrained Models



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Understand the use of transfer learning in image classification
- Modify an Imagenet-pretrained network to perform a new classification task
- Perform fine-tuning of a Resnet50 after modifying its architecture

The ImageNet-1K dataset



Explore Imagenet classes on this [notebook](#)

Index 68: sidewinder

Index 85: quail

Index 122: American lobster

Index 216: clumber

Index 273: dingo

Index 302: ground beetle

Index 403: aircraft carrier

Index 420: banjo

Index 775: sarong

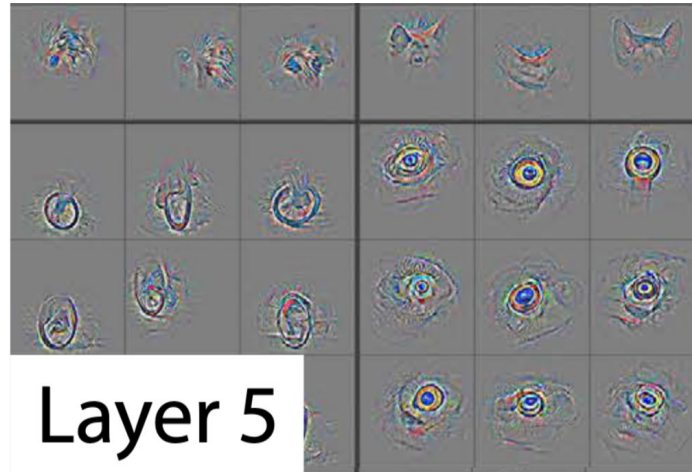
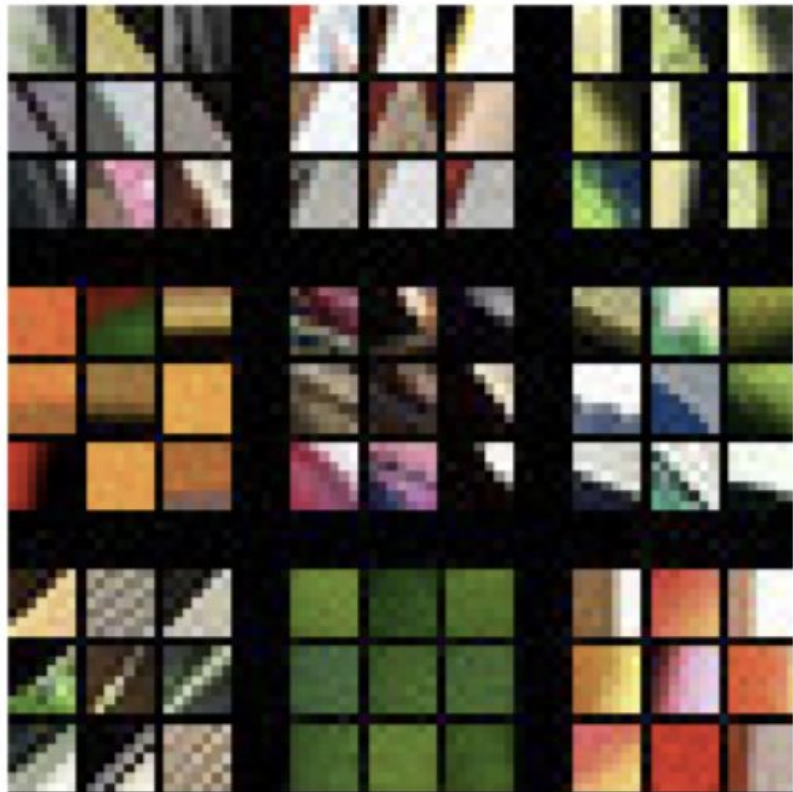
Index 946: cardoon

Image source: <https://cs.stanford.edu/people/karpathy/cnnembed/>

What an Imagenet-pretrained model already knows



Layer 1



Layer 5



Images from [visualizing and understanding convolutional networks](#)

We can transfer these learned features to a new classification task

Fresh vs rotten fruit classification



ImageNet index for 'Granny Smith': 948
ImageNet index for 'banana': 954
ImageNet index for 'orange': 950

Imagenet has no training data for the rotten variants of fruit

Inference on out-of-training-distribution data

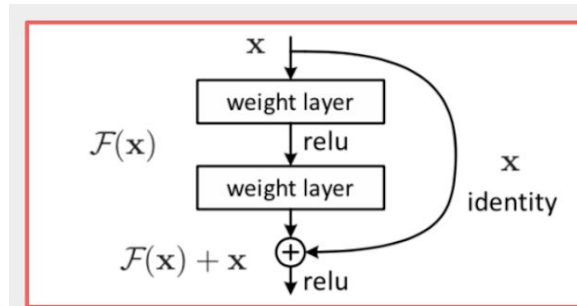
pomegranate, probability = 0.50
buckeye, horse chestnut, conker, probability = 0.21
fig, probability = 0.12
bagel, beigel, probability = 0.03
French loaf, probability = 0.02

face powder, probability = 0.12
golf ball, probability = 0.08
croquet ball, probability = 0.08
baseball, probability = 0.07
ping-pong ball, probability = 0.07



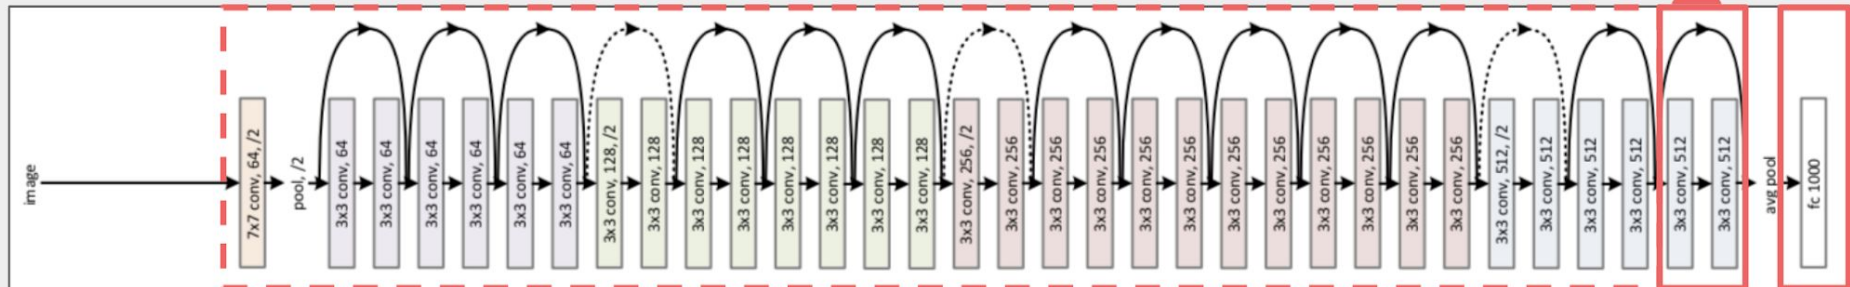
Changing an ImageNet-pretrained Resnet50

An Imagenet-pretrained ResNet has a **1000 output units**
We need to remap this to **6 output units**

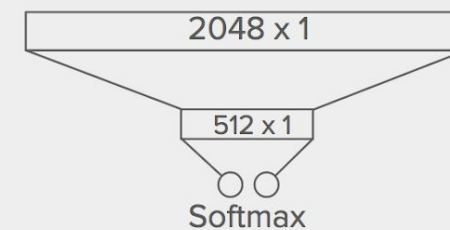


Residual Learning Block

ResNet50 Diagram



Re-architect fully-connected layers



Using the original image transformations

```
import torch
import torchvision.models as models
from torchvision.models import ResNet50_Weights

# Create an instance of ResNet50_Weights pretrained on
# ImageNet-1k
weights = ResNet50_Weights.IMAGENET1K_V2

# Load pretrained ResNet50 with specified weights
model = models.resnet50(weights=weights)

# Access the original data transformations used for
training
original_transformation = weights.transforms()
original_transformation
```

```
ImageClassification(
  crop_size=[224]
  resize_size=[232]
  mean=[0.485, 0.456, 0.406]
  std=[0.229, 0.224, 0.225]
  interpolation=InterpolationMode.BILINEAR
)
```


Step 1: inspect the final layers of the model

```
import torch
import torchvision.models as models
from torchvision.models import ResNet50_Weights

# Create an instance of ResNet50_Weights pretrained on ImageNet-1k
weights = ResNet50_Weights.IMAGENET1K_V2

# Load pretrained ResNet50 with specified weights
model = models.resnet50(weights=weights)

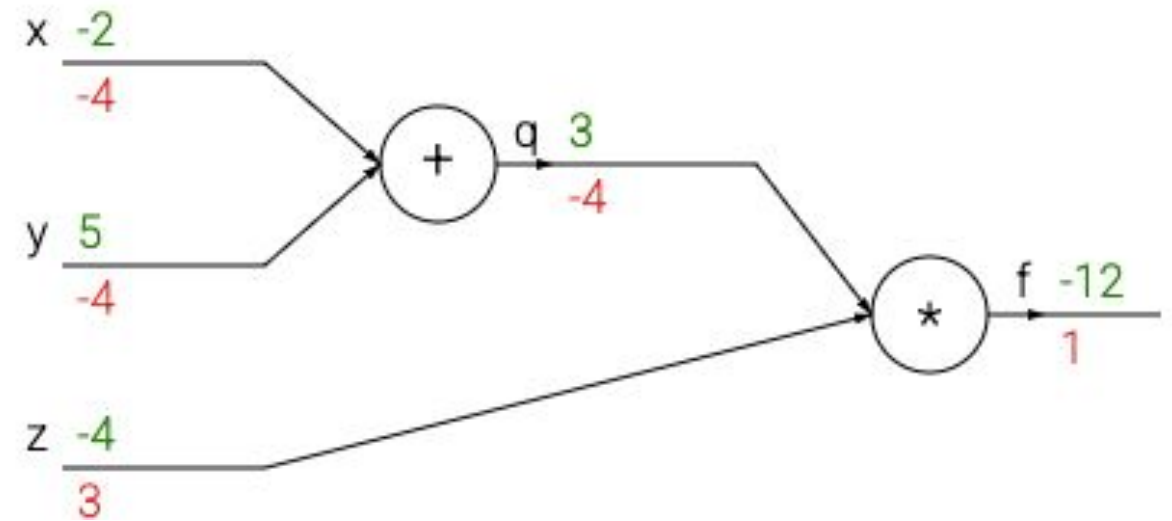
# This will print out the architecture, including names of layers
model
```

```
(fc): Linear(in_features=2048, out_features=1000, bias=True)
```

Step 2: freeze the base model

Image from <https://cs231n.github.io/optimization-2/>

```
# Freeze the base model  
for param in model.parameters():  
    param.requires_grad = False  
  
# Optionally, unfreeze specific layers  
for param in model.fc.parameters():  
    # Example: unfreeze the fully connected layer  
    param.requires_grad = True
```



The gradient computation graph is **only** computed on layers with `requires_grad=True`
Only these weights are updated during training

Step 3: add extra layers with the new target

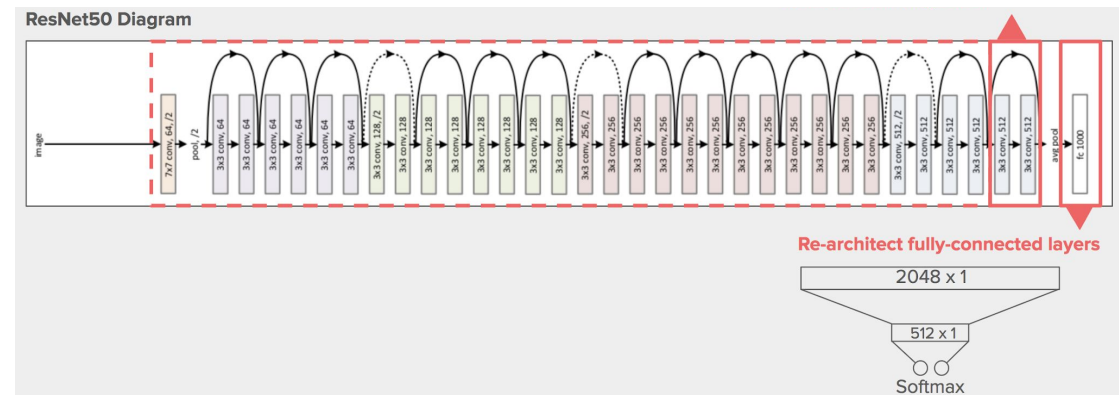
```
num_new_classes = 6
```

Approach 1: Add layer after classifier

```
def add_layer_model():  
    model = resnet50(weights=weights)  
    model.fc = nn.Sequential(  
        model.fc,  
        nn.ReLU(),  
        nn.Linear(1000, num_new_classes)  
    )  
    return model
```

Approach 2: Replace classifier

```
def replace_layer_model():  
    model = resnet50(weights=weights)  
    num_features = model.fc.in_features  
    model.fc = nn.Linear(num_features, num_new_classes)  
    return model
```



Fine-tuning by unfreezing the convolutional base

Unfreeze specific layers

```
for param in model.fc.parameters():
```

Example: unfreeze the fully connected layer

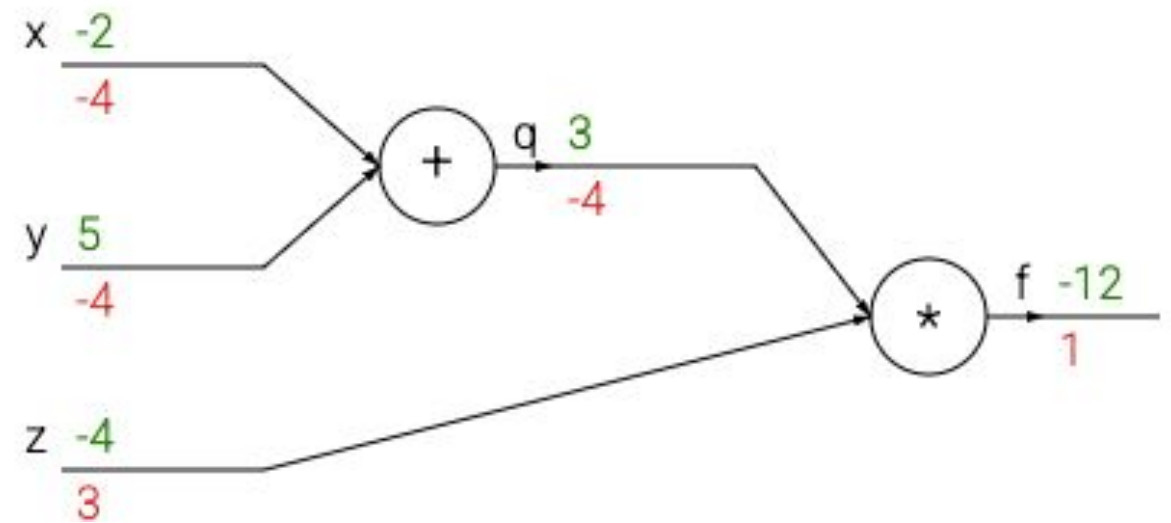
```
    param.requires_grad = True
```

*... # train for some epochs,
acquire low validation loss*

Unfreeze the whole model

to perform fine tuning

```
model.requires_grad_(True)
```



- After updating the final layers with non-random weights, we unfreeze the whole network and train it again
- This two stage process avoids “**catastrophic forgetting**”

Summary

Trained neural networks have hierarchical knowledge

- They learn basic shapes on their first layers and complex patterns on their last layers
- Models pretrained on rich datasets like ImageNet learn universal visual patterns that are useful for many tasks

Model Adaptation

- Modifying network architecture requires strategic choices
- We can add new final layers to preserve more ImageNet knowledge or replace the original final layers for a different representation

Progressive training prevents “catastrophic forgetting”

- Start by freezing the base model, train new layers, then fine-tune all layers for optimal results

SPONSORED BY THE

Further reading and references

ImageNet

- <https://en.wikipedia.org/wiki/ImageNet>

Visualizing and understanding convolutional networks

- <https://arxiv.org/pdf/1311.2901>

Resnet50 in torchvision models

- <https://pytorch.org/vision/main/models/generated/torchvision.models.resnet50.html>