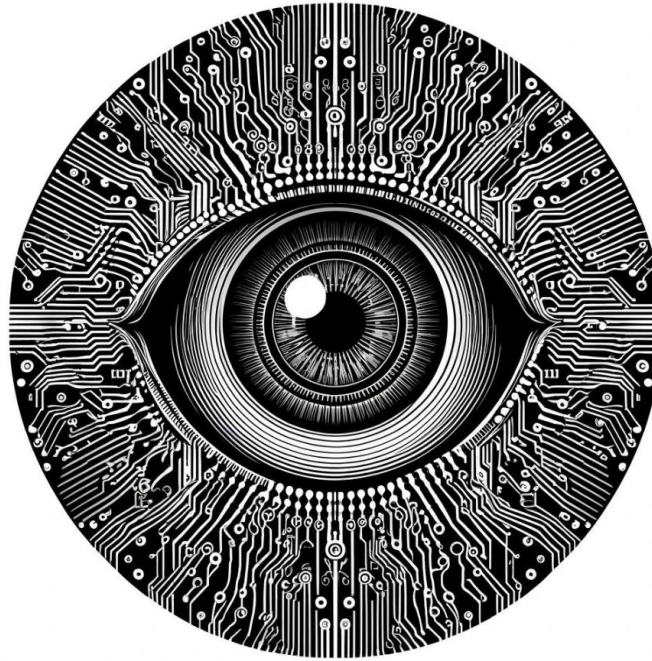


Contrastive Language-Image Pretraining (CLIP)



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Understand how CLIP models are trained
- Classify images with CLIP using a zero-shot approach
- Create image and text embeddings with a pretrained CLIP model
- Discuss CLIP's uses and limitations

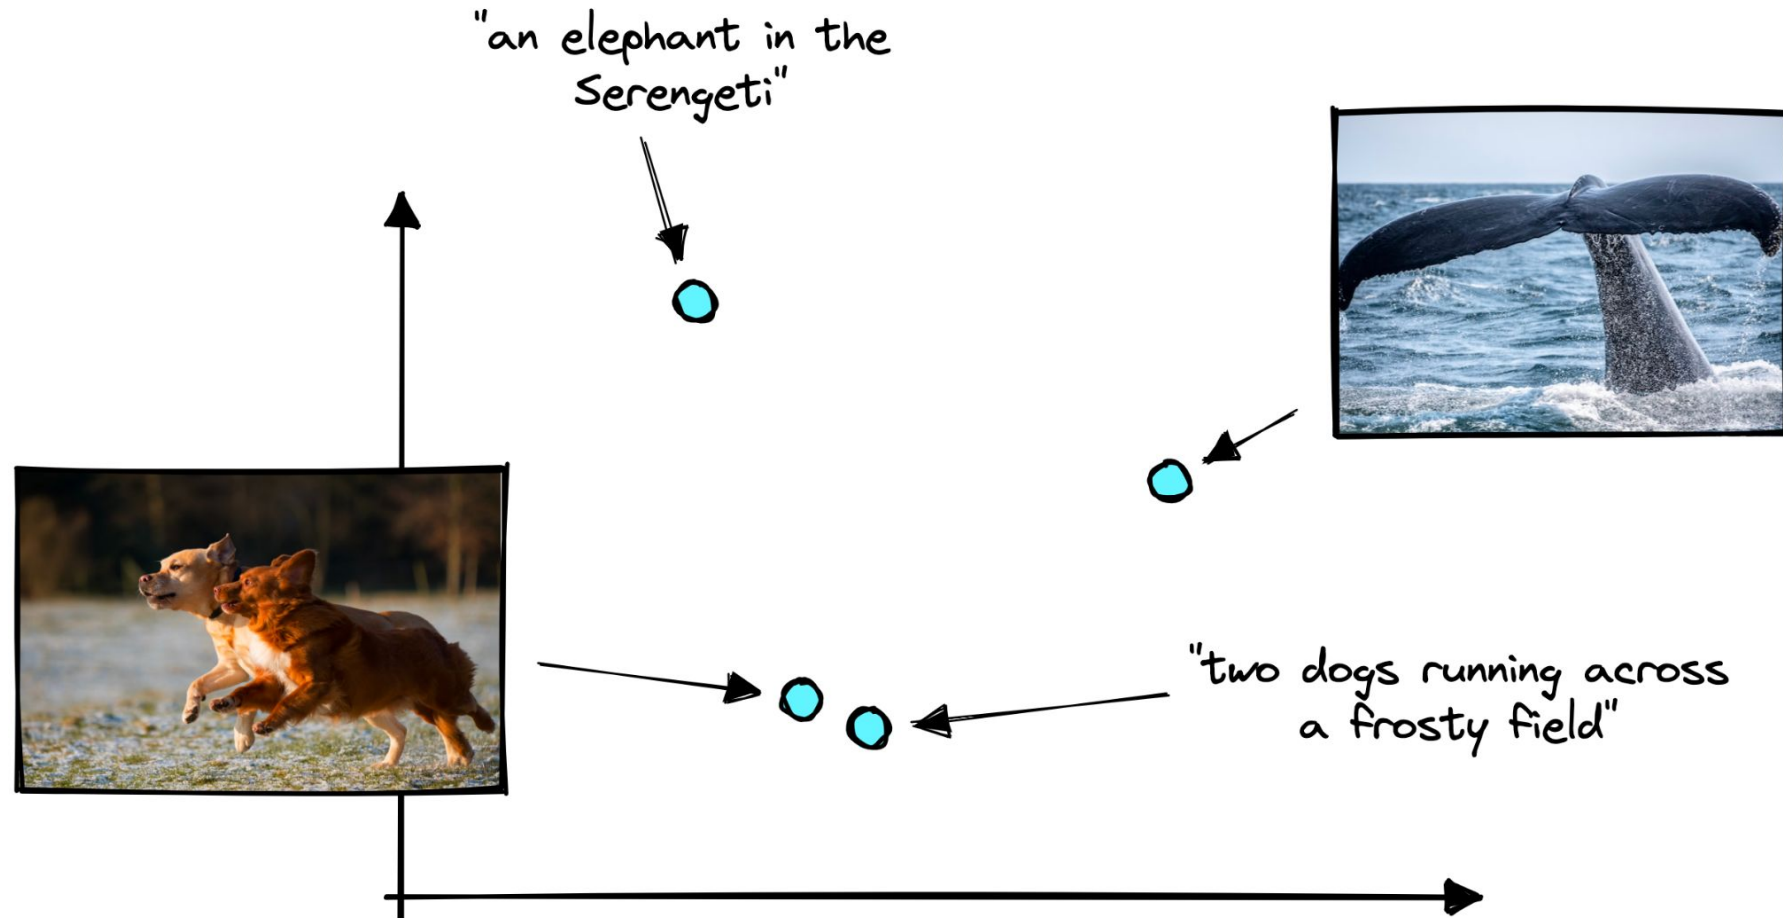
CLIP: ‘Contrastive Language Image Pretraining’

Learning Transferable Visual Models From Natural Language Supervision

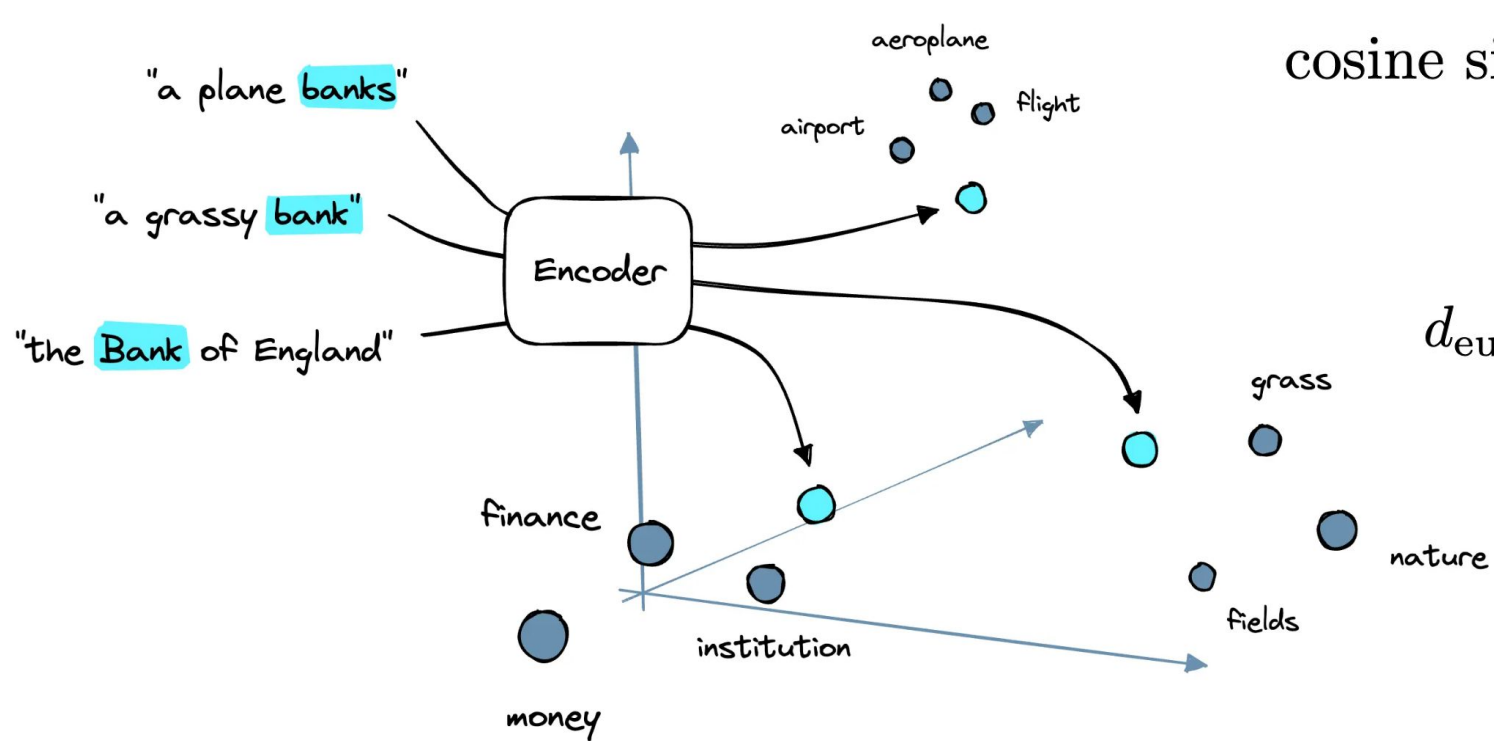
Alec Radford^{*1} Jong Wook Kim^{*1} Chris Hallacy¹ Aditya Ramesh¹ Gabriel Goh¹ Sandhini Agarwal¹
Girish Sastry¹ Amanda Askell¹ Pamela Mishkin¹ Jack Clark¹ Gretchen Krueger¹ Ilya Sutskever¹

- Connects text and images in a shared embedding space
- Created by OpenAI in 2021
- Trained on **400M image-textual description pairs** scraped from the internet
- Predicts the most relevant text snippet given an image, enabling “zero-shot classification”

Aligning text and image embeddings



Text encoders

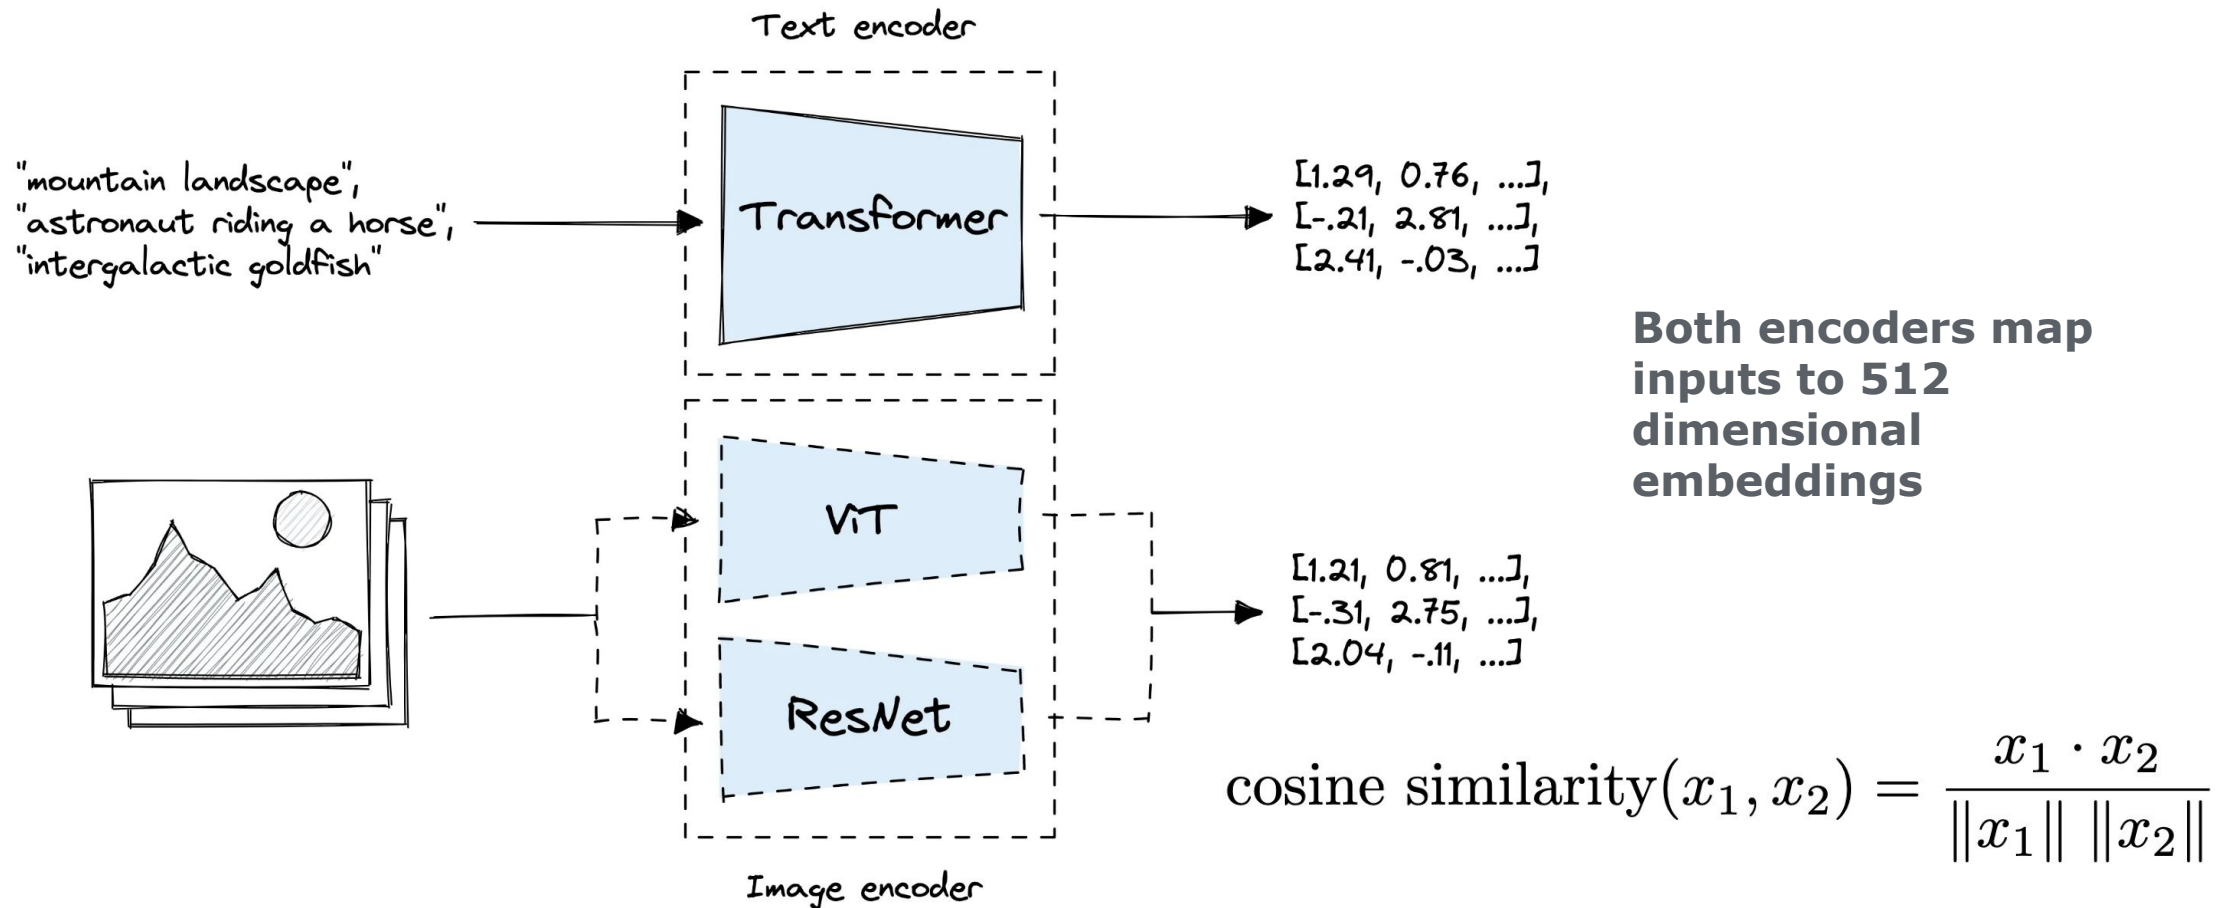


$$\text{cosine similarity}(x_1, x_2) = \frac{x_1 \cdot x_2}{\|x_1\| \|x_2\|}$$

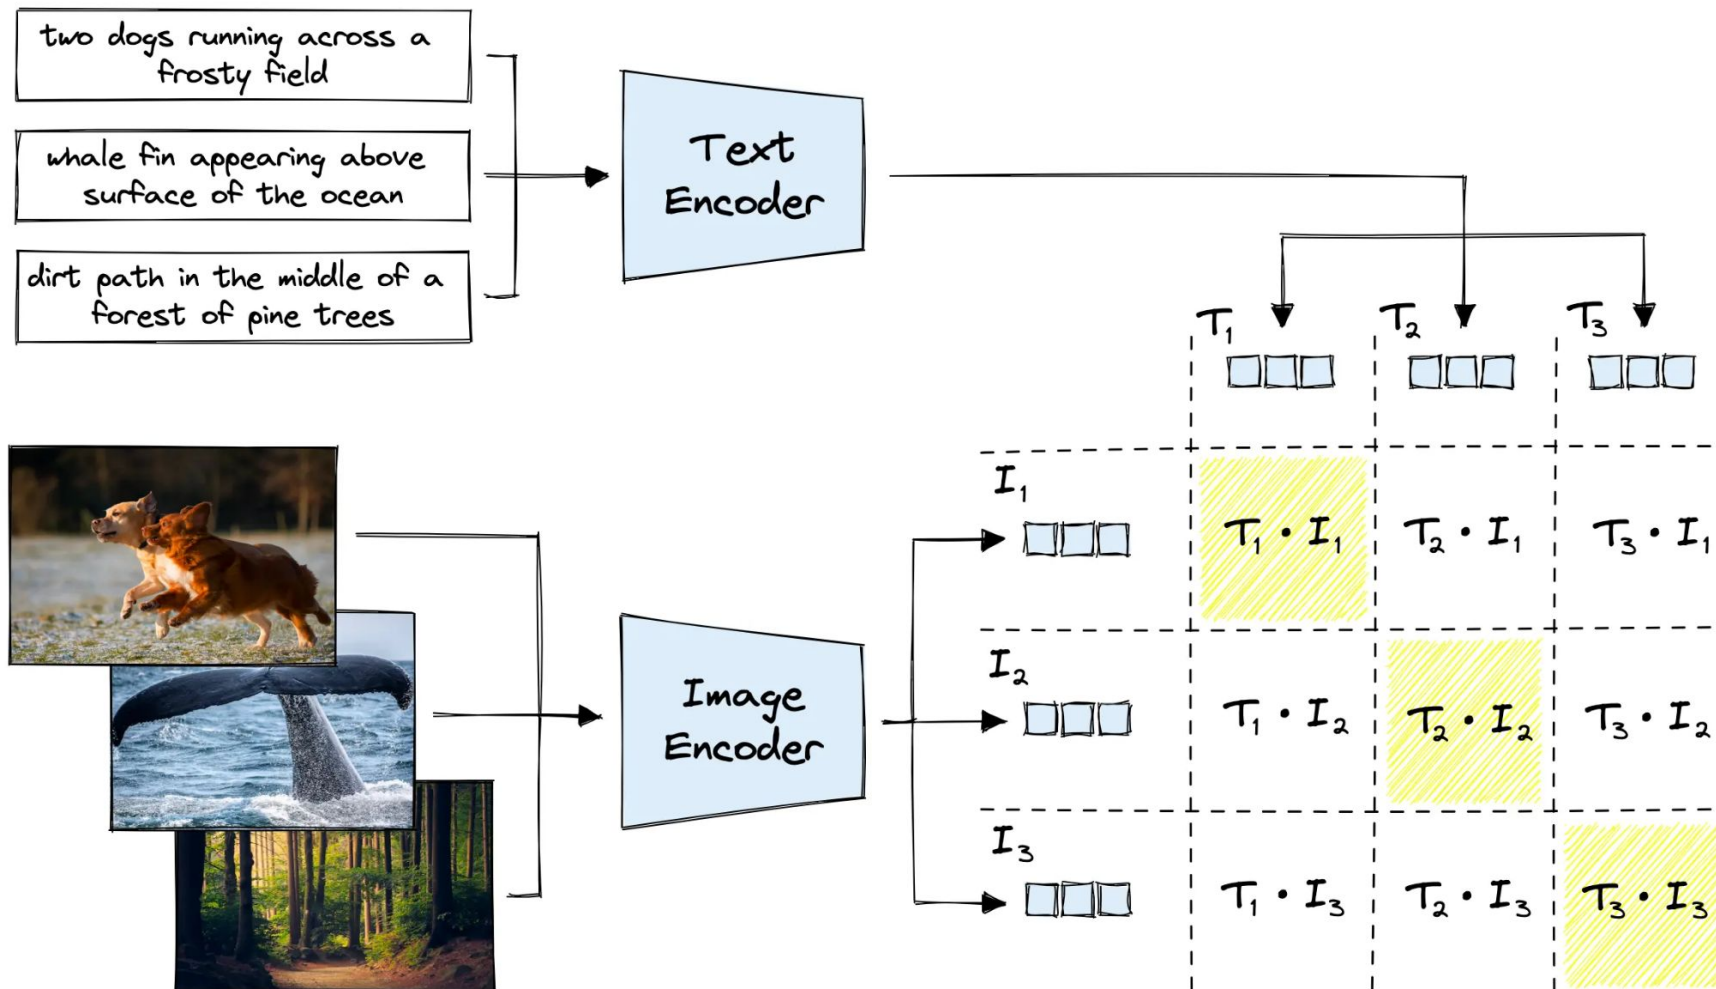
$$d_{\text{euclid}}(x_1, x_2) = \sqrt{\sum_{i=1}^n (x_{1i} - x_{2i})^2}$$

Image from <https://www.pinecone.io/learn/series/image-search/vision-transformers/>

CLIP's architecture



Maximizing cosine similarity of matching text and image embeddings



Training algorithm

```
# image_encoder - ResNet or Vision Transformer
# text_encoder - CBOW or Text Transformer
# I[n, h, w, c] - minibatch of aligned images
# T[n, l] - minibatch of aligned texts
# W_i[d_i, d_e] - learned proj of image to embed
# W_t[d_t, d_e] - learned proj of text to embed
# t - learned temperature parameter

# extract feature representations of each modality
I_f = image_encoder(I) #[n, d_i]
T_f = text_encoder(T) #[n, d_t]

# joint multimodal embedding [n, d_e]
I_e = l2_normalize(np.dot(I_f, W_i), axis=1)
T_e = l2_normalize(np.dot(T_f, W_t), axis=1)

# scaled pairwise cosine similarities [n, n]
logits = np.dot(I_e, T_e.T) * np.exp(t)

# symmetric loss function
labels = np.arange(n)
loss_i = cross_entropy_loss(logits, labels, axis=0)
loss_t = cross_entropy_loss(logits, labels, axis=1)
loss = (loss_i + loss_t)/2
```

Figure 3. Numpy-like pseudocode for the core of an implementation of CLIP.

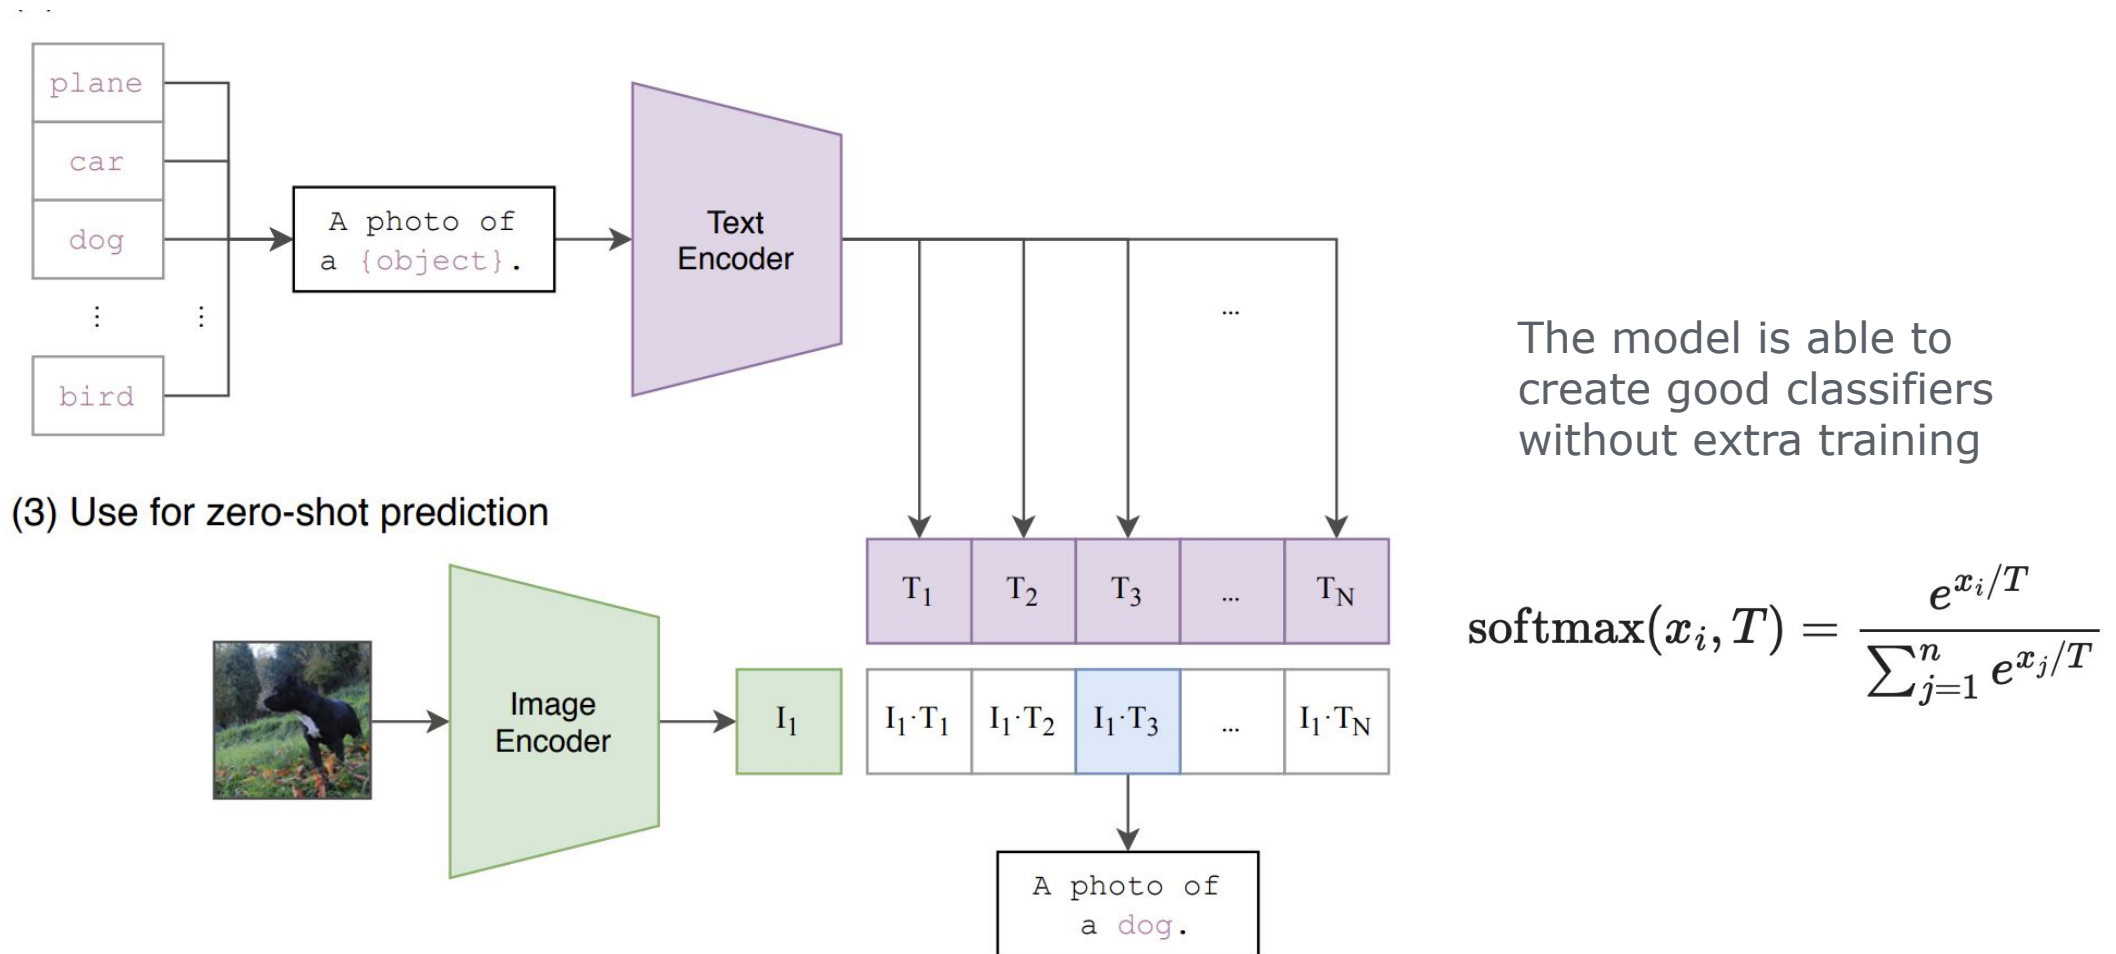
	T ₁	T ₂	T ₃	...	T _N
I ₁	I ₁ ·T ₁	I ₁ ·T ₂	I ₁ ·T ₃	...	I ₁ ·T _N
I ₂	I ₂ ·T ₁	I ₂ ·T ₂	I ₂ ·T ₃	...	I ₂ ·T _N
I ₃	I ₃ ·T ₁	I ₃ ·T ₂	I ₃ ·T ₃	...	I ₃ ·T _N
⋮	⋮	⋮	⋮	⋮	⋮
I _N	I _N ·T ₁	I _N ·T ₂	I _N ·T ₃	...	I _N ·T _N

$$H(p, q) = - \sum_i p(i) \log q(i)$$

$$H(q, p) = - \sum_i q(i) \log p(i)$$

$$\text{symmetric CE loss} = \frac{H(p, q) + H(q, p)}{2}$$

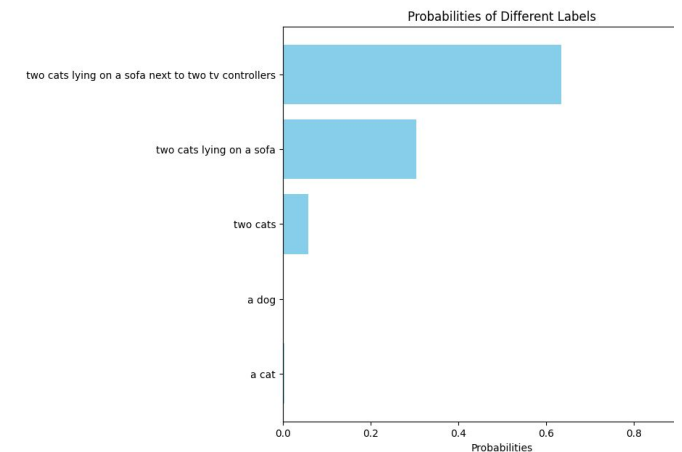
Zero-shot classification with CLIP



Producing embeddings with CLIP (1/2)

```
import torch
from transformers import CLIPProcessor, CLIPModel
from PIL import Image

# Load model and inputs
model =
CLIPProcessor.from_pretrained("openai/clip-vit-base-patch32")
processor =
CLIPModel.from_pretrained("openai/clip-vit-base-patch32")
image = Image.open("cats.jpg")
cat_text = ['a cat', 'a dog', 'two cats',
            'two cats lying on a sofa',
            ""two cats lying on a sofa
            next to two tv controllers""]
```



Producing embeddings with CLIP (2/2)

Get embeddings

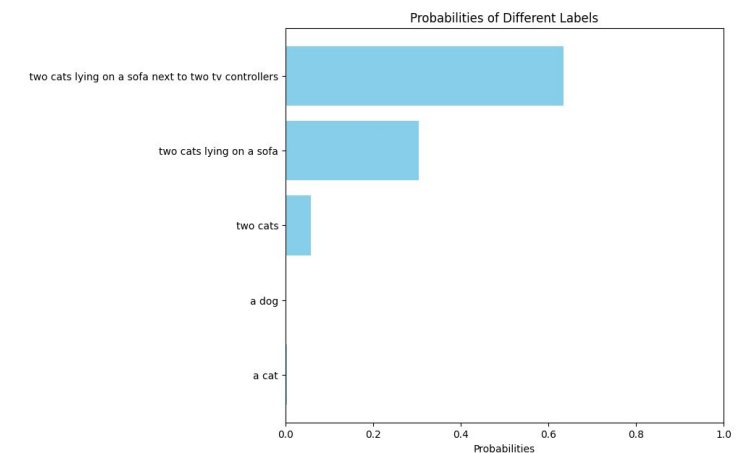
```
inputs = processor(text=cat_text, images=image, return_tensors="pt",  
padding=True)  
outputs = model(**inputs)
```

Calculate similarities

```
sims = torch.nn.functional.cosine_similarity(  
    outputs.image_embeds[:, None],  
    outputs.text_embeds[None, :],  
    dim=-1  
)
```

Print results

```
for text, sim in zip(cat_text, sims[0]):  
    print(f"{text}: {sim:.3f}")
```

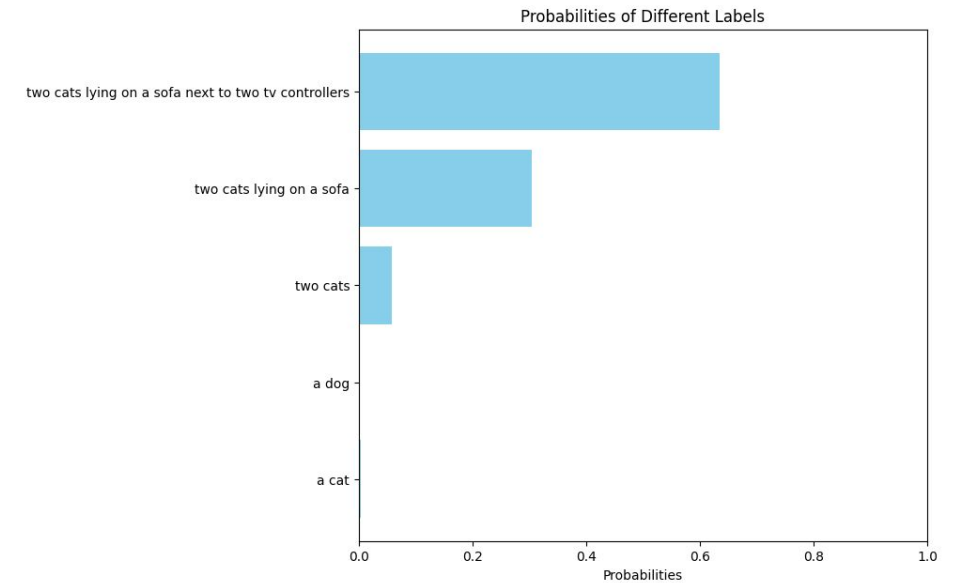


Zero-shot classification with CLIP

















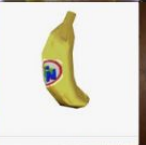






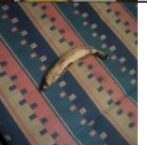
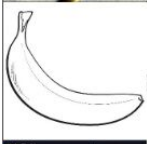






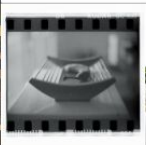


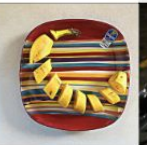

```
# Notice what happens with the output probabilities
# when we make the labels more specific
cat_text = ['a cat',
            'a dog',
            'two cats',
            'two cats lying on a sofa',
            'two cats lying on a sofa next to two tv controllers']
```

```
inputs = processor(text=cat_text,
                  images=[cats_img],
                  return_tensors="pt", padding=True)
cat_outputs = model(**inputs)
cat_logits_per_image = cat_outputs.logits_per_image
cat_probs = (cat_logits_per_image/temperature).softmax(dim=1)
```

$$\text{softmax}(x_i, T) = \frac{e^{x_i/T}}{\sum_{j=1}^n e^{x_j/T}}$$



Transferable representations: CLIP against a ResNet101 pretrained on Imagenet

	Dataset Examples						ImageNet ResNet101	Zero-Shot CLIP	Δ Score
ImageNet							76.2	76.2	0%
ImageNetV2							64.3	70.1	+5.8%
ImageNet-R							37.7	88.9	+51.2%
ObjectNet							32.6	72.3	+39.7%
ImageNet Sketch							25.2	60.2	+35.0%
ImageNet-A							2.7	77.1	+74.4%

Note that CLIP was trained on a dataset of **400 million images** scraped from the Internet

The Imagenet1K dataset has only **1.28 million training images**

CLIP is a much more capable “foundation model”

Resnet101

~44.5 million parameters

CLIP ViT-L/14@336px:

~428 million parameters

Training Resnet101: about 90 V100 GPU hours (4 days)

Training CLIP ViT-L/14: about ~16,000 V100 GPU hours (666 days)

Limitations against fully supervised models

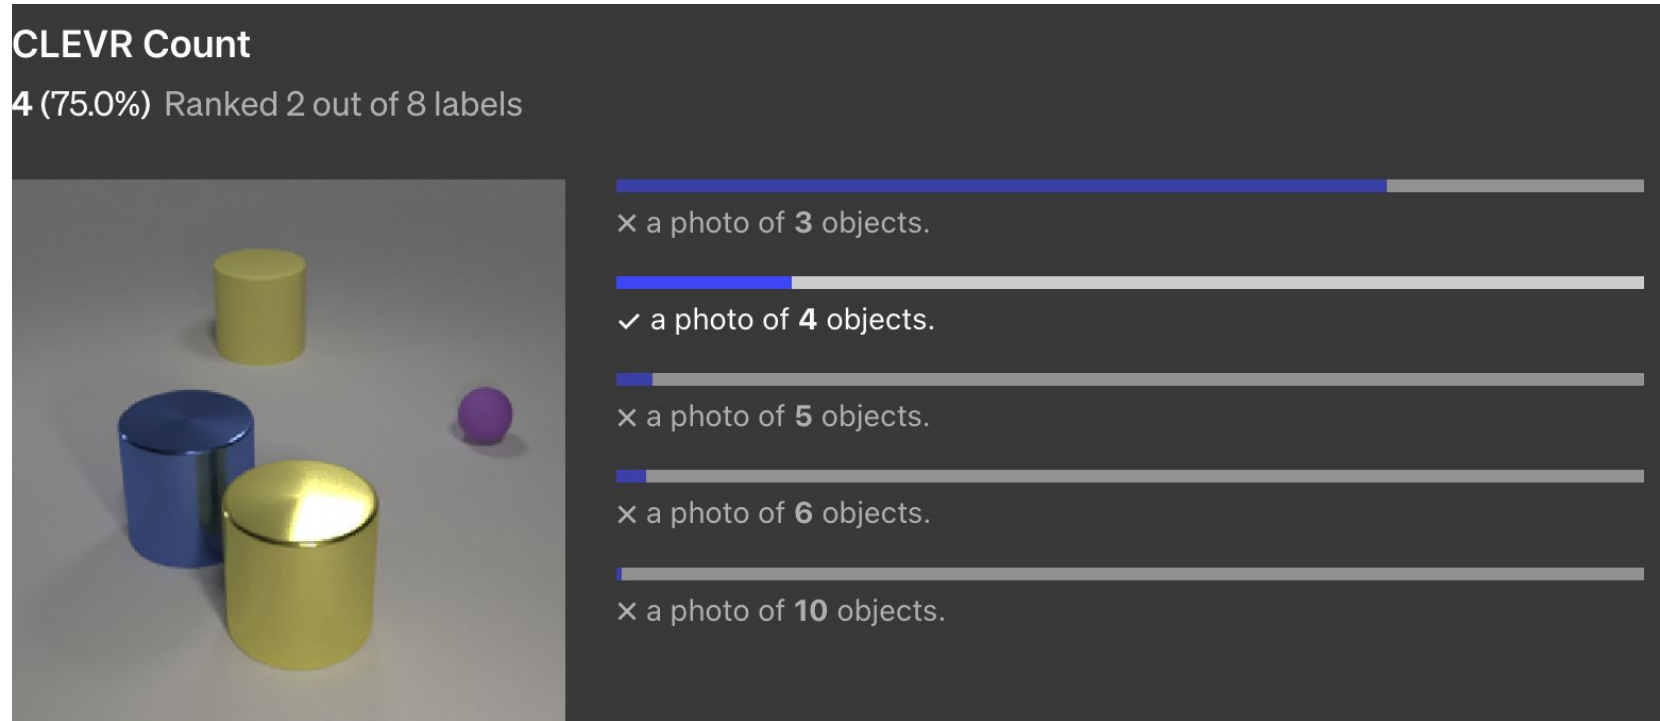


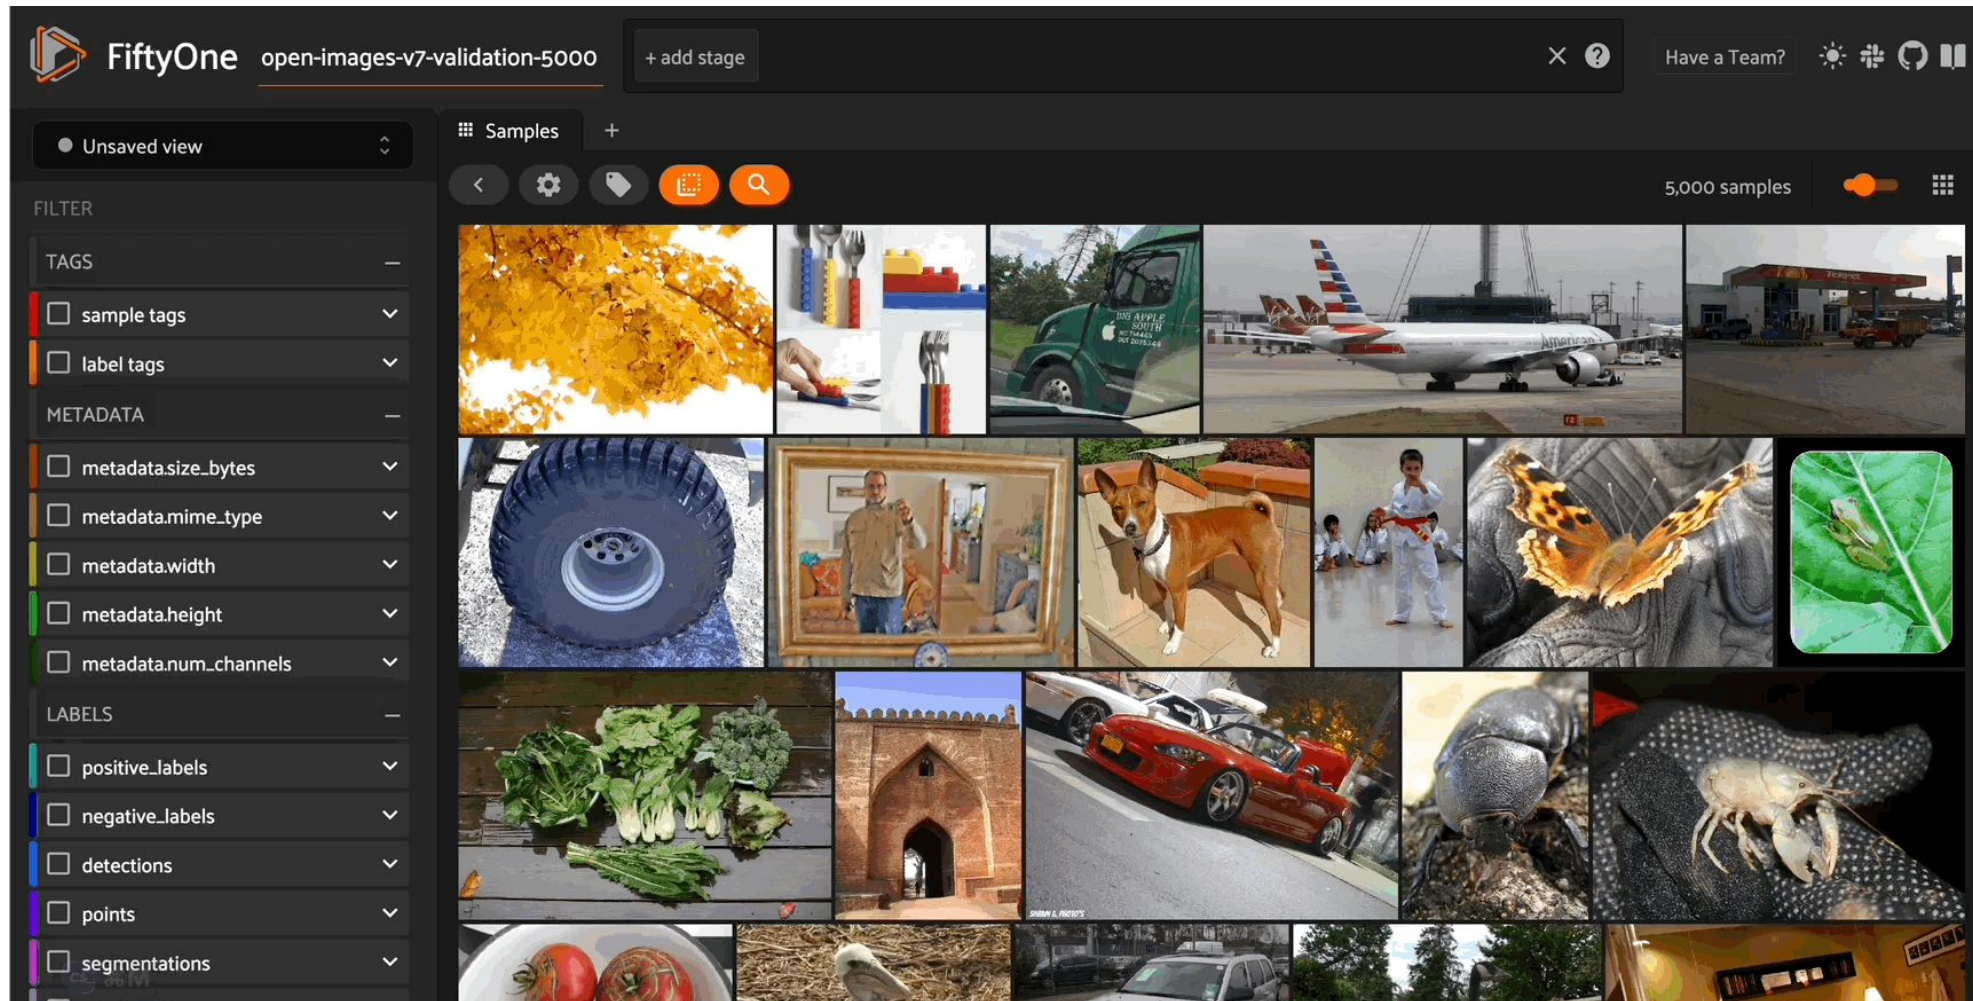
Image from
<https://openai.com/index/clip/>

The authors comment on the limitations section of the paper that CLIP often fails to perform as well as a fully supervised model fine tuned for the task.

Notable examples:

- digit recognition (MNIST)
- tumor classification (PatchCamelyon)
- satellite imaging (EuroSAT)
- object counting (CLEVR Count)

Semantic search with CLIP



CLIP guides image generation of diffusion models

Prompt:

"Create an image
of an astronaut
riding a horse in
pencil drawing
style"





Image by OpenAI's Dall-E 3

Summary

CLIP learns joint embeddings

- CLIP creates a shared embedding space for images and text
- These multimodal embeddings have applications on semantic search and image generation

CLIP excels at zero-shot classification

- CLIP performs well on unseen tasks without additional training

CLIP has limitations

- CLIP is not always better than models fine-tuned for specific tasks
- Retraining CLIP with a ViT base is expensive both computationally and data-wise

Further reading

Learning Transferable Visual Models From Natural Language Supervision

- <https://arxiv.org/abs/2103.00020>

A Google Search Experience for Computer Vision Data

- <https://voxel51.com/blog/a-google-search-experience-for-computer-vision-data/>

Multi-modal ML with OpenAI's CLIP

- <https://www.pinecone.io/learn/series/image-search/clip/>

SPONSORED BY THE