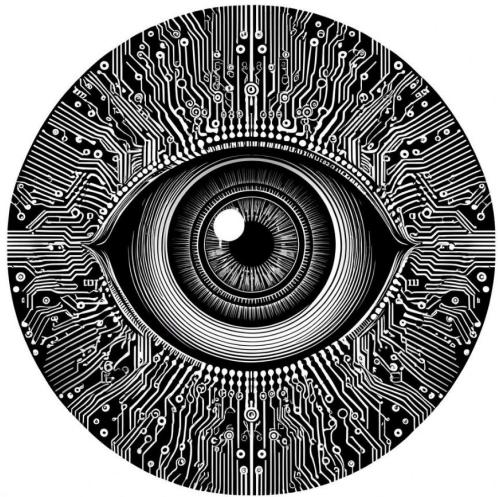


# Workshop 3 - Introduction to Image Classification



Antonio Rueda-Toicen

# About me

- AI Researcher at [Hasso Plattner Institute](#), AI Engineer & DevRel for [Voxel51](#)
- Organizer of the [Berlin Computer Vision Group](#)
- Instructor at [Nvidia's Deep Learning Institute](#) and Berlin's [Data Science Retreat](#)
- Preparing a [MOOC for OpenHPI](#) (to be published in May)



[LinkedIn](#)

# Agenda

- [Building a feedforward network for classification in PyTorch](#)
- [Metrics for classification and experiment tracking](#)
- [PyTorch's Dataset and DataLoader](#)

## Notebooks

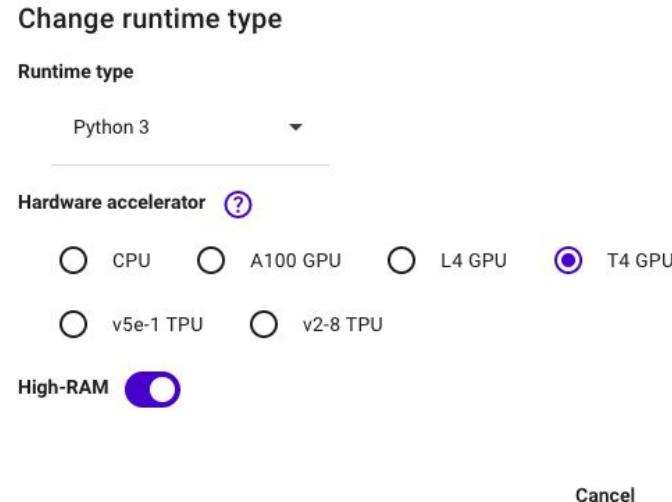
- Classifying breeds of cats and dogs
  - [Colab notebook](#)
- Digit recognition on the MNIST dataset (competition starter)
  - [Kaggle notebook](#)

# Setup for today

- [Setting up a Weights and Biases API token](#)
- [Setting up free GPU and TPU access in Kaggle Notebooks](#)

# Google Colab Runtime Setup for Training

- Runtime -> Change Runtime type



# Google Colab Setup for wandb API keys

- Go to [wandb.ai/authorize](https://wandb.ai/authorize) to find your API key
- Go to Secrets -> “Add New Secret”  
in your Colab environment

☰ Secrets

ⓘ Configure your code by storing environment variables, file paths, or keys. Values stored here are private, visible only to you and the notebooks that you select.

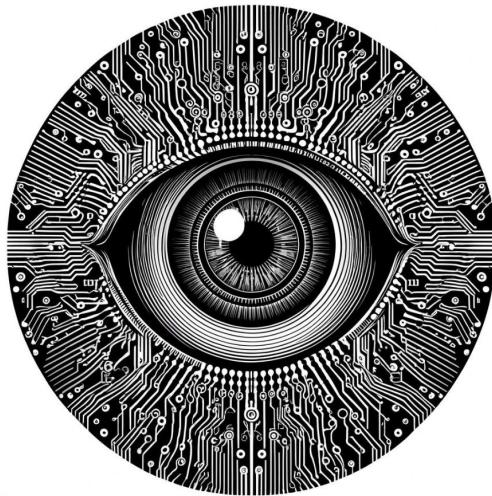
<> Secret name cannot contain spaces.

{x} Notebooks Name Value Actions

Notebook access	Name	Value	Actions
OFF	FIFTYONE_CV/	.....	👁️ 🗑️ 🗑️
OFF	FIFTYONE_CV/	.....	👁️ 🗑️ 🗑️
OFF	HF_TOKEN	.....	👁️ 🗑️ 🗑️
OFF	MAPILLARY_TC	.....	👁️ 🗑️ 🗑️
OFF	OPENAI_API_K	.....	👁️ 🗑️ 🗑️
OFF	kaggle_token	.....	👁️ 🗑️ 🗑️
OFF	kaggle_username	.....	👁️ 🗑️ 🗑️
OFF	nvidia-ngc-api	.....	👁️ 🗑️ 🗑️
OFF	qdrant_api	.....	👁️ 🗑️ 🗑️
OFF	synapse_client_	.....	👁️ 🗑️ 🗑️
OFF	synapse_token	.....	👁️ 🗑️ 🗑️
ON	wandb_api	.....	👁️ 🗑️ 🗑️

+ Add new secret

# Classifying Pets According to Breed



# What is the breed of my dog?



# The Oxford-IIIT Pets dataset

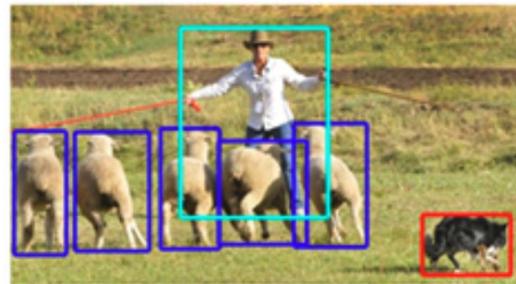


[Oxford-IIIT Pet Dataset](#)

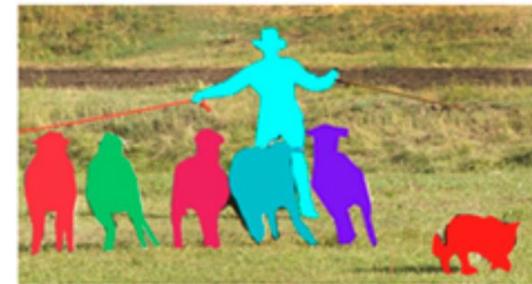
# Localized vs Non-localized Image Classification



Classification



Object Detection



Instance Segmentation

# Apps that classify pet breeds

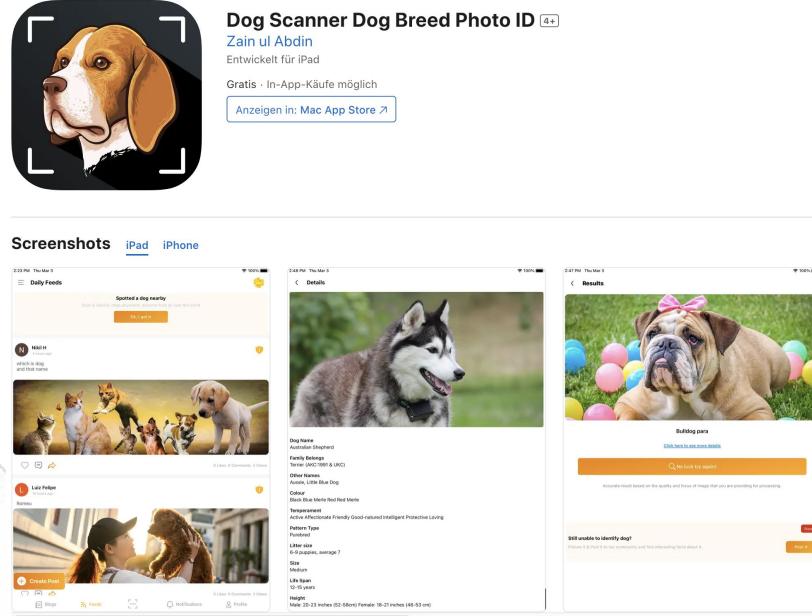
App Store Vorschau

Öffne den Mac App Store, um Apps zu kaufen und zu laden.

**Dog Scanner Dog Breed Photo ID** 4+

Zain ul Abdin  
Entwickelt für iPad

Gratis · In-App-Käufe möglich  
[Anzeigen in: Mac App Store](#)



The screenshot shows the app's main interface. It features a camera icon with a frame over a dog's face. Below the camera are three cards: 'Daily Feeds' (listing 'Spotted a dog nearby'), 'Details' (listing 'Dog Name: Australian Shepherd', 'Family Belongs: Herding (Sheepdog & HIC)', 'Other Names: Husky', 'Color: Black/Blue Merle Red/Fawn Merle', 'Age: 12-15 years', 'Height: Male: 25-28 inches (63.5cm) Female: 18-21 inches (45.5cm)'), and 'Results' (listing 'Bulldog paws' with a 'Click here to see more details' button). At the bottom, there's a note: 'Sorry unable to identify dog!' followed by 'Please take a clear photo of your dog and try again later'.

**DogSnap:Dog breed scanner&Care**

Breed Identifier,Photo Scanner  
Designed for iPhone. Not verified for macOS.

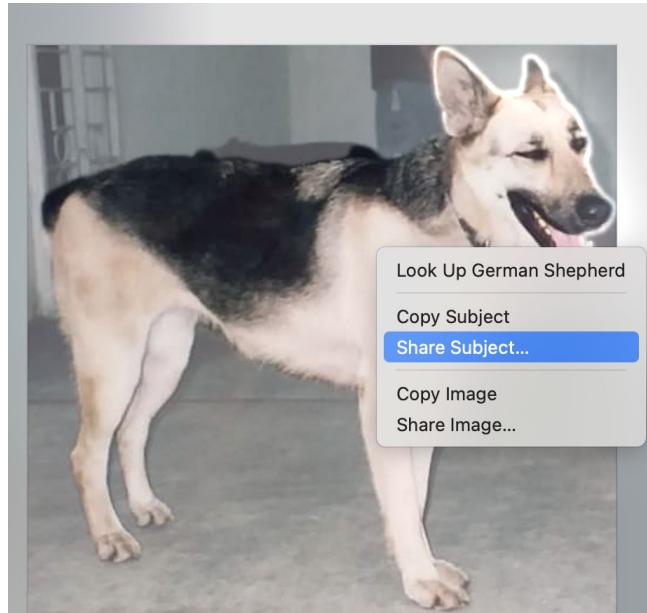
[Get](#) In-App Purchases

AGE	CATEGORY	DEVELOPER	LANGUAGE	SIZE
4+ Years Old	Education	HICLUB, INC.	EN English	155,4 MB



The promotional images for DogSnap show the app's key features: 1. Intelligent Identification For Dog: A golden retriever with text overlay. 2. What Breed Is Your Dog?: An interface showing a white dog with a camera icon. 3. Human-dog translator: A screen showing a dog and a person communicating. 4. Bring photo of dog to life: A screen showing a dog wearing a bear mask.

# Image lens in MacOS's Finder





# DATASET

Tags Names Classes Description ...

## SAMPLE 0

ID Filepath Tags Label Field

## SAMPLE 1

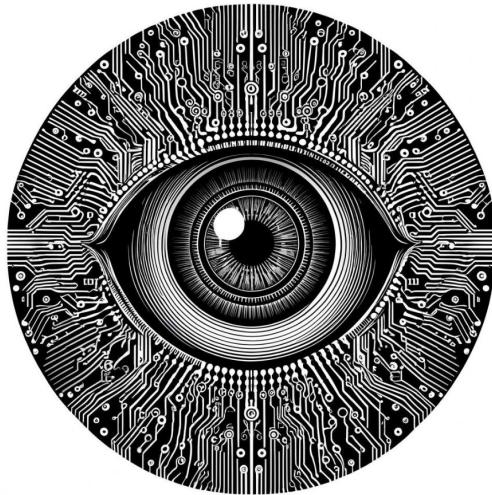
ID Filepath Tags Label Field

## SAMPLE 2

ID Filepath Tags Label Field

How we create a  
FiftyOne dataset

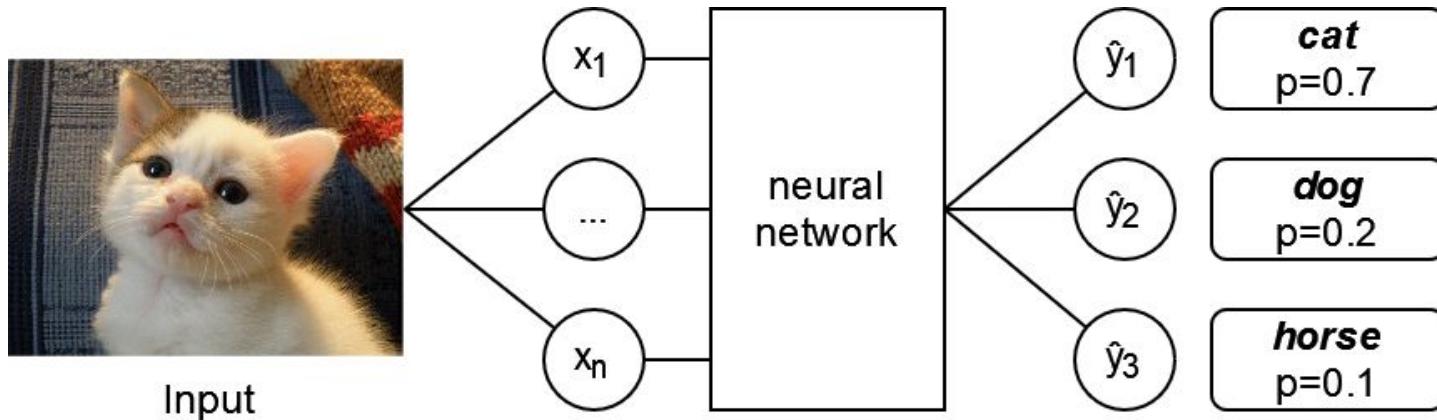
# Building a Feedforward Network for Classification in PyTorch



# Learning goals

- Create a multilayer neural network for classification in PyTorch
- Gain familiarity with the nn.Module syntax for network creation
- Understand usage of the softmax activation function
- Develop intuitions on Categorical Cross Entropy
- Use the Adam variant of stochastic gradient descent

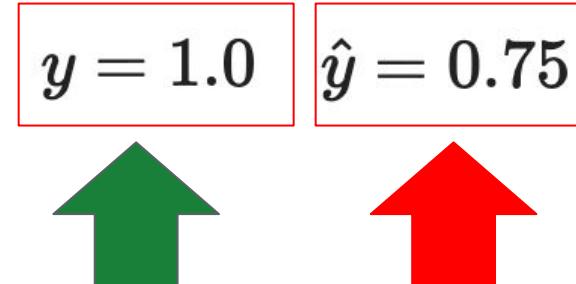
# Multiclass classification: n probabilities for n classes



**Classification**  $P(\hat{y}_{\text{cat}}) + P(\hat{y}_{\text{dog}}) + P(\hat{y}_{\text{horse}}) = 1$

## Cross entropy loss

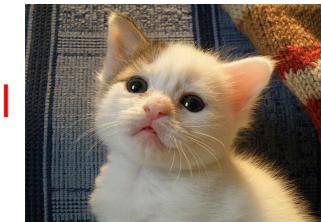
$$H(p, q) = -\frac{1}{n} \sum_{i=1}^n p_i(x) \log q_i(x)$$



Labels  $y$  map to  $p_i(x)$ , predictions  $\hat{y}$  map to  $q_i(x)$ . Suppose that we show the network, only the following image ( $n=1$ ).

“one hot encoding” 🔥 = only one true label

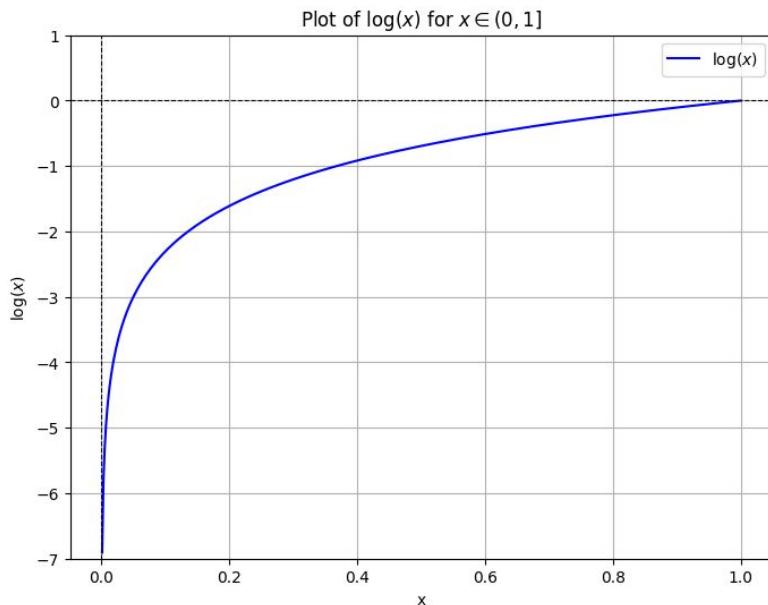
The label for the image is encoded as  $y = [1, 0, 0]$  ( $i = 0 = \text{cat}$ ,  $i = 1 = \text{dog}$ ,  $i = 2 = \text{horse}$ )



- First the network outputs  $\hat{y} = [0.75, 0.25, 0.0]$
- Then the network outputs  $\hat{y} = [0.99, 0.01, 0.0]$  (if trained correctly)

# Intuition for the log function in cross entropy loss

$$y = [1, 0, 0] \quad \text{where } i = 0 \text{ (cat), } i = 1 \text{ (dog), } i = 2 \text{ (horse)}$$



$$\ln(1) = \log(1) = 0$$



$$y_0 = 1 \quad \text{and} \quad \hat{y}_0 = 1$$

Cross Entropy Loss =  $-1 \cdot \log(1) = 0$  (perfect prediction, no loss)

$$y_0 = 1 \quad \text{and} \quad \hat{y}_0 = 0.01$$

Cross Entropy Loss =  $-(1 \cdot \log(0.01)) = 4.65$  (bad prediction, high loss)

# Implementing cross entropy loss from probabilities for a single sample

```
import torch

# Example predicted probabilities from model
y_hat = torch.tensor([0.75, 0.25, 0.0])

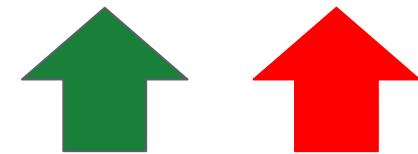
# True label: y = [1, 0, 0]
y = torch.tensor([1, 0, 0], dtype=torch.float32)

def cross_entropy(y, y_hat):
    epsilon = 1e-15 # equal to 1 times 10^-15 (very small number)
    # Adding epsilon prevents log(0) which is undefined
    log_probs = torch.log(y_hat + epsilon)
    # Notice that the sum would only give a value different
    # than zero on the index of true_label
    # we are computing a single sample
    # so n = 1 (batch dimension)
    return -torch.sum(y * log_probs)

loss = cross_entropy(y, y_hat)
# Gives us tensor(0.287), not perfect, but not the worst possible
```



y = [1, 0, 0] y\_hat = [0.75, 0.25, 0.0]

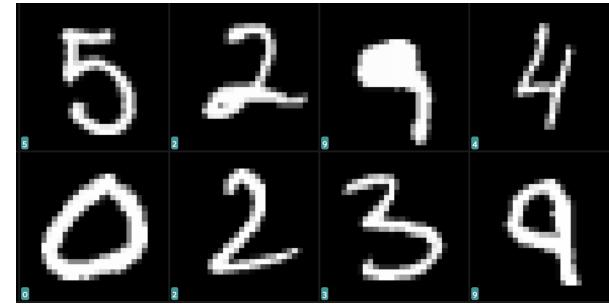


$$H(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^n y_i \log \hat{y}_i$$

# A minimal network for classification

```
import torch.nn as nn

# Model
model = nn.Sequential(
    # Flatten the input image
    nn.Flatten(),
    # The number of input features is the number of pixels in the image
    nn.Linear(in_features=28 * 28, out_features = 128),
    # We add a non-linearity
    nn.ReLU(),
    # We create 10 scores, aka 'logits', one for each class that we have
    # Notice that there is no ReLU after nn.Linear
    nn.Linear(in_features=128, out_features = 10))
```



**Tip:** do not confuse these “logits” with the function described on  
<https://en.wikipedia.org/wiki/Logit>

$$f(x) = \ln\left(\frac{x}{1-x}\right)$$

✗ these “logits” are **not** this ^

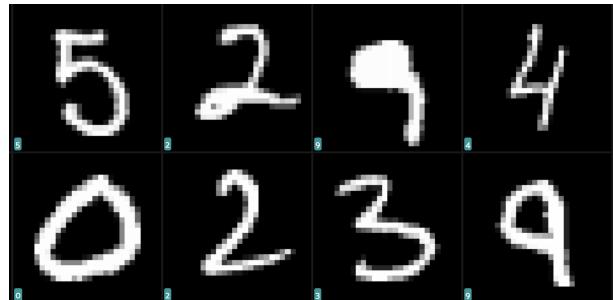
# Using the nn.Module syntax

```
import torch.nn as nn

class MNISTClassifier(nn.Module):
    def __init__(self):
        super().__init__()
        # Instead of Sequential, we define each module as a class attribute
        self.flatten = nn.Flatten()
        self.fc1 = nn.Linear(in_features=28 * 28, out_features=128)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(in_features=128, out_features=10)

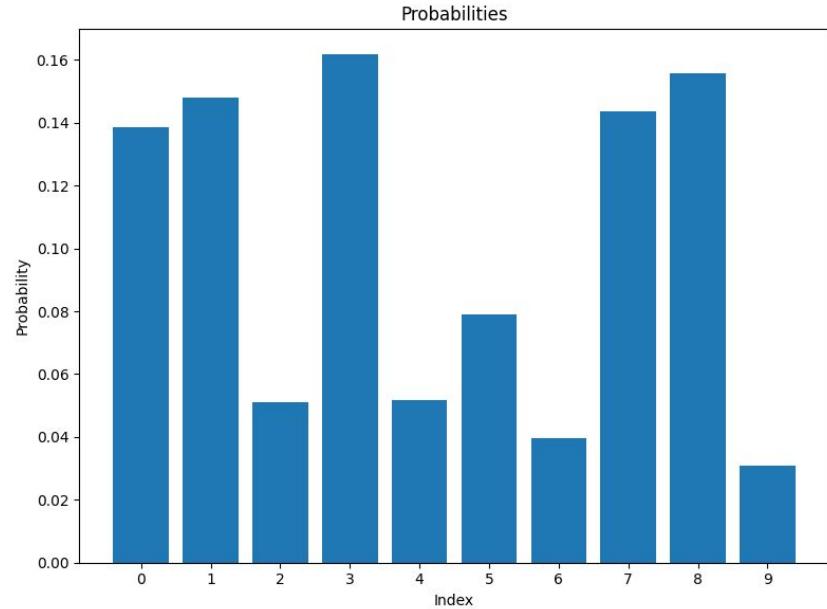
    def forward(self, x):
        # Define the forward pass and set a breakpoint
        import pdb; pdb.set_trace()
        x = self.flatten(x)
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x) # notice no relu here
        return x

# Create an instance of the model
model = MNISTClassifier()
```

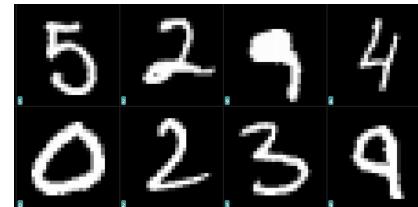


# The softmax function

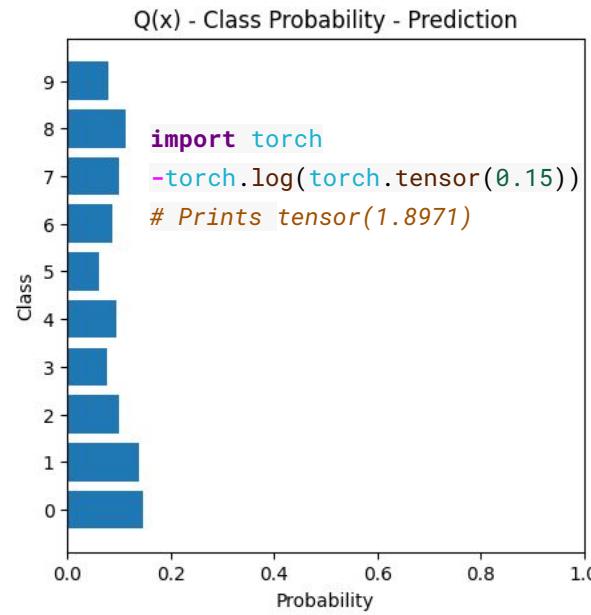
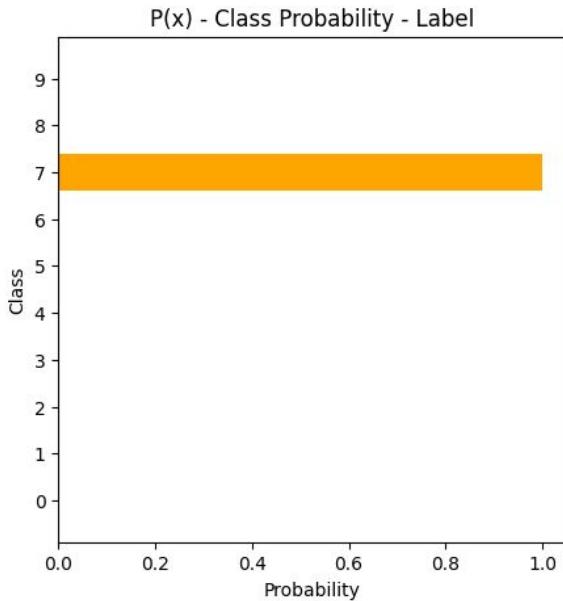
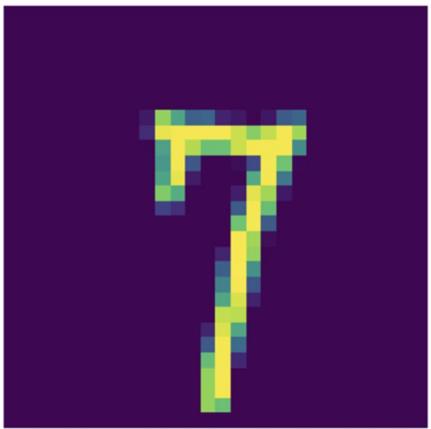
$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$



```
import torch.nn.functional as F
logits = torch.tensor([ 0.7645,  0.8300, -0.2343,  0.9186, -0.2191,  0.2018, -0.4869,  0.8000,  0.8815, -0.7336])
probabilities = F.softmax(logits, dim=0)
print(probabilities.sum()) # prints tensor(1.00)
```

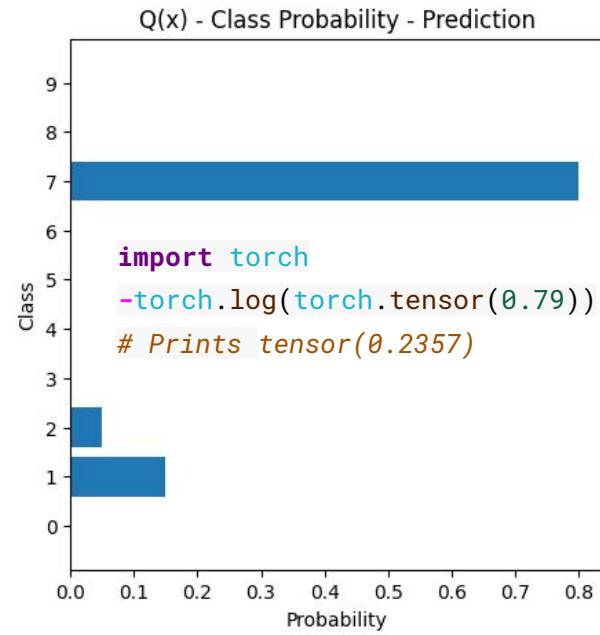
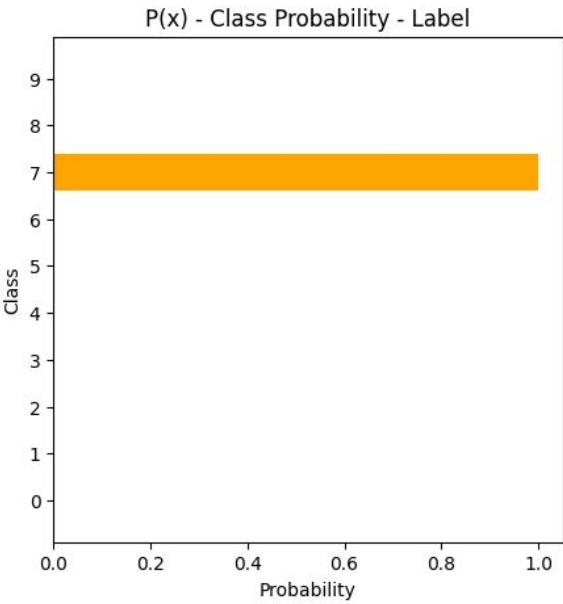
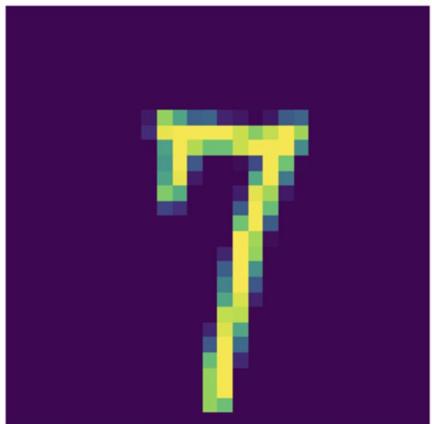


# High cross entropy loss - high disagreement between $P(x) = y$ and $Q(x) = \hat{y}$



$$H(p, q) = -\frac{1}{n} \sum_{i=1}^n p_i(x) \log q_i(x) = H(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^n y_i \log \hat{y}_i$$

# Low cross entropy loss - low disagreement between $y$ and $\hat{y}$



$$H(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^n y_i \log \hat{y}_i$$

Cross Entropy =  $-\log(\hat{y}_i)$   
with index  $i$  being the one true class

# Quirks of nn.CrossEntropyLoss

```
# Compute output
logits = model(input_image)

# Turn output into probabilities, useful for interpretation
probabilities = F.softmax(logits)

# CrossEntropyLoss will compute log-softmax on the raw logits
# because it is more numerically stable
ce_loss = nn.CrossEntropyLoss()
ce_loss(logits, labels)
```

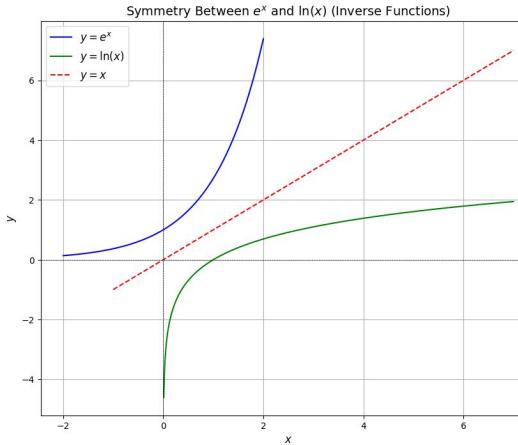
$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

$$\text{LogSoftmax}(\mathbf{x})_i = x_i - \log\left(\sum_{j=1}^n \exp(x_j)\right)$$

# Turning LogSoftmax into interpretable probabilities

$$p_i = \exp(\log p_i) = e^{\ln p_i}$$

```
import torch.nn.functional as F
logits = torch.tensor([ 0.7645,  0.8300, -0.2343,
                      0.9186, -0.2191,  0.2018,
                     -0.4869,  0.8000,  0.8815,
                     -0.7336])
log_probabilities = F.log_softmax(logits, dim=0)
print(log_probabilities.sum()) # prints tensor (-24.6857)
probabilities = torch.exp(log_probabilities)
print(probabilities.sum()) # prints tensor(1.00)
```



$$\text{Softmax}(x_i) = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

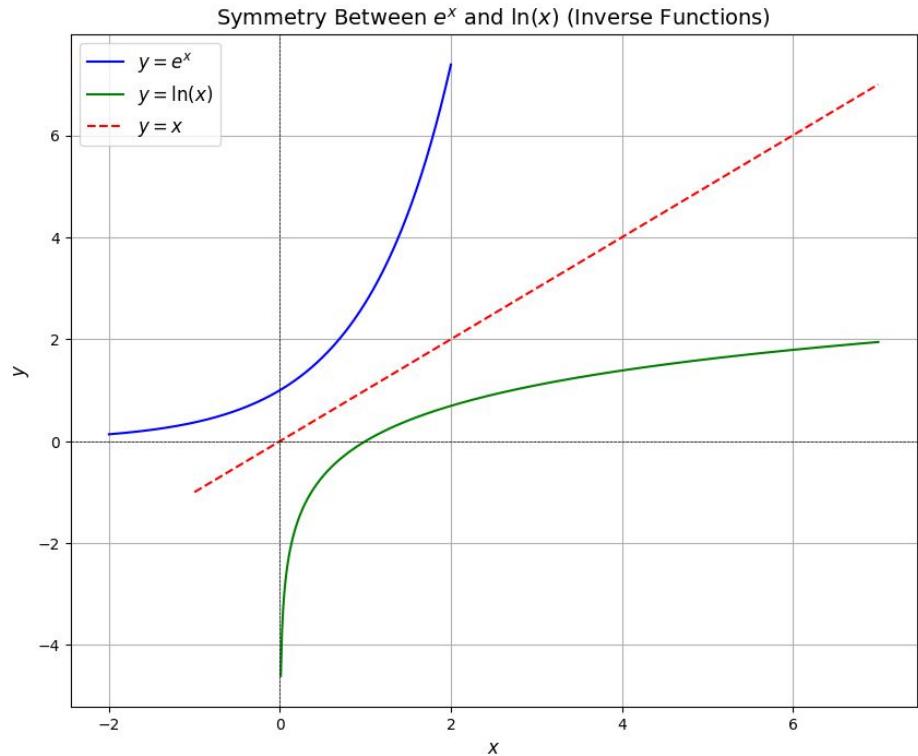
$$\text{LogSoftmax}(\mathbf{x})_i = x_i - \log\left(\sum_j^n \exp(x_j)\right)$$

# Why do we use $\log(x) = \ln(x)$ ?

$$p_i = \exp(\log p_i) = e^{\ln p_i}$$

$$\frac{d}{dx} e^x = e^x$$

$$\frac{d}{dx} \ln(x) = \frac{1}{x}$$



# Cross entropy loss from logits

```
import torch
import torch.nn.functional as F

def cross_entropy_loss(logits, labels):
    """
    Calculate cross entropy loss from logits using log softmax

    Args:
        logits: Tensor of shape (batch_size, num_classes) containing raw model
        outputs
        labels: Tensor of shape (batch_size,) containing class indices (0-9 for
        MNIST)

    Returns:
        loss: Scalar tensor with the mean loss
    """
    # Apply log softmax to get log probabilities
    log_probs = F.log_softmax(logits, dim=1)

    # Calculate negative log likelihood loss
    # This efficiently computes cross entropy without explicitly creating one-hot
    vectors
    loss = F.nll_loss(log_probs, labels)
```

**“negative of the log of the likelihood (of the correct class” = `F.nll_loss`**

$$\text{Cross Entropy} = -\log(\hat{y}_i)$$

with index  $i$  being the one true class

# Extra: Adam vs Stochastic Gradient Descent

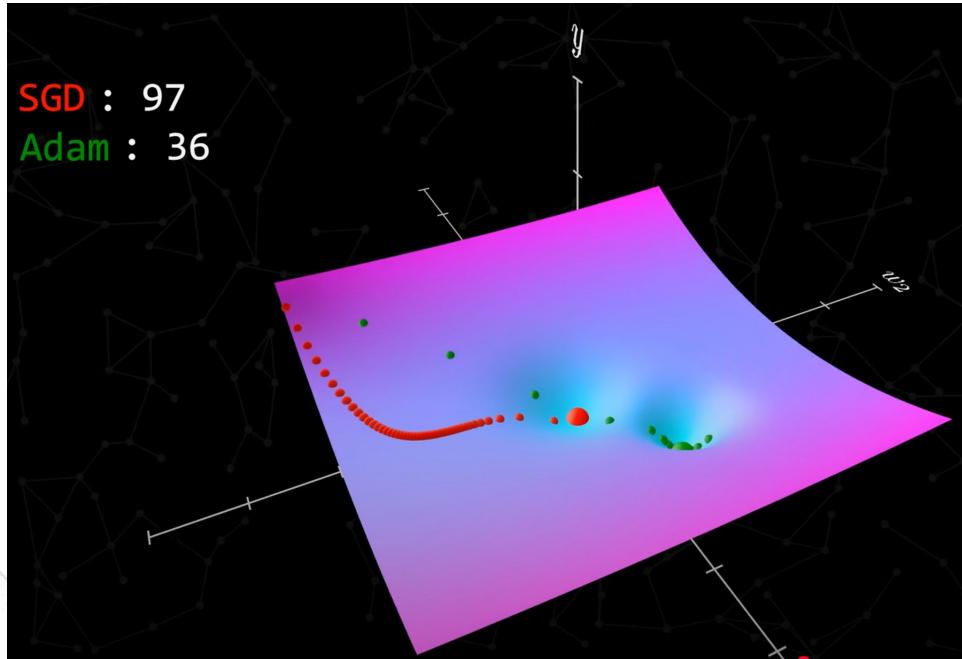


Image from "[Who is Adam and what is he optimizing?](#)"

Adaptive Gradients with Momentum (ADAM) is a variant of stochastic gradient descent that performs well in many problems.

```
from torch.optim import Adam
optimizer = Adam(
    model.parameters(),
    lr=0.001)

# Compute gradients and update weights
ce_loss.backward()
optimizer.step()
```

# Summary

## Feedforward neural networks for image classification

- Convert pixel inputs into class probabilities
- The number of output units determines the number of scores that we can produce
- The Adam optimizer is a variant of SGD that works well in practice

## Classification pipeline

- We use one-hot encoding when we want to predict a single label per image
- We convert the raw output scores (logits) to probabilities using softmax

## Cross entropy loss

- Measures classification prediction quality (agreement of probabilities in  $\hat{y}$  and  $y$ )
- Used with log-softmax in PyTorch for numerical stability

# Further reading

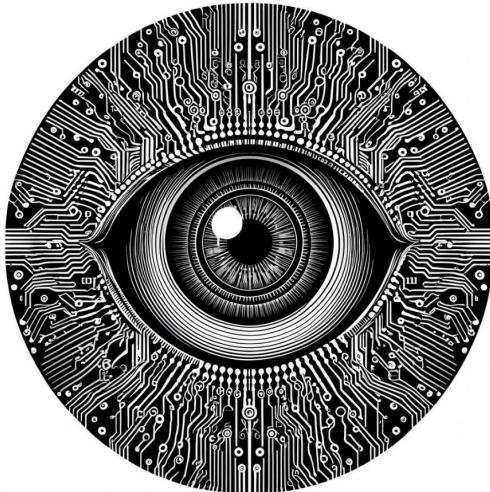
## PyTorch's CrossEntropyLoss

- <https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>

## Who is Adam and what is he optimizing?

- <https://www.youtube.com/watch?v=MD2fYip6QsQ>

# Performance Metrics for Classification and Experiment Tracking



# Learning goals

- Understand the use of performance metrics for classification
- Identify elements in a confusion matrix
- Explore the effect of threshold moving on classification metrics
- Compute accuracy, precision, recall and f1 score for classifiers
- Checkpointing models with the best performance
- Logging experiments into wandb

# Confusion matrices



$$y = P(\text{cat}) = 1.0$$

$$\hat{y} = P(\text{cat}) = 0.95$$

True Positive (TP)  $\Rightarrow y = 1, \hat{y} = 1$

True Negative (TN)  $\Rightarrow y = 0, \hat{y} = 0$

False Positive (FP)  $\Rightarrow y = 0, \hat{y} = 1$

False Negative (FN)  $\Rightarrow y = 1, \hat{y} = 0$

'positive' = 1, 'negative' = 0 (not interchangeable)

$$\begin{bmatrix} \text{TP: 50} & \text{FN: 10} \\ \text{FP: 5} & \text{TN: 100} \end{bmatrix}$$

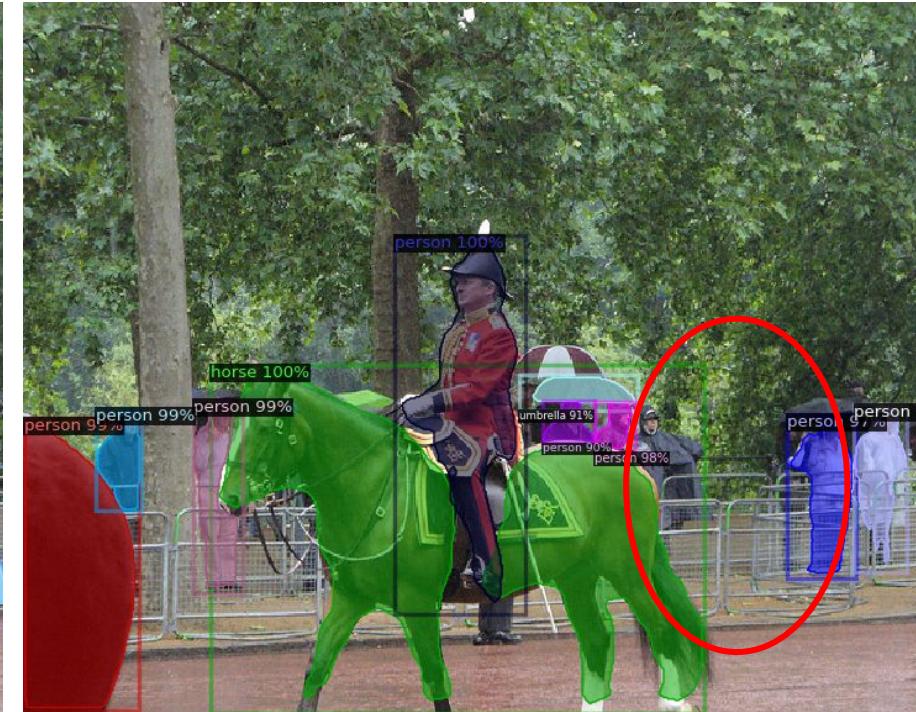
A classifier outputs probabilities, so how do we define that  $y_{\text{hat}} = 1$ ?

Threshold for positive:  $P > 0.5$

# Moving the threshold for $\hat{y} = 1$



Threshold for positive:  $P > 0.5$



Threshold for positive:  $P > 0.9$

# Interpretability with performance metrics



Threshold for positive:  $P > 0.5$

$$\text{accuracy} = \frac{\#\text{correct predictions}}{\#\text{predictions}}$$

assume that a classifier predicts the correct class with confidence 0.8 seven times and with confidence 0.1 three times

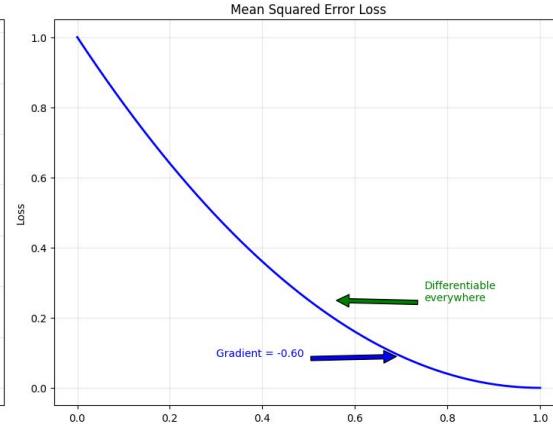
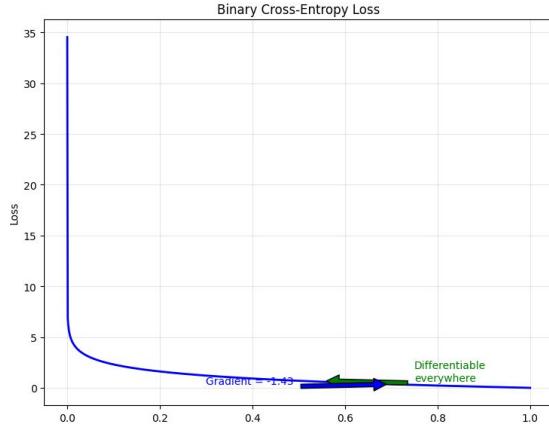
## Cross Entropy Loss

$$-\frac{1}{10} [7 \ln(0.8) + 3 \ln(0.1)]$$

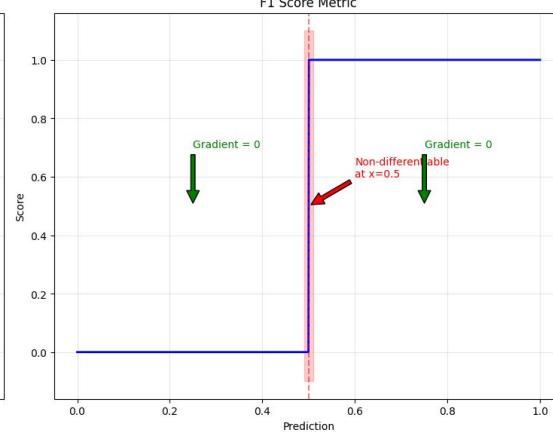
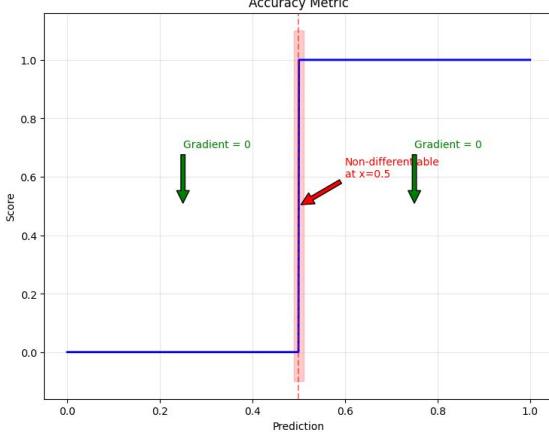
Numerically,

- $\ln(0.8) \approx -0.2231$ , so  $7 \ln(0.8) \approx -1.5617$ .
- $\ln(0.1) \approx -2.3026$ , so  $3 \ln(0.1) \approx -6.9078$ .
- Total is  $-8.4695$ ; divided by 10, it's about 0.847.

# Performance metrics vs loss functions

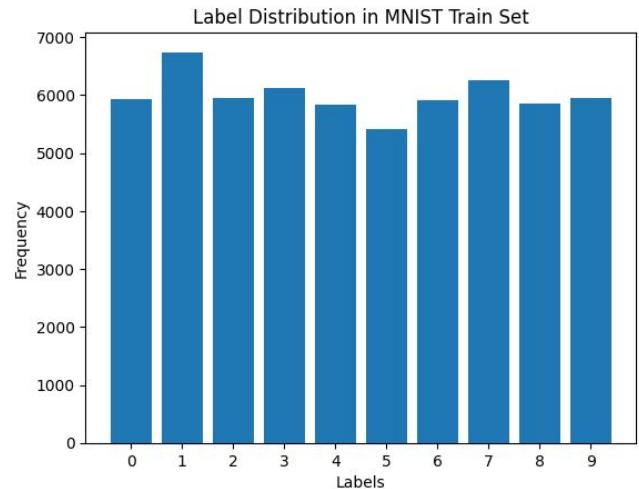


**Gradient Examples:**  
BCE gradient at pred=0.7:  
**-1.429**  
MSE gradient at pred=0.7:  
**-0.600**  
**Accuracy 'gradient' at pred=0.7:**  
**0.000**  
**F1 'gradient' at pred=0.7:**  
**0.000**



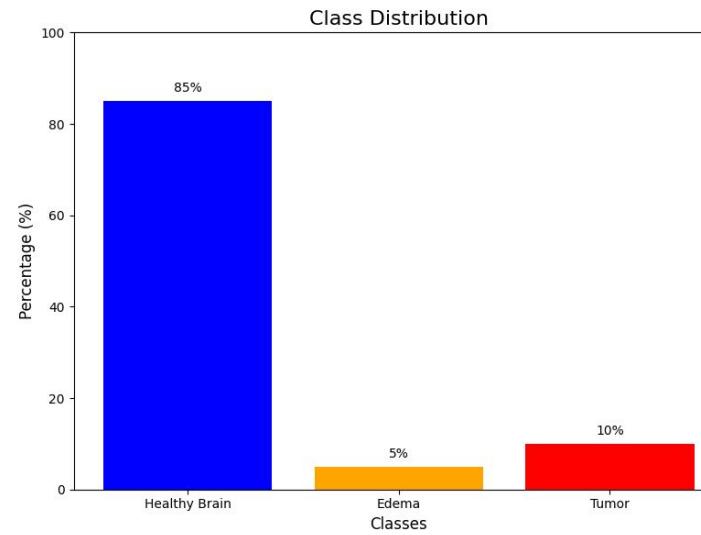
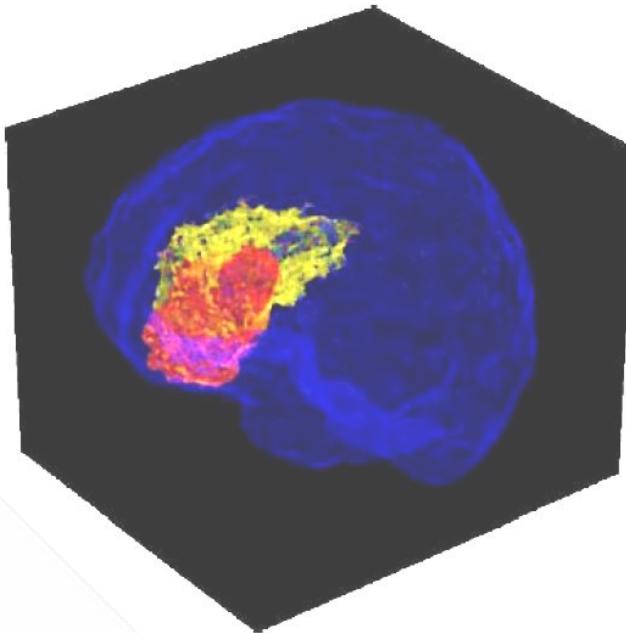
# Multiclass accuracy is only good for balanced datasets

```
def get_batch_accuracy(output, y):
    # Here we give as prediction
    # the label where we have the top confidence
    # No threshold, just top confidence
    pred = probs.argmax(dim=1, keepdim=True)
    correct = pred.eq(y.view_as(pred)).sum().item()
    return correct / probs.shape[0]
```



$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Number of Predictions}} = \frac{TP + TN}{TP + FP + FN + TN}$$

# Accuracy is misleading in imbalanced datasets



'no tumor' in the image,  
majority class prediction

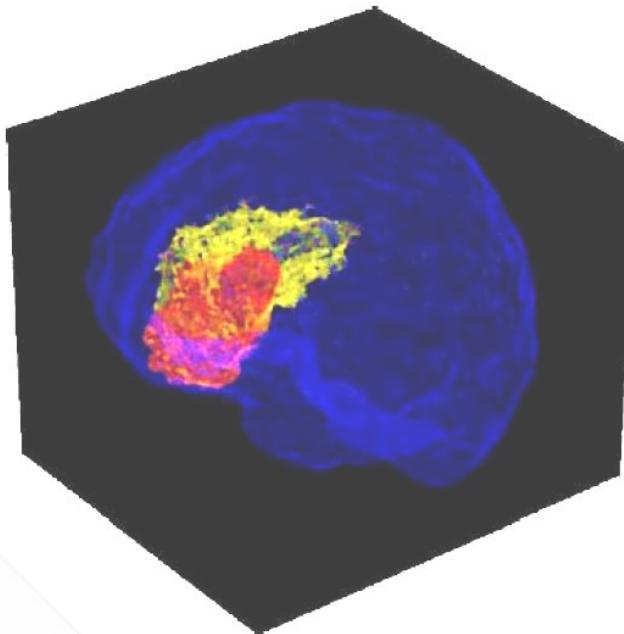
Correct Predictions = 900

Total Predictions = 1000

$$\text{Accuracy} = \frac{900}{1000} = 0.9 \text{ (90\%)}$$

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Number of Predictions}} = \frac{TP + TN}{TP + FP + FN + TN}$$

# Tradeoffs in precision and recall



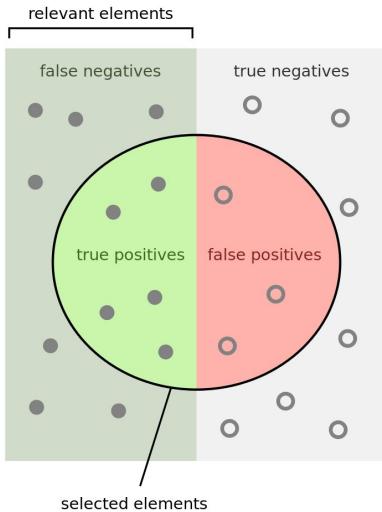
Labeling just one voxel as tumor, with the prediction being correct = **perfect precision**

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Labeling the whole brain as tumor = **perfect recall**

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

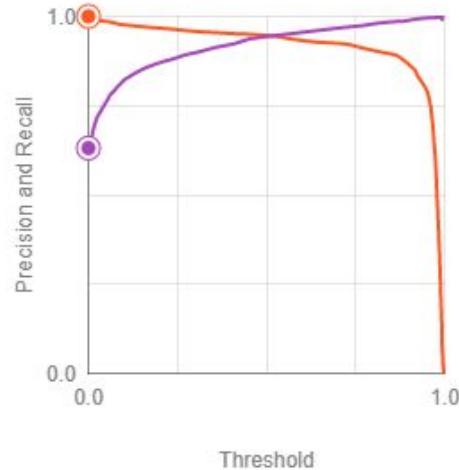
## Thresholds for $y_{\text{hat}} = 1$ , precision, recall, and f1 metrics



How many selected items are relevant?      How many relevant items are selected?

$$\text{Precision} = \frac{\text{How many selected items are relevant?}}{\text{How many relevant items are selected?}}$$
$$\text{Recall} = \frac{\text{How many relevant items are selected?}}{\text{How many relevant items are selected?}}$$

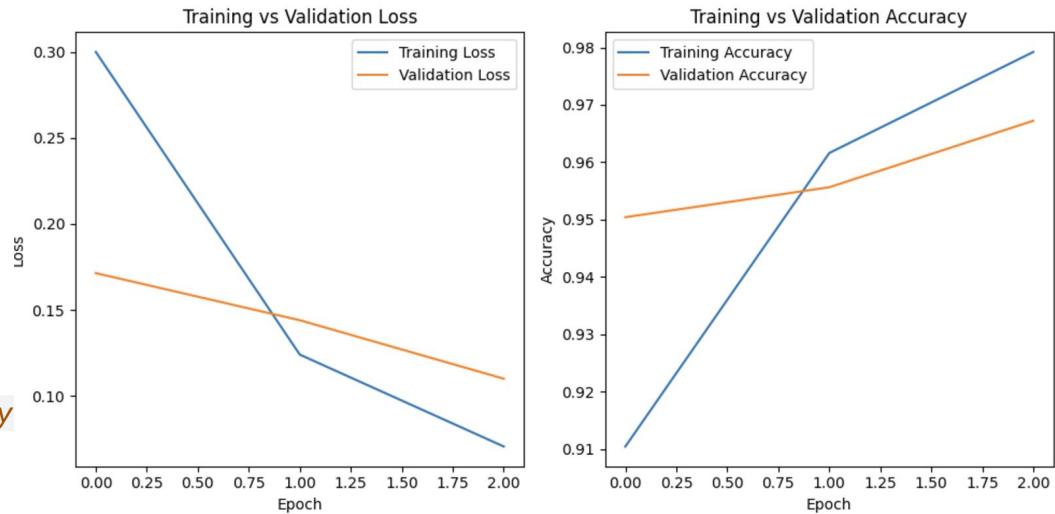
Precision and Recall vs Threshold



$$F_1 = \frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

# Checkpointing models with the best performance

```
for epoch in range(epochs):  
  
    train_loss, train_acc = train()  
    valid_loss, valid_acc = validate()  
  
    # Save model if validation  
    # accuracy improves  
    if valid_acc > best_valid_accuracy:  
        best_valid_accuracy = valid_acc  
  
    # We save the model as a dictionary  
    torch.save(model.state_dict(),  
               'best_model.pth')
```



# Logging experiments and models on wandb

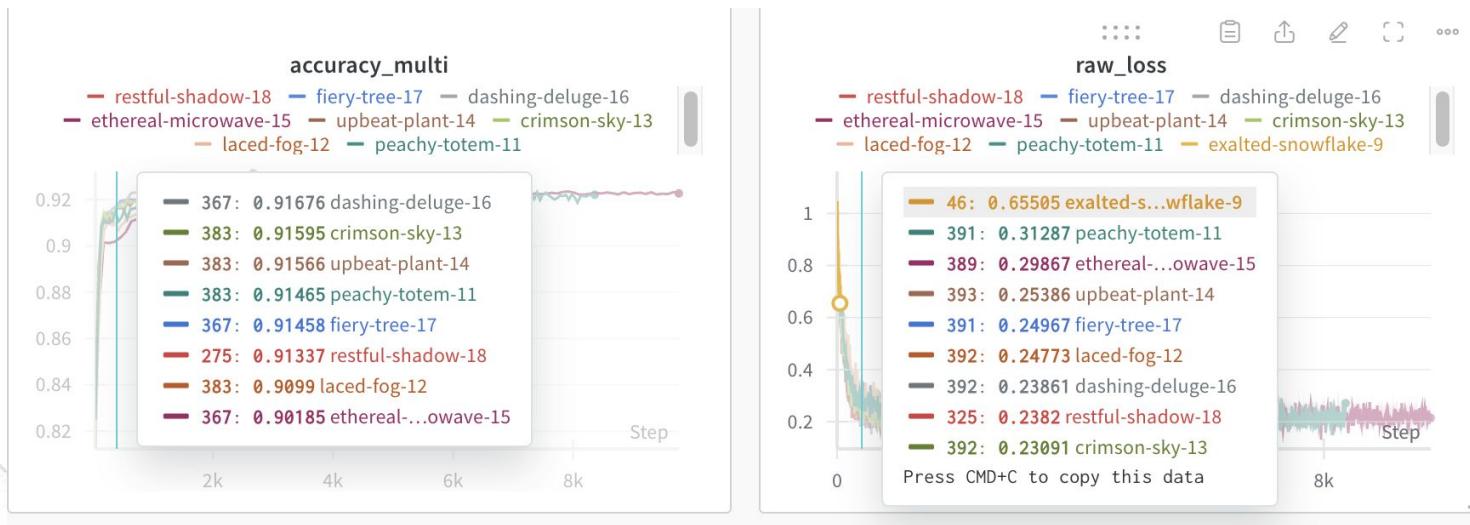


Image from wandb.ai

# Summary

**Confusion matrices are the basis of performance metrics for classification**

- Threshold moving for the positive class affects them and makes the model more precise or sensitive
- The balance of a dataset impacts the metric used: accuracy for balanced datasets, precision and recall for imbalanced ones

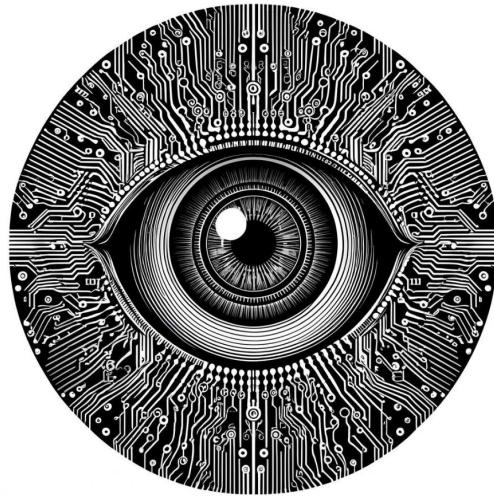
**Performance metrics for classification are non-differentiable**

- We use them for interpretability, not to directly update the weights

**After aligning on a performance metric, we use it to checkpoint our models**

- Experiment tracking allows us to benchmark our training runs

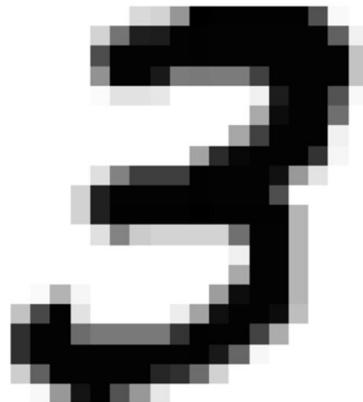
# PyTorch Datasets and Dataloaders



# Learning goals

- Create `Dataset` objects in PyTorch to wrap images and target values together
- Implement `DataLoader` PyTorch objects to feed data to a model
- Understand the connection between `DataLoader` and Stochastic Gradient Descent

# The batch dimension is important for our models



torch image tensors follow the format: [N, C, H, W]  
where:

- N = batch size (number of images)
- C = channels (e.g., 1 for grayscale, 3 for RGB)
- H = height in pixels
- W = width in pixels

```
print(torch_tensor_gray.shape)  
# torch.Size([1, 1, 28, 28])
```

# Feeding torch tensors to a PyTorch model

```
# Define the train loader with batch size and shuffling
batch_size = 32
train_loader = DataLoader(dataset=train_dataset,
                         batch_size=batch_size,
                         shuffle=True)

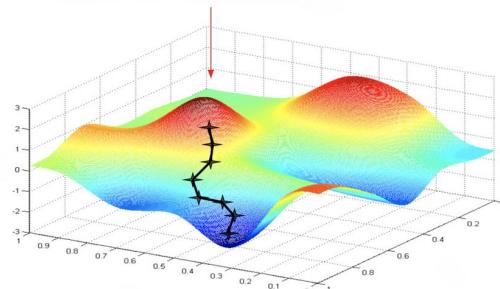
# Move model to device (GPU/CPU)
model = model.to(device)

# Training loop
for images, labels in train_loader:
    # Remember that data has to be explicitly sent to the GPU
    images, labels = images.to(device), labels.to(device)
    output = model(images)
    batch_loss = loss_function(output, labels)
```

**Tip: use the nvidia-smi bash command to check the RAM available in your GPU**

# Shuffling the training set only

```
# Create data loaders
batch_size = 32
# Notice that we shuffle the training loader, but not the validation or test loaders.
# This practice of shuffling the training set is one of the techniques
# that have been shown to improve training.
train_loader = DataLoader(train_subset, batch_size=batch_size, shuffle=True)
valid_loader = DataLoader(valid_subset, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_set, batch_size=batch_size, shuffle=False)
```



$$w_{ij} = w_{ij} - (\text{learning rate} * \frac{dL}{dw_{ij}})$$

# Dataloader depends on Dataset

```
# Define the MNIST Dataset class as a subclass of Dataset
class MNISTDataset(Dataset):
    # Content of the dataset
    # gets called on dataset_instance = MNISTDataset(<params>)
    def __init__(self, path_to_csv, labels=True, transform=None):
        ...
    def __len__(self):
        # What gets called on len(dataset_instance)
        return len(self.data)

    def __getitem__(self, idx):
        # What gets called on dataset_instance[index]
        ...
        return image_tensor, label_tensor
```

# Practice creating a Dataset



KAGGLE · GETTING STARTED PREDICTION COMPETITION · ONGOING

## Digit Recognizer

Learn computer vision fundamentals with the famous MNIST data

<https://www.kaggle.com/competitions/digit-recognizer>

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99	100																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																
4	6	3	2	4	8	1	5	6	6	7	5	2	4	6	1	3	4	7	7	9	6	0	3	1	9	8	3	2	0	7	8	2	1	2	9	4	6	8	3	2	4	7	1	1	4	7	8	0	0	4	9	5	4	6	3	4	3	6	5	6	1	2	7	9	4	0	3	0	6	1	7	3	1	9	1	1	7	3	7	2	0	9	5	5	4	3	1	4	9	2	7	5	2	0	4	7	7	5	9	6	3	9	0	2	2	6	3	1	0	1	5	6	5	7	3	8	3	9	8	1	6	1	9	2	2	4	8	6	2	1	1	3	3	2	4	6	0	3	5	7	9	6	3	2	5	4	8	4	3	9	2	0	1	8	1	9	6	2	7	2	4	4	2	4	1	7	0	2	8	4	9	0	5	3	5	8	1	2	8	6	0	3	9	2	8	6	1	3	5	5	0	0	0	8	2	6	9	2	5	7	4	1	6	7	5	9	8	4	6	1	0	9	3	2	9	4	4	0	6	9	6	9	4	7	8	0	1	3	7	9	7	8	3	3	0	5	5	2	0	8	2	5	2	9	0	5	1	7	4	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0	8	4	5	7	8	7	2	4	0	4	4	5	1	8	9	3	2	2	4	3	2	7	4	2	0	3	9	4	8	1	7	8	0	1	2	0	3	7	9	1	5	3	3	2	8	9	1	5	0	1	6	3	1	1	6	7	1	4	1	2	2	8	3	1	2	3	5	7	4	3	4	4	3	4	1	2	6	6	5	4	1	9	4	7	8	1	8	4	9	4	2	3	4	7	7	5	7	6	5	8	8	9	7	5	4	7	8	9	1	4	5	9	6	1	3	9	5	9	2	6	0

# References

## Datasets and Dataloaders

- [https://pytorch.org/tutorials/beginner/basics/data\\_tutorial.html](https://pytorch.org/tutorials/beginner/basics/data_tutorial.html)

## MNIST on PyTorch's datasets

- <https://pytorch.org/vision/0.20/generated/torchvision.datasets.MNIST.html>

## MNIST on Fiftyone

- <https://try.fiftyone.ai/datasets/mnist/samples>

# Summary

## Dataset objects wrap images and labels together

- Provide a standard way to access data through `__getitem__` and `__len__` methods

## Dataloaders split datasets in batches

- Dataloaders shuffle the training data and maintain sequential order for validation and test sets

## Dataloaders are fundamental to implement Stochastic Gradient Descent

- Their random sampling and shuffling of training samples provide the ‘stochastic’ part of Stochastic Gradient Descent

# Resources

- [Github Repository](#)
- [YouTube playlist](#)
- [Discord channel](#)

**#practical-computer-vision-workshops**