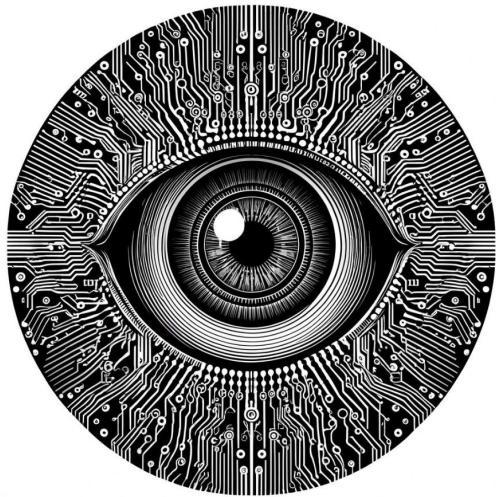


Workshop 9 - Regularization with Data Augmentation



Antonio Rueda-Toicen

About me

- AI Researcher at [Hasso Plattner Institute](#), AI Engineer & DevRel for [Voxel51](#)
- Organizer of the [Berlin Computer Vision Group](#)
- Instructor at [Nvidia's Deep Learning Institute](#) and Berlin's [Data Science Retreat](#)
- Preparing a [MOOC for OpenHPI](#) (now open for registration)



[LinkedIn](#)

How to use our Discord channel during the workshop

- Our channel is **#practical-computer-vision-workshops**. Please ask all questions there instead of the Zoom chat. Through Discord we can have better and more detailed discussions.
- **Step 1 - Use the Discord invite on the Voxel51 website**
<https://discord.com/invite/fiftyone-community>
- **Step 2 - Access channel** (use the direct link below or search)
<http://bit.ly/3YmvPXG>

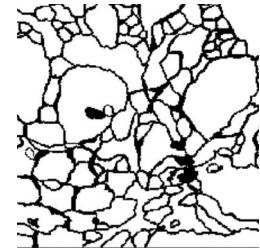
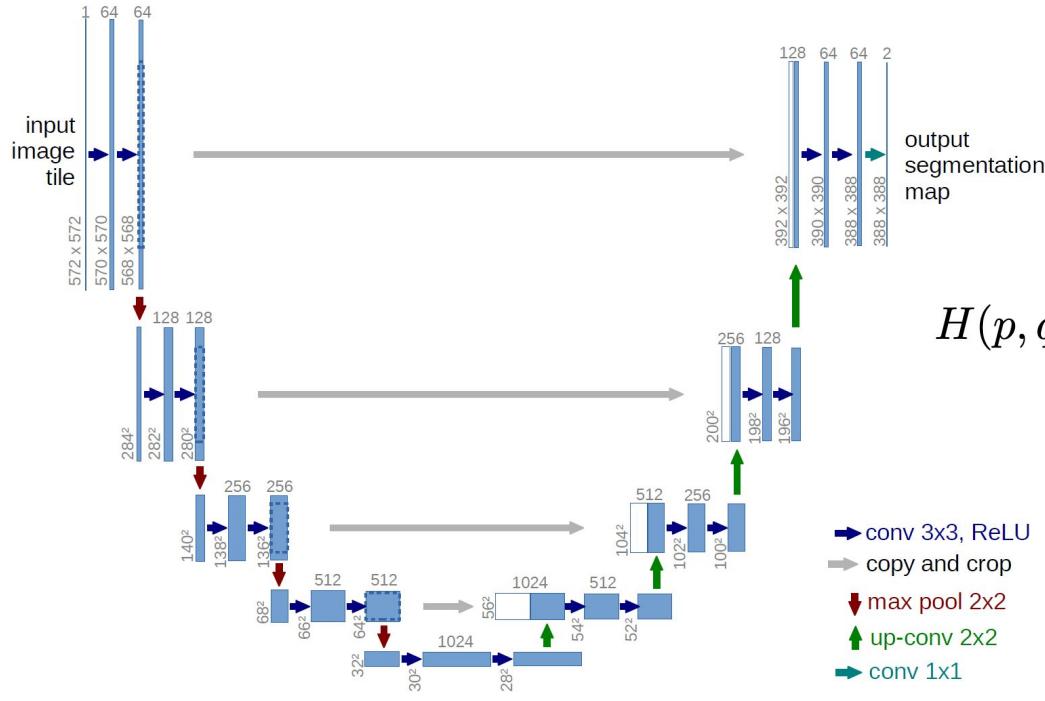
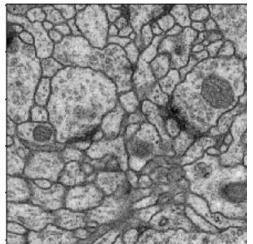
Agenda

- Upsampling and channel mixing with convolutions (review)
- Image data augmentation

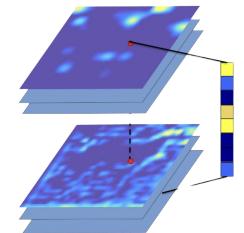
Notebooks

- Semantic segmentation with U-net (***now with augmentations! :D**)
 - Implementation with albumentations
 - Implementation with PyTorch's transforms.v2

Semantic segmentation with U-Net



$$H(p, q) = - \sum_i p(i) \log q(i)$$



U-Net: Convolutional Networks for Biomedical Image Segmentation

Segmentation of neuronal structures in EM stacks challenge - ISBI 2012



The content of this page has not been vetted since shifting away from MediaWiki. If you'd like to help, check out the [how to help guide!](#)

The ISBI 2012 challenge server has been decommissioned, but you can still [download the training and test data!](#)



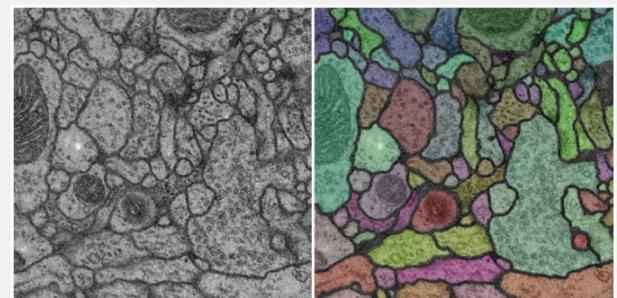
The [IEEE International Symposium on Biomedical Imaging \(ISBI\)](#) is the premier forum for the presentation of technological advances in theoretical and applied biomedical imaging and image computing. In the context of the [ISBI 2012](#) conference which was held in Barcelona, Spain, from 2 to 5 May 2012, we organized a [challenge workshop](#) on segmentation of neuronal structures in EM stacks.

In this [challenge](#), a full stack of EM slices will be used to train [machine learning](#) algorithms for the purpose of automatic segmentation of neural structures.

The images are representative of actual images in the real-world, containing some noise and small image alignment errors. None of these problems led to any difficulties in the manual labeling of each element in the image stack by an [expert human neuroanatomist](#). The aim of the challenge is to compare and rank the different competing methods based on their pixel and object classification accuracy.

Relevant dates

- Deadline for submitting results: February 1st / March 1st, 2012
- Notification of the evaluation: February 21st / March 2nd, 2012

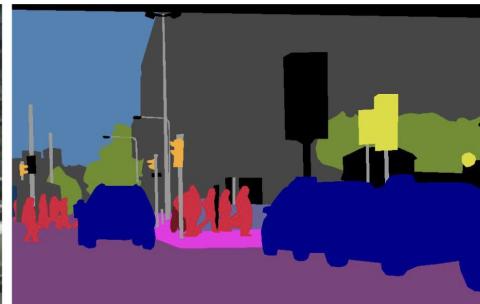


Example of ssTEM image and its corresponding segmentation

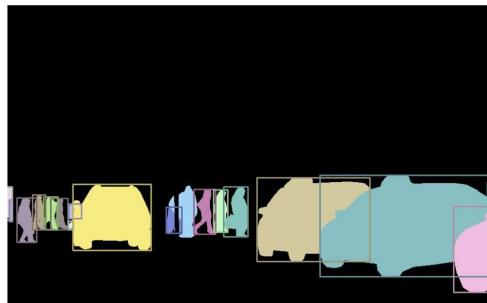
Semantic, Instance, and Panoptic Segmentation



(a) image



(b) semantic segmentation



(c) instance segmentation



(d) panoptic segmentation

Image from [Panoptic Segmentation](#)

Affine

 Augmentation to apply affine transformations to images.

Note: Transform parameters are image size-dependent. For accurate results, test with your own images or similar-sized ones.

[View Transform Details](#)

[Upload Image](#)

[Execute Transform](#)

Scale ⓘ

```
[  
    0.5,  
    2  
]
```

Translate Percent ⓘ

```
[  
    -0.05,  
    0.05  
]
```

Translate Px ⓘ

```
null
```

Rotate ⓘ

[-45, 45]

Shear ⓘ

```
[  
    -15,  
    15  
]
```

Interpolation ⓘ

cv2.INTER_LINEAR ▾

Original

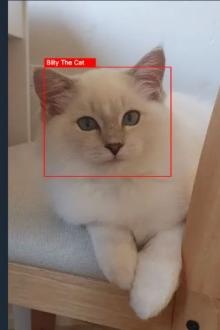
Image 733x484



Mask 733x484



BBoxes 733x484



Keypoints 733x484



Transformed

Image 733x484



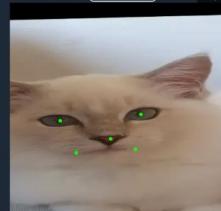
Mask 733x484



BBoxes 733x484



Keypoints 733x484



<https://explore.albumations.ai/transform/Affine>

ElasticTransform

 Apply elastic deformation to images, masks, bounding boxes, and keypoints.

Note: Transform parameters are image size-dependent. For accurate results, test with your own images or similar-sized ones.

[View Transform Details](#)

[Upload Image](#)

[Execute Transform](#)

Alpha ⓘ

300

Sigma ⓘ

10

Interpolation ⓘ

cv2.INTER_LINEAR ▾

Approximate ⓘ

False ▾

Same Dxdy ⓘ

True ▾

Mask Interpolation ⓘ

cv2.INTER_NEAREST ▾

Noise Distribution ⓘ

gaussian ▾

Keypoint Remapping Method ⓘ

mask ▾

Border Mode ⓘ

cv2.BORDER_CONSTANT ▾

Fill ⓘ

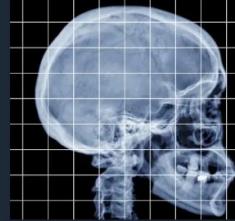
0

Fill Mask ⓘ

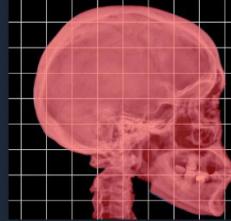
0

Original

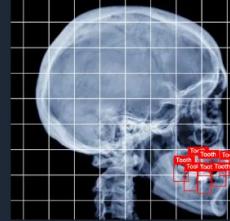
Image 555x600



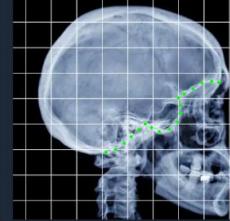
Mask 555x600



BBoxes 555x600

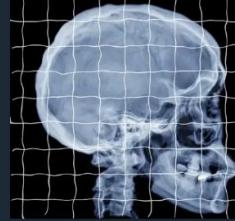


Keypoints 555x600

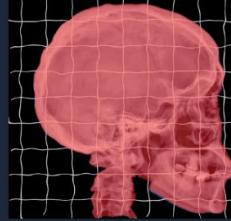


Transformed

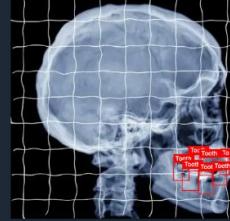
Image 555x600



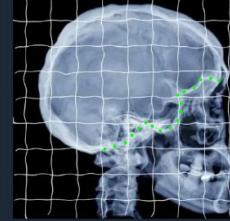
Mask 555x600



BBoxes 555x600

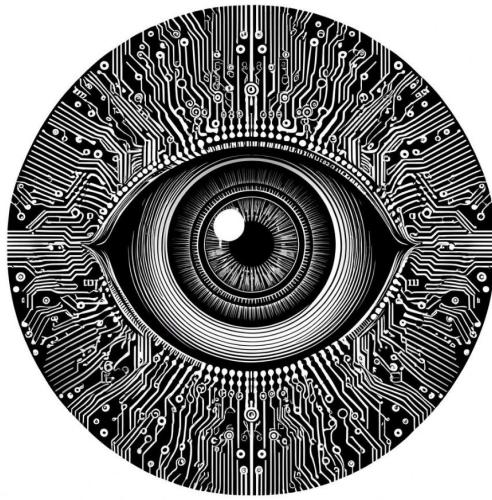


Keypoints 555x600



<https://explore.albuminations.ai/transform/ElasticTransform>

Upsampling and Channel Mixing with Convolutions

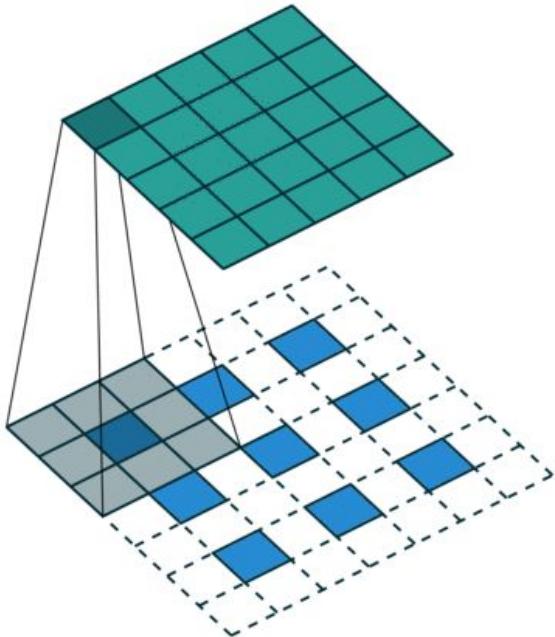


Antonio Rueda-Toicen

Learning goals

- Master upsampling and channel mixing with convolutions
- Use bilinear interpolation to resize images
- Understand hypercolumns and feature combination for image generation

Upsampling filters (aka transposed convolution / upconvolution / atrous convolution / deconvolution)

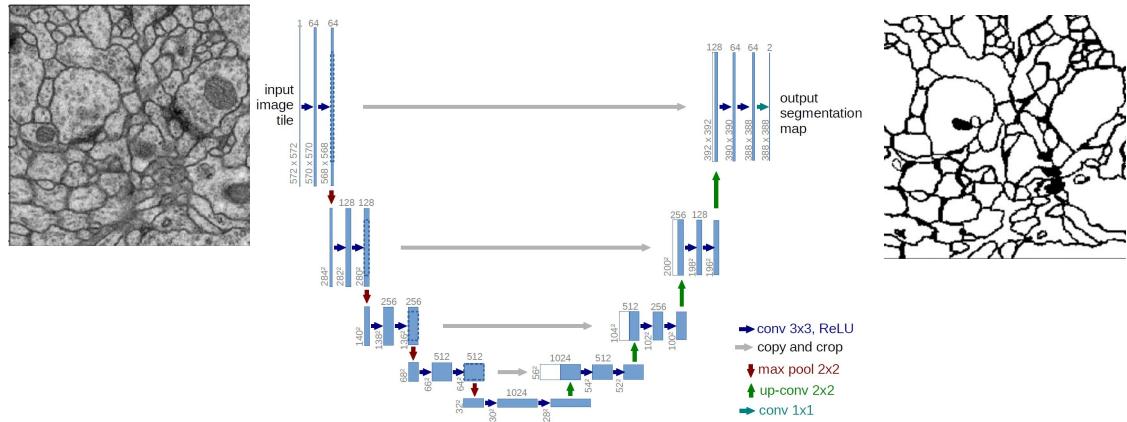
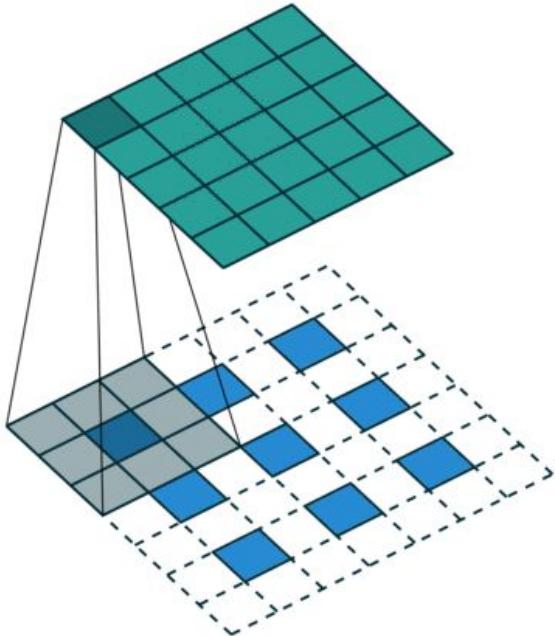


Upconvolution of 3x3 and stride = 1 padded with zeros

Effect: the output is a 5x5 image with 'synthetic' pixels

[Activity: Conv2D in Pytorch \(Colab notebook\)](#)

Usage of 'up-convolution': decoder paths in image generation models

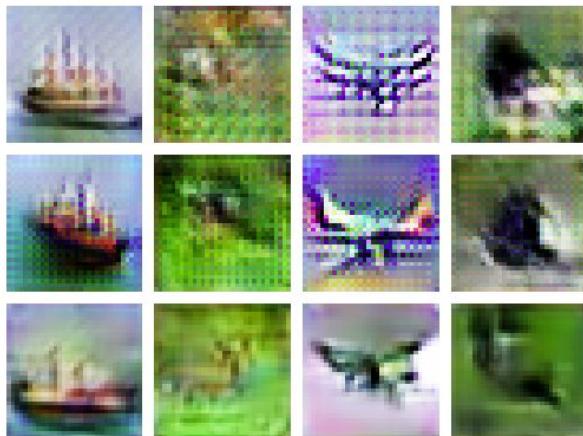


Upconvolution of 3x3 and stride = 1 padded with zeros

Effect: the output is a 5x5 image with 'synthetic' pixels

[Activity: Conv2D in Pytorch \(Colab notebook\)](#)

The checkerboard artifact from “deconvolution”



Deconv in last two layers.
Other layers use resize-convolution.
Artifacts of frequency 2 and 4.

Deconv only in last layer.
Other layers use resize-convolution.
Artifacts of frequency 2.

All layers use resize-convolution.
No artifacts.

Original Image
Size: 224x224



Upscaled Image (Transposed Convolution)
Size: 447x447

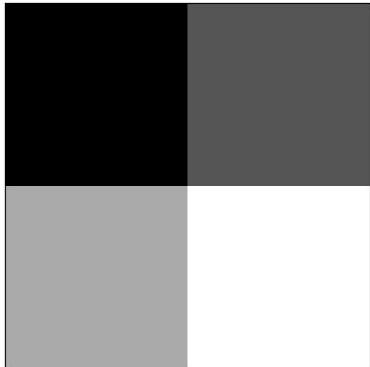


Image from [Deconvolution and Checkerboard artifacts](#)

Problem: “checkerboard” artifacts

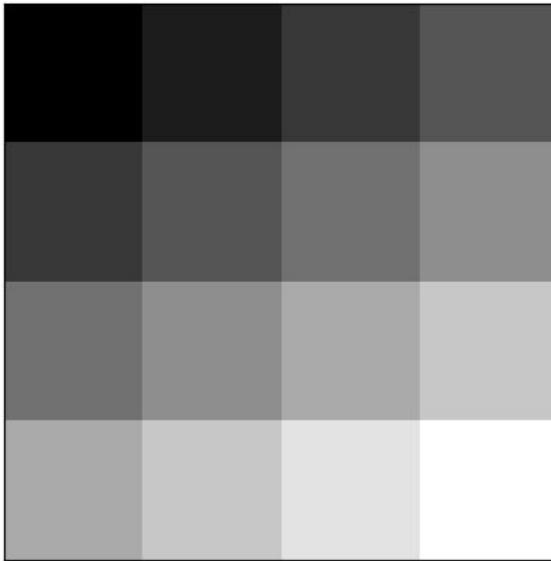
Bilinear interpolation

input



$$Q = \begin{bmatrix} 10 & 20 \\ 30 & 40 \end{bmatrix}$$

output



$$f(x, y) = [1 - w_x \quad w_x] \begin{bmatrix} Q_{11} & Q_{12} \\ Q_{21} & Q_{22} \end{bmatrix} \begin{bmatrix} 1 - w_y \\ w_y \end{bmatrix}$$

$$w_x = \frac{x - x_1}{x_2 - x_1}$$

$$w_y = \frac{y - y_1}{y_2 - y_1}$$

$(x_1, y_1) = (0, 0)$ coordinates of Q_{11}

$(x_2, y_2) = (1, 1)$ coordinates of Q_{22}

➡ $(x, y) = (0.6, 0.7)$ target point

Bilinear interpolation in PyTorch

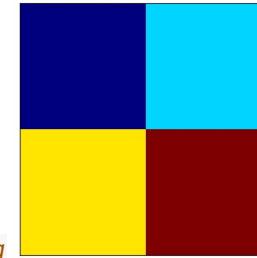
```
import torch
import torch.nn.functional as F

# Create a 1x1x2x2 tensor (batch_size x channels x height x width)
# Notice that the batch size and channel dimensions are created by wrapping
# the height and width tensor with two pairs of extra square brackets
input = torch.tensor([[[[10, 20],
                      [30, 40]]]],
                     dtype=torch.float32)

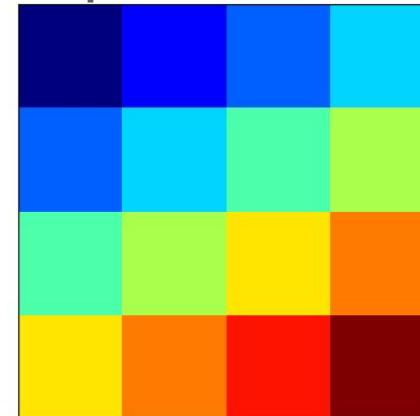
# Upscale to 4x4
output = F.interpolate(input, size=(4, 4), mode='bilinear',
align_corners=True)

import matplotlib.pyplot as plt
# The colormap is just for illustration of corner alignment
plt.imshow(input.squeeze(), cmap="jet")
plt.imshow(output.squeeze(), cmap="jet")
```

input (with colormap)
shape: 2x2



output (with colormap)
shape: 4x4



1x1 convolutions mix image channels

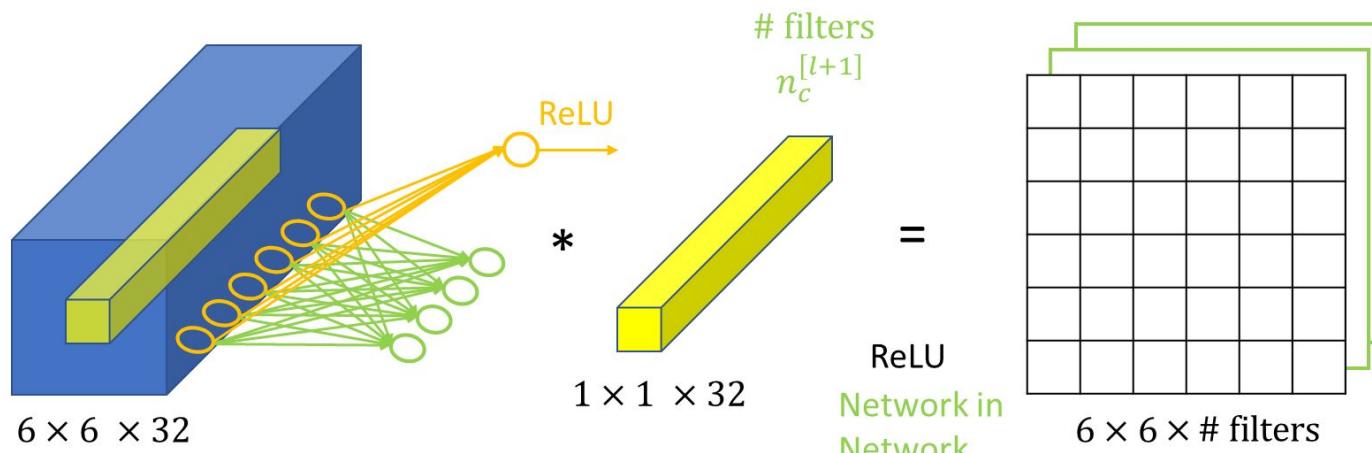
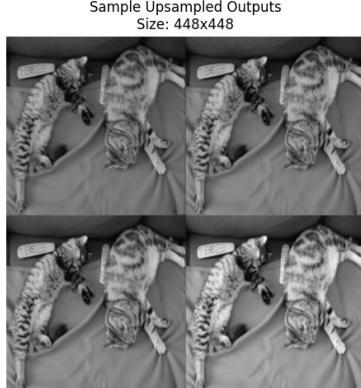


Image from "[What does a 1x1 convolution do?](#)" by Andrew Ng

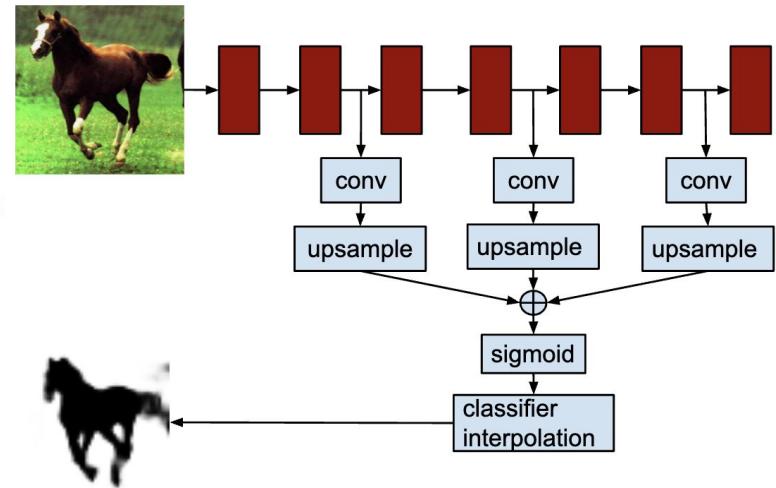
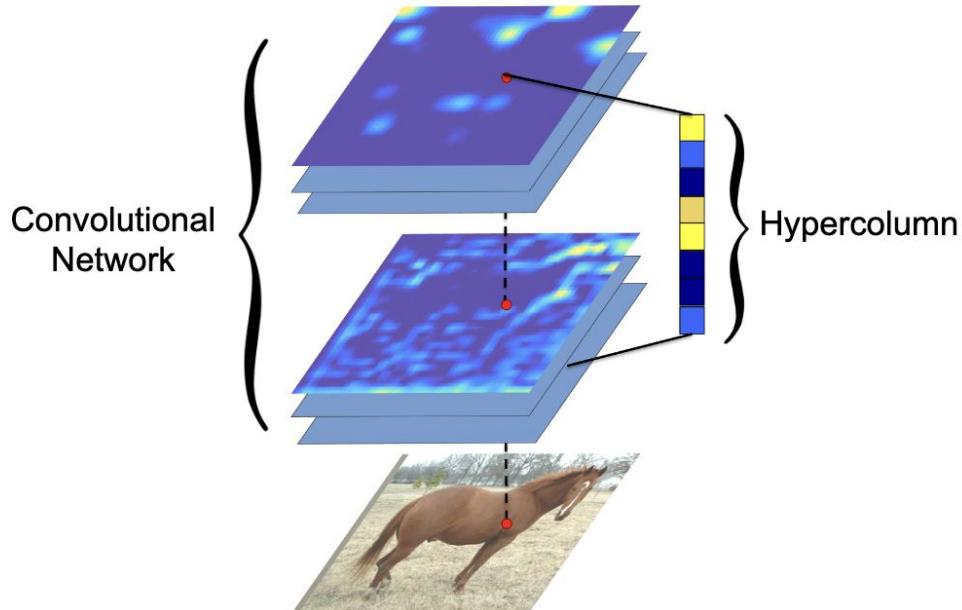
Upsampling and channel mixing with 1x1 convs



```
layers = nn.Sequential(  
    nn.Conv2d(  
        in_channels=1,  
        out_channels=64,  
        kernel_size=3,  
        padding=1,  
        bias=False  
    ),  
    nn.Upsample(  
        scale_factor=2,  
        mode='bilinear',  
        align_corners=True  
    ),  
    nn.Conv2d(  
        in_channels=64,  
        out_channels=1,  
        kernel_size=1,  
        bias=False  
    )  
)
```

Note: no checkerboard artifacts

Combining “hypercolumns” is a common use of 1×1 convolutions



Images from “[Hypercolumns for Object Segmentation and Fine-grained Localization](#)”

Summary

1x1 convolutions mix channel information

- Mixing channels allows us to create new images by combining learned features

Transposed convolutions enable upsampling (with checkerboard artifacts)

- Bilinear interpolation and channel mixing is a better alternative

Further reading and references

A guide to convolution arithmetic for deep learning

- <https://arxiv.org/abs/1603.07285>

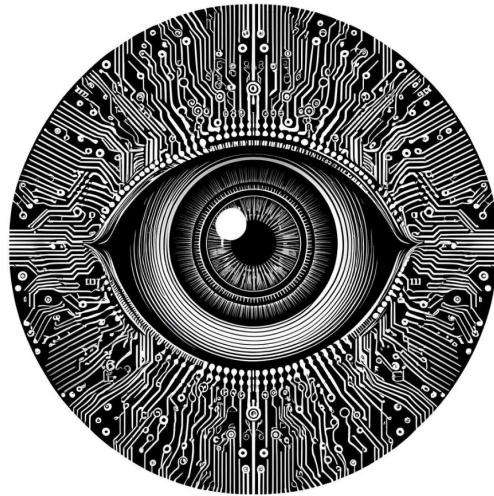
Network in network (1x1 convolutions)

- <https://arxiv.org/abs/1312.4400>

Hypercolumns for object segmentation and fine-grained localization

- https://openaccess.thecvf.com/content_cvpr_2015/papers/Hariharan_Hypercolumns_for_Object_2015_CVPR_paper.pdf

Image Data Augmentation



Antonio Rueda-Toicen

SPONSORED BY THE



Federal Ministry
of Education
and Research

Learning goals

- Use data augmentation to extend training datasets
- Gain familiarity with PyTorch's transforms.v2: affine transformations, vertical and horizontal flipping, and random crops

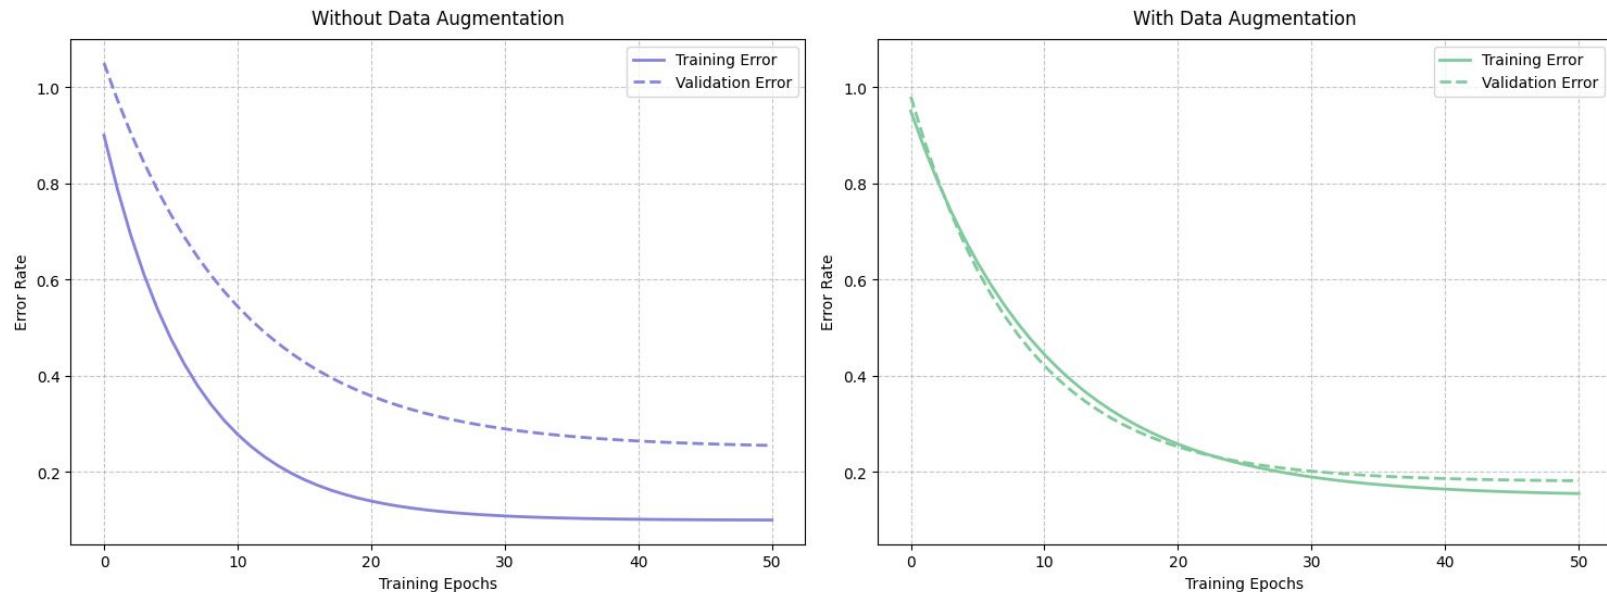
Data augmentation



Image from <https://docs.fast.ai/vision.augment>

Data augmentation as regularization

Impact of Data Augmentation on Model Training



Affine transformations

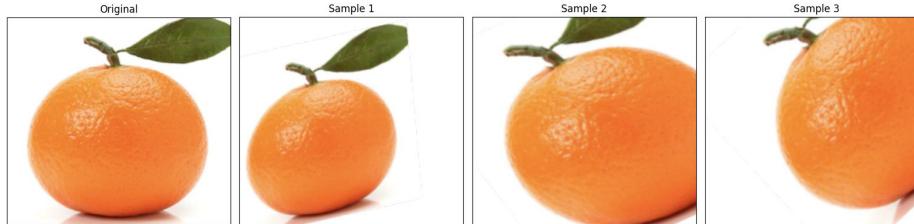
$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} + \begin{pmatrix} \Delta_x \\ \Delta_y \end{pmatrix}$$



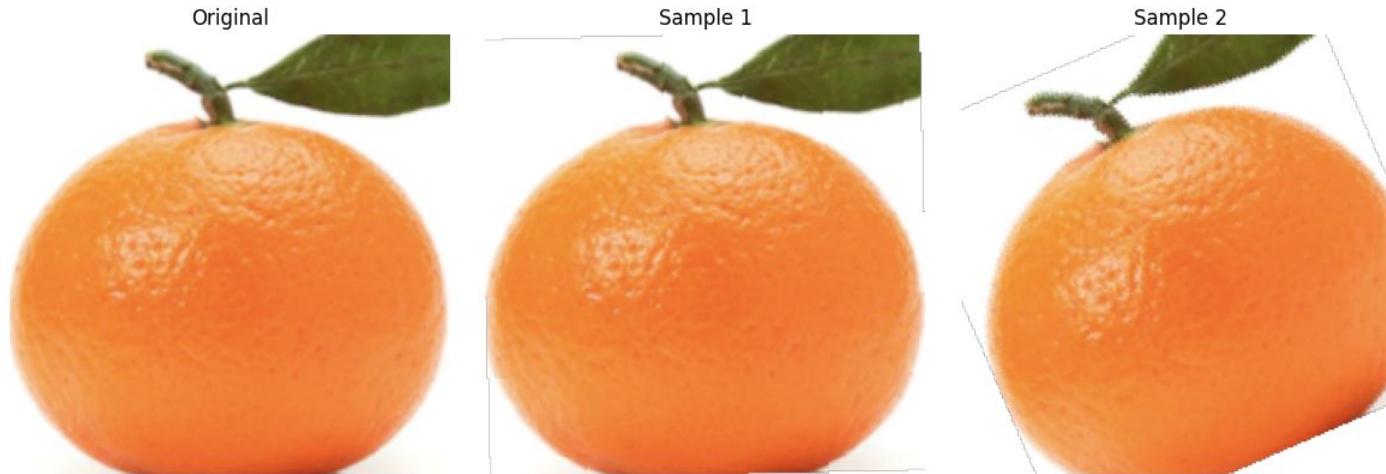
Linear transformation
(Rotation, scaling, shear)



Translation
(Moving along the axes)



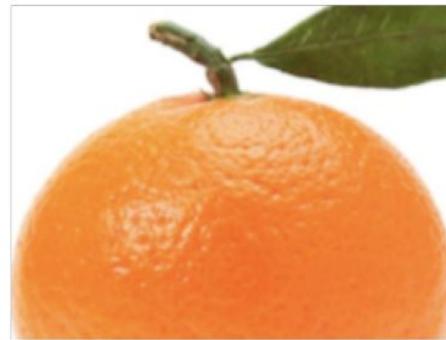
The RandomAffine transformation - rotation



$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

`transforms.RandomAffine(degrees=45)`

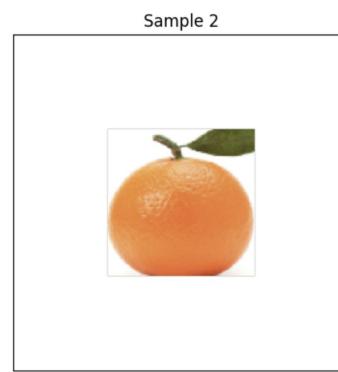
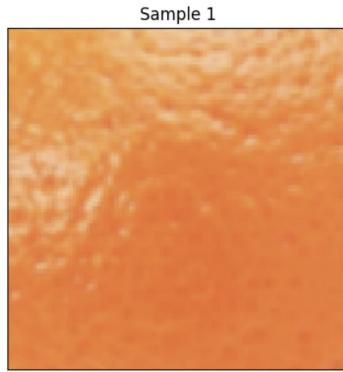
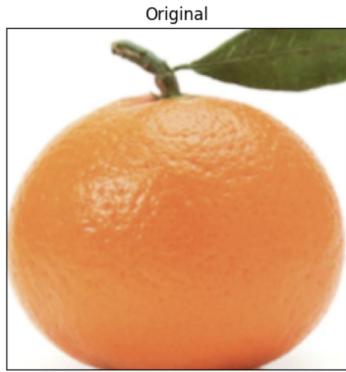
The RandomAffine transformation - translation



$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} + \begin{pmatrix} t_x \\ t_y \end{pmatrix}$$

`transforms.RandomAffine(degrees=0, translate=(0.4, 0.4))`

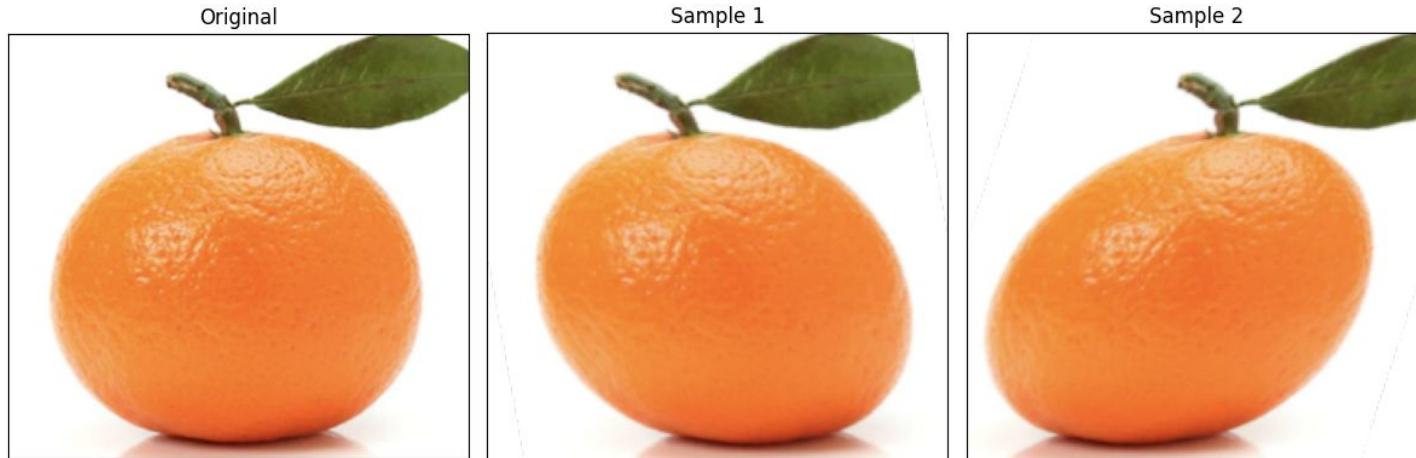
The RandomAffine transformation - scaling



$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} s_x & 0 \\ 0 & s_y \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

```
transforms.RandomAffine(degrees=0, scale=(0.25, 2.5))
```

The RandomAffine transformation - shear



$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} 1 & s_x \\ s_y & 1 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

transforms.RandomAffine(degrees=0, shear=45)

Horizontal and vertical flipping

Original



Sample 1

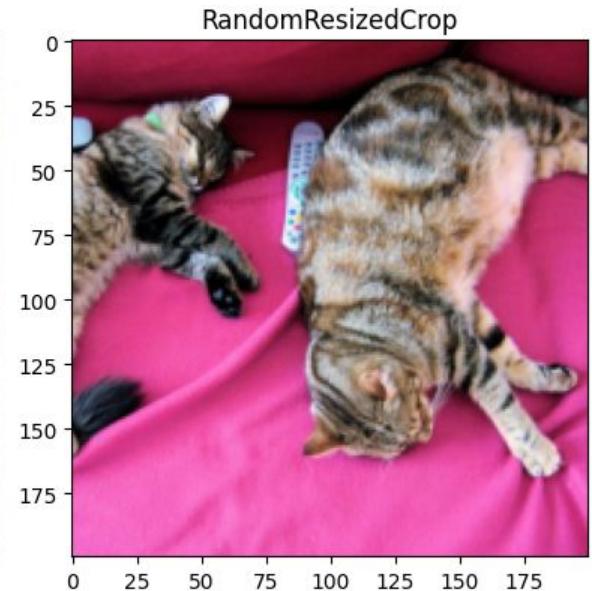
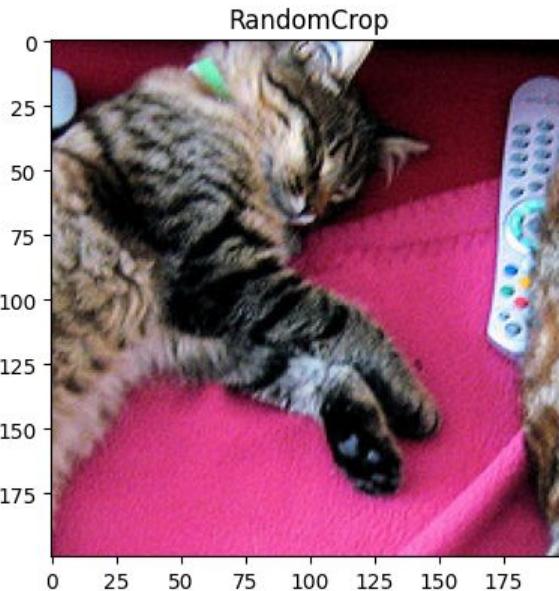
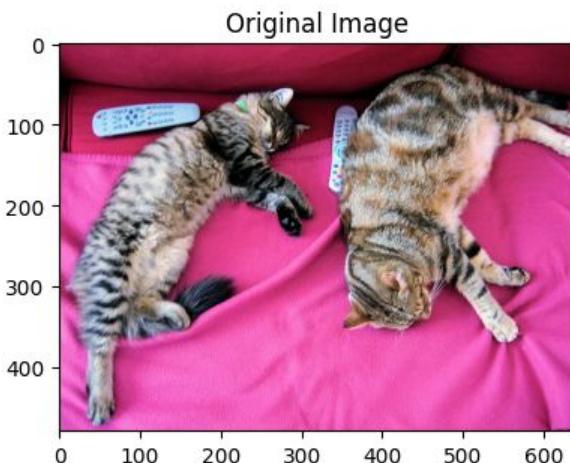


Sample 2



```
transforms.Compose([
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomVerticalFlip(p=0.5)
])
```

RandomCrop and RandomResizedCrop

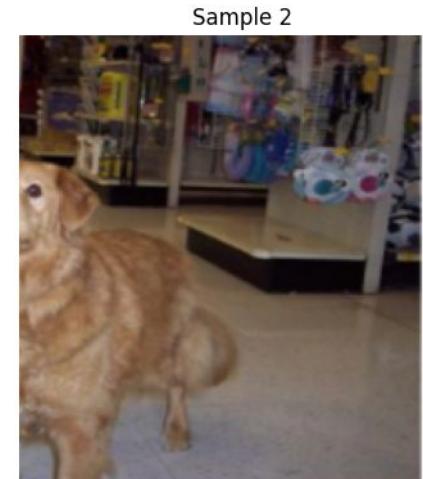


```
transforms.RandomCrop(size=224)
```

```
transforms.RandomResizedCrop(  
    size=224,  
    scale=(0.08, 1.0), # Range of size of crop  
    ratio=(3/4, 4/3)   # Range of aspect ratio  
)
```

Blind data augmentation can damage your model

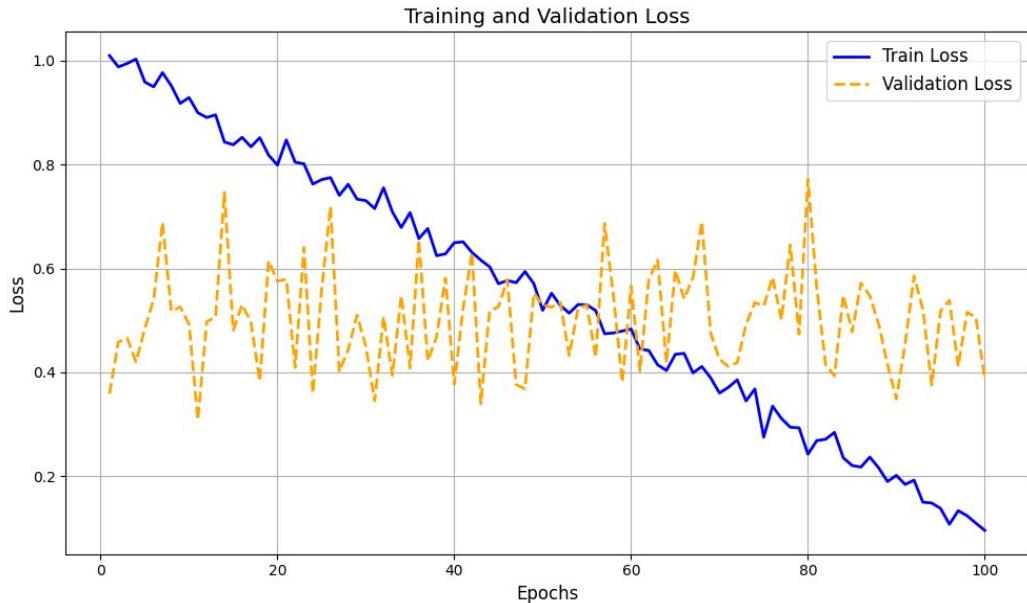
\hat{y} = golden retriever (unit 207 on Imagenet)



$$\text{Cross Entropy Loss} = -\log(0.003) = 5.809$$

$$-\log(0.6) = 0.511$$

The validation set should not be augmented during training

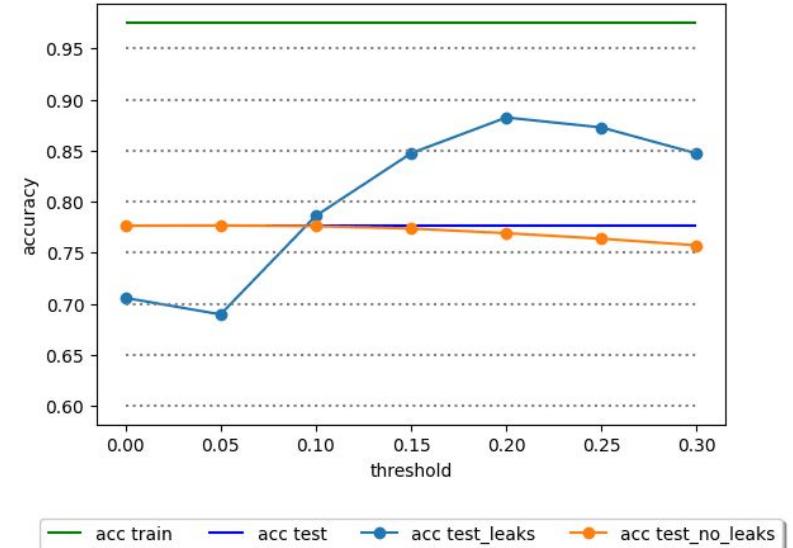


```
import torchvision.transforms.v2 as T

imagenet_mean = [0.485, 0.456, 0.406]
imagenet_std = [0.229, 0.224, 0.225]

transform = T.Compose([
    # Resize to 224x224
    T.Resize((224, 224)),
    # Convert to torch Image
    T.ToImage(),
    # Convert to scaled float tensor
    T.ToDtype(torch.float32, scale=True),
    # Apply ImageNet normalization
    T.Normalize(mean=imagenet_mean,
                std=imagenet_std)
])
```

Augmentation should happen after train, validation, and test splitting



Images from "[On Leaky Datasets and a Clever Horse](#)"

Summary

Data augmentation creates synthetic input data

- Data augmentation can improve the robustness of a model and prevent overfitting

PyTorch transforms are powerful and composable

- `transforms.v2` provides rotation, scale, crop, and flip (+other) operations
- Each transform has specific parameters for fine control

Avoiding data leakage: best practices

- Apply augmentations after train, validation, and test split. Never augment validation data. Test augmentation effects through observation and loss values

Further reading and references

Data augmentation is still data curation

- <https://voxel51.com/blog/data-augmentation-is-still-data-curation/>

Getting started with PyTorch's transforms v2

- https://pytorch.org/vision/main/auto_examples/transforms/plot_transforms_getting_started.html

Illustration of transforms

- https://pytorch.org/vision/main/auto_examples/transforms/plot_transforms_illustrations.html#sphx-glr-auto-examples-transforms-plot-transforms-illustrations-py

Resources

- [Github Repository](#)
- [YouTube playlist](#)
- [Discord channel](#)

#practical-computer-vision-workshops